

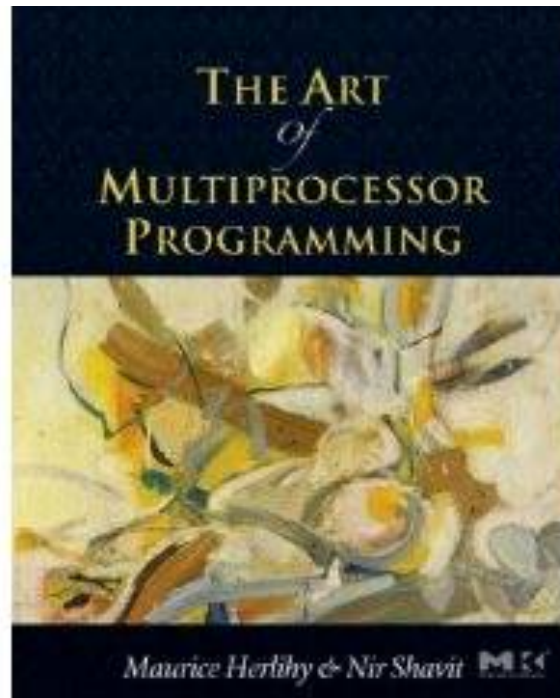
# Concurrent Linked Lists

**Acknowledgement:**

**Slides adopted from the companion slides for the book  
"The Art of Multiprocessor Programming"  
by Maurice Herlihy and Nir Shavit**

# What We'll Cover Today

Chapter 9 of:



Digital copy can be obtained via WUSTL library:

<http://catalog.wustl.edu/search/>

# Today: Concurrent Objects

- Adding threads should not lower throughput
  - Contention effects
  - Mostly fixed by Queue locks
- Should increase throughput
  - Not possible if inherently sequential
  - Surprising things are parallelizable

# Coarse-Grained Synchronization: the Good

- Each method locks the object
  - Avoid contention using queue locks
  - Easy to reason about
    - In simple cases

# Coarse-Grained Synchronization: the Bad

- Sequential bottleneck
  - Threads “stand in line”
- Adding more threads
  - Does not improve throughput
  - Struggle to keep it from getting worse

# This Lecture

- Introduce four “patterns”
  - Bag of tricks ...
  - Methods that work more than once ...
- For highly-concurrent objects
  - Concurrent access
  - More threads, more throughput

# This Lecture

- Coarse-grained locking
- Fine-grained locking
- Optimistic synchronization
- Lazy synchronization
- Lock-free synchronization

# First:

## Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
  - Independently-synchronized components
- Methods conflict when they access
  - The same component ...
  - At the same time



## Second: Optimistic Synchronization

- Search without locking ...
- If you find it, lock and check ...
  - OK: we are done
  - Oops: start over
- Evaluation
  - Usually cheaper than locking, but
  - Mistakes are expensive

# Third:

## Lazy Synchronization

- Postpone hard work
- Removing components is tricky
  - Logical removal
    - Mark component to be deleted
  - Physical removal
    - Do what needs to be done

# Fourth:

## Lock-Free Synchronization

- Don't use locks at all
  - Use `compareAndSet()` & relatives ...
- Advantages
  - No Scheduler Assumptions/Support
- Disadvantages
  - Complex
  - Sometimes high overhead

# Linked List

- Illustrate these patterns ...
- Using a list-based Set
  - Common application
  - Building block for other apps

# Set Interface

- Unordered collection of items
- No duplicates
- Methods
  - **add(x)** put **x** in set
  - **remove(x)** take **x** out of set
  - **contains(x)** tests if **x** in set

# List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x) ;  
    public boolean remove(T x) ;  
    public boolean contains(T x) ;  
}
```

# List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x) ;  
    public boolean remove(T x) ;  
    public boolean contains(T x) ;  
}
```



**Add item to set**

# List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(Tt x);  
}
```

**Remove item from set**



# List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```



Is item in set?

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public volatile Node next;  
}
```

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public volatile Node next;  
}
```



**item of interest**

# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public volatile Node next;  
}
```



**Usually hash code**

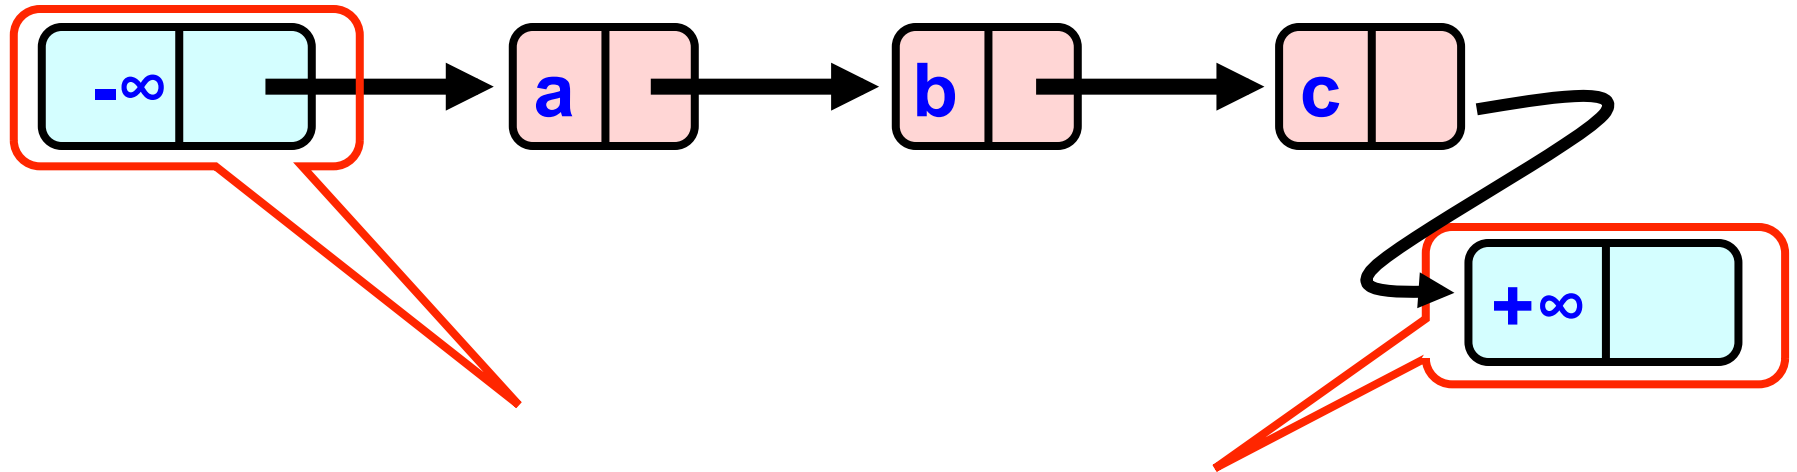
# List Node

```
public class Node {  
    public T item;  
    public int key;  
    public Node next;  
}
```

**Reference to next node**



# The List-Based Set



Sorted with Sentinel nodes  
(min & max possible keys)

Once you find a key larger than the key you are searching for, you are done.

# Reasoning about Concurrent Objects

- Invariant
  - Property that always holds
- Established because
  - True when object is **created**
  - Truth **preserved** by each method
    - Each **step** of each method

# Specifically ...

- Invariants preserved by
  - **add()**
  - **remove()**
  - **contains()**
- Most steps are trivial
  - Usually one step tricky
  - Often linearization point




# Interference

- Invariants make sense only if we assume **freedom from interference**: methods considered are the only modifiers to the data structure.
- Language encapsulation helps
  - List nodes not visible outside class
- Freedom from interference needed even for removed nodes
  - Some algorithms traverse removed nodes
  - Careful with **malloc()** & **free()**!
- We rely on garbage collection

# Abstract Data Types

- Concrete representation:



- Abstract Type:

$\{a, b\}$

# Abstract Data Types

- Meaning of representation given by *abstraction map*, carrying lists that satisfy representation invariant to set.

$$S( \boxed{\phantom{a}} \boxed{\phantom{a}} \rightarrow \boxed{a} \boxed{\phantom{a}} \rightarrow \boxed{b} \boxed{\phantom{a}} \rightarrow \boxed{\phantom{a}} \boxed{\phantom{a}} ) = \{a, b\}$$

# Representation Invariant

- Which concrete values meaningful?
  - Sorted?
  - Duplicates?
- Rep invariant
  - Characterizes legal concrete reps
  - Preserved by methods
  - Relied on by methods

# Representation Invariant

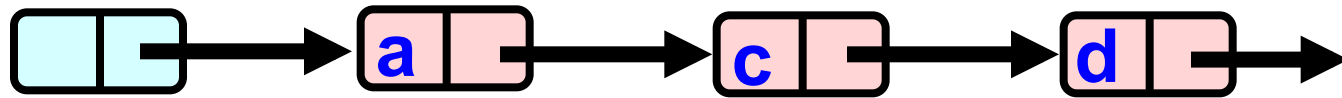
- Sentinel nodes
  - tail reachable from head
- Sorted
- No duplicates

# Abstraction Map

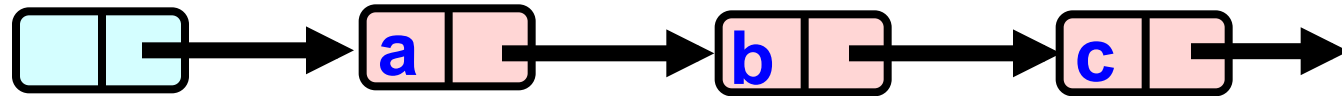
- $S(\text{head}) =$ 
  - $\{ x \mid \text{there exists } a \text{ such that}$ 
    - $a \text{ reachable from head and}$
    - $a.\text{item} = x$
  - $\}$

# Sequential List Based Set

add()

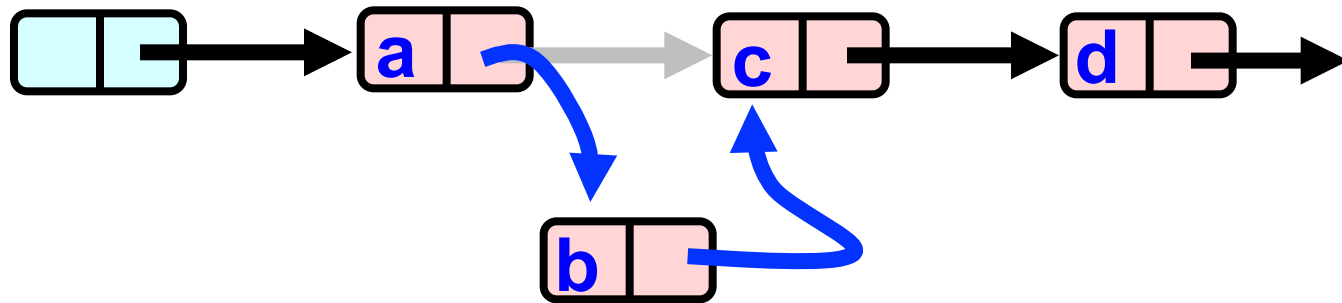


remove()

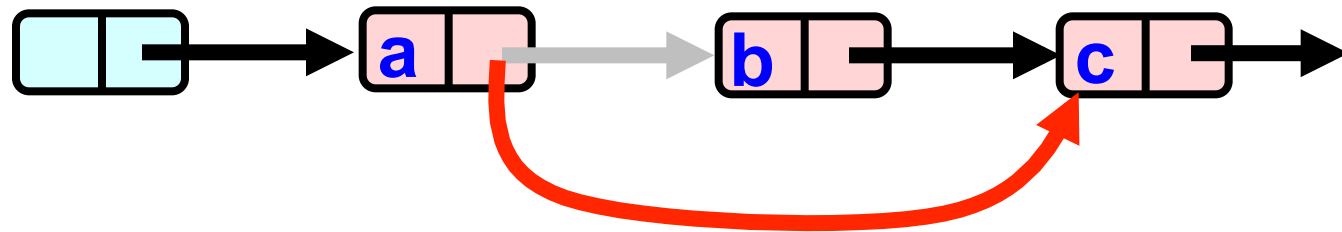


# Sequential List Based Set

add()



remove()





# Coarse-Grained Locking

- Easy, same as synchronized methods
  - "One lock to rule them all ... "
- Simple, clearly correct
  - Deserves respect!
- Works poorly with contention
  - Queue locks help
  - But bottleneck still an issue

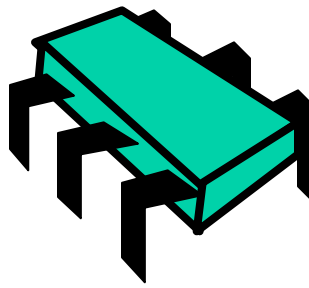
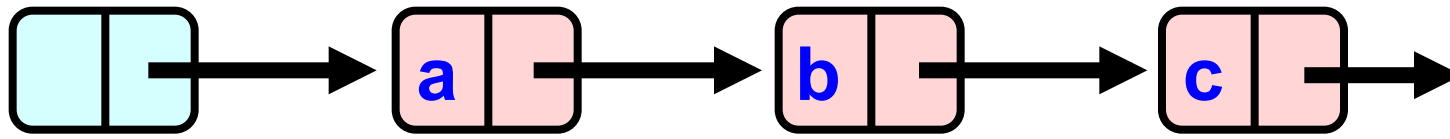
# Fine-grained Locking

- Requires **careful** thought
  - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”

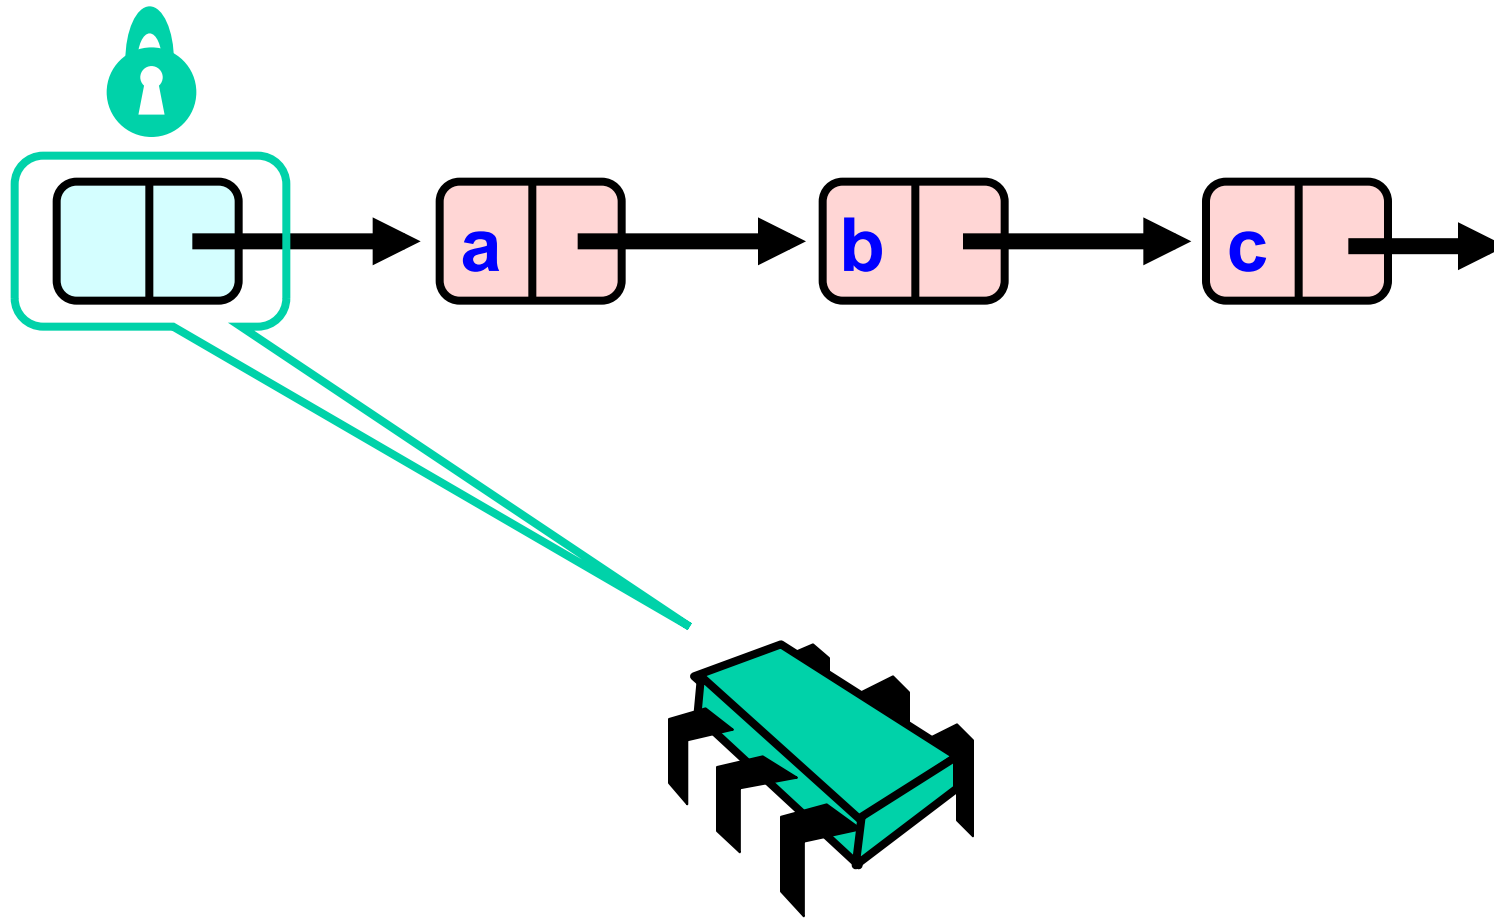
# Fine-grained Locking

- Requires **careful** thought
  - “Do not meddle in the affairs of wizards, for they are subtle and quick to anger”
- Split object into pieces
  - Each piece has own lock
  - Methods that work on disjoint pieces need not exclude each other

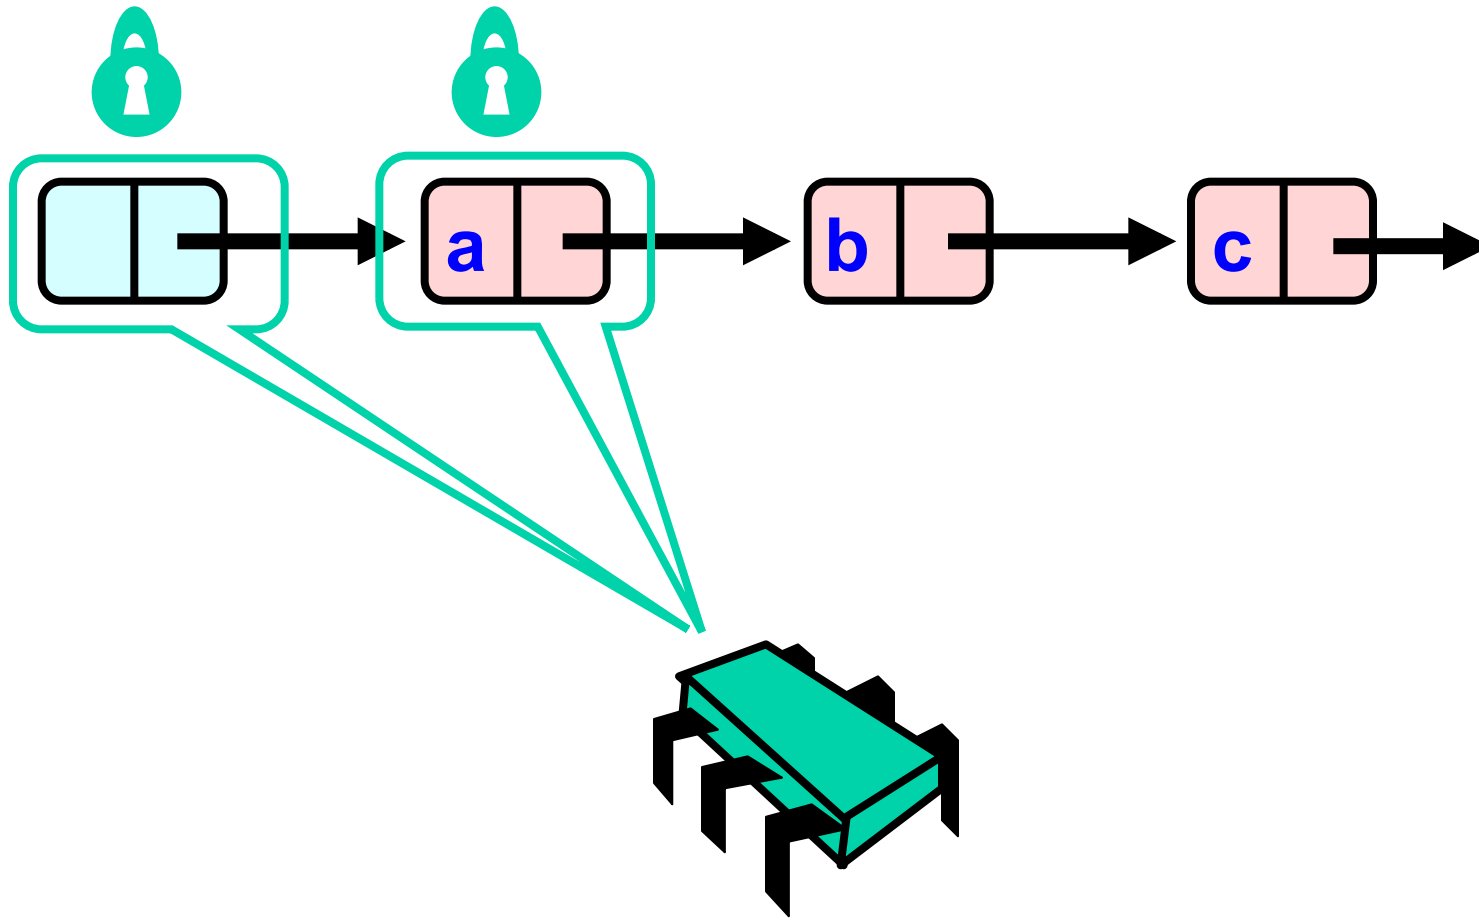
# Hand-over-Hand locking




# Hand-over-Hand locking



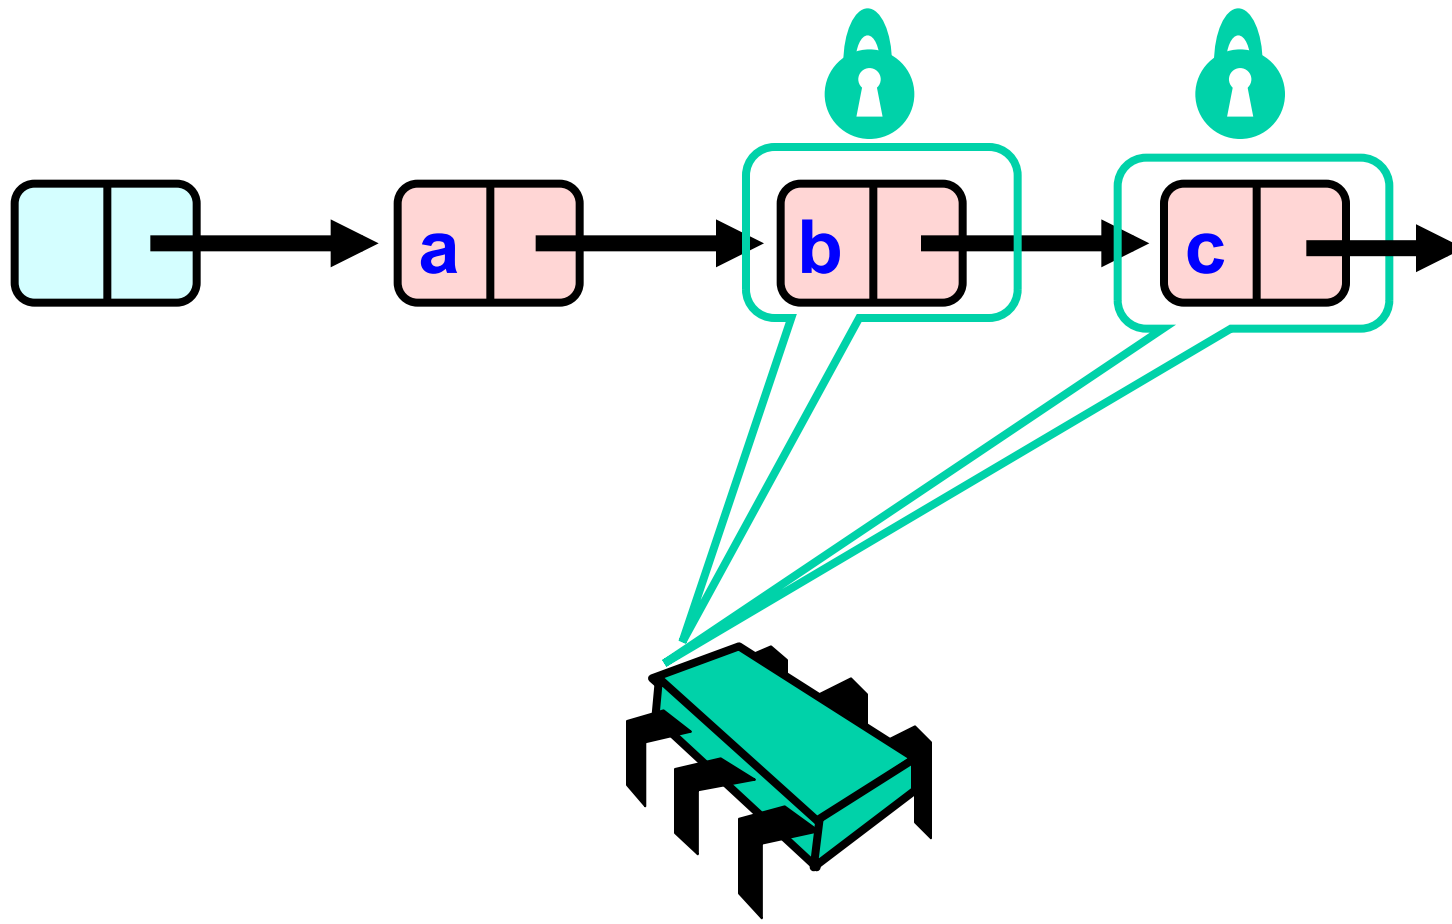
# Hand-over-Hand locking



# Hand-over-Hand locking

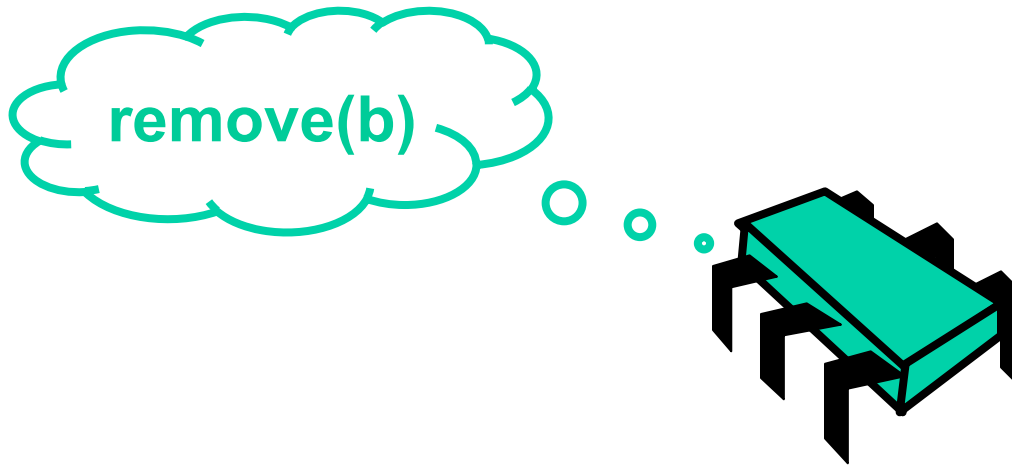
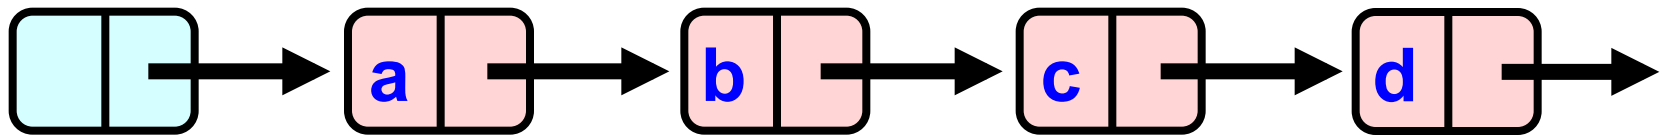


# Hand-over-Hand locking

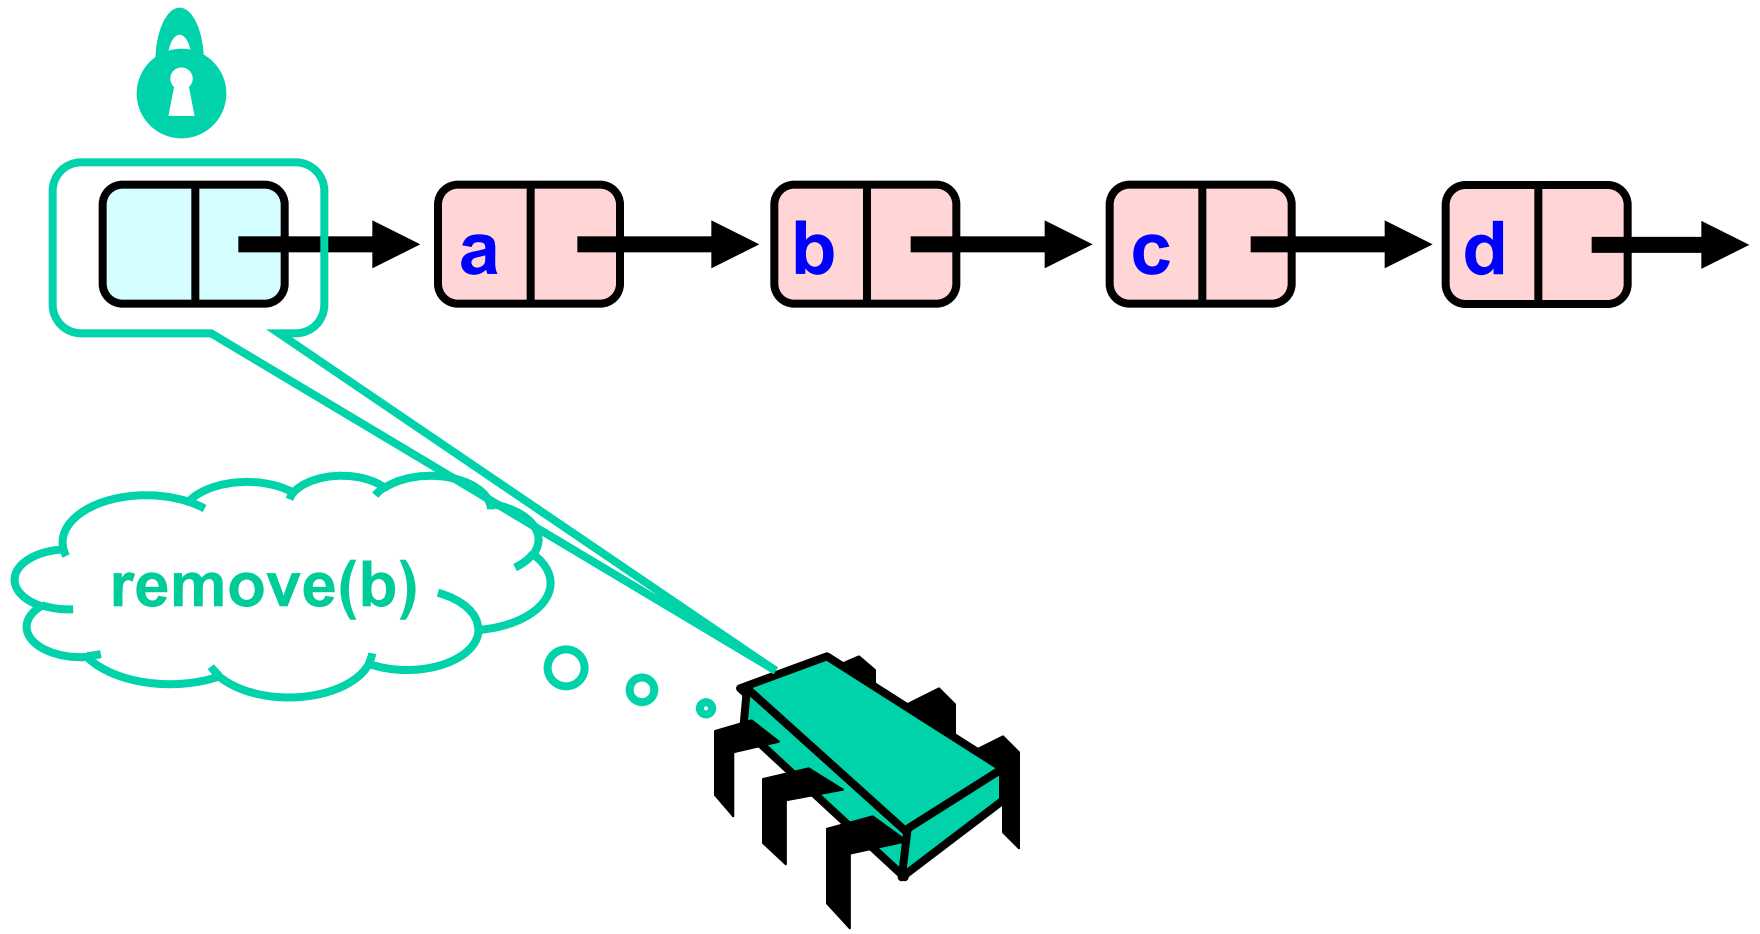




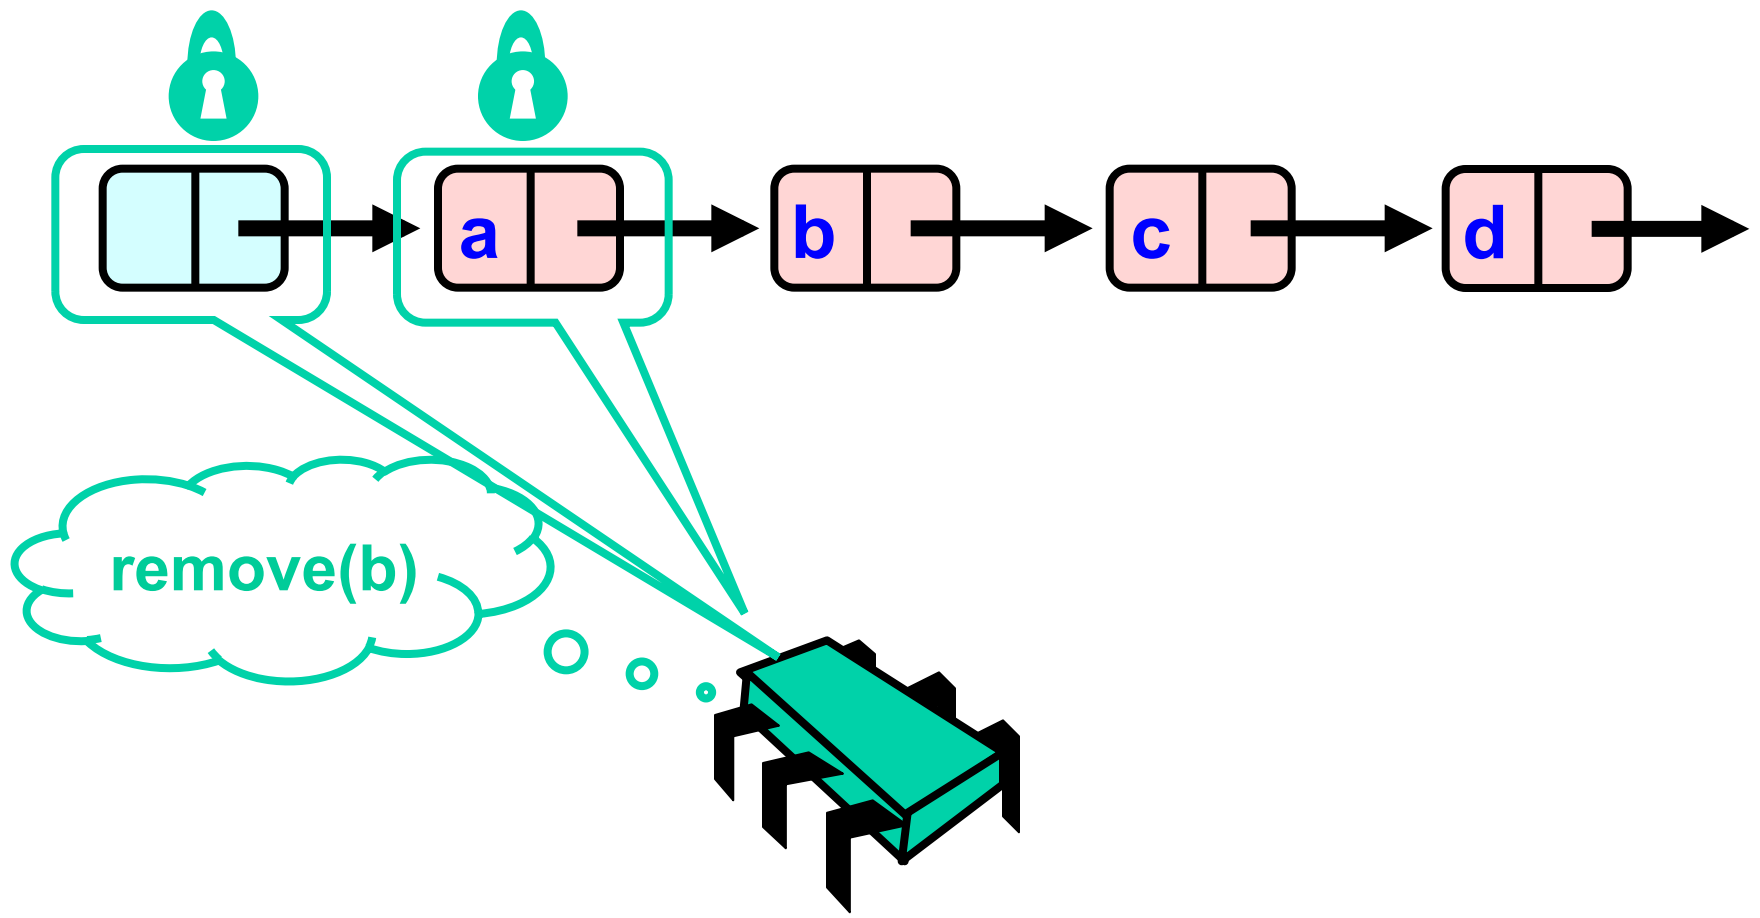
# Removing a Node



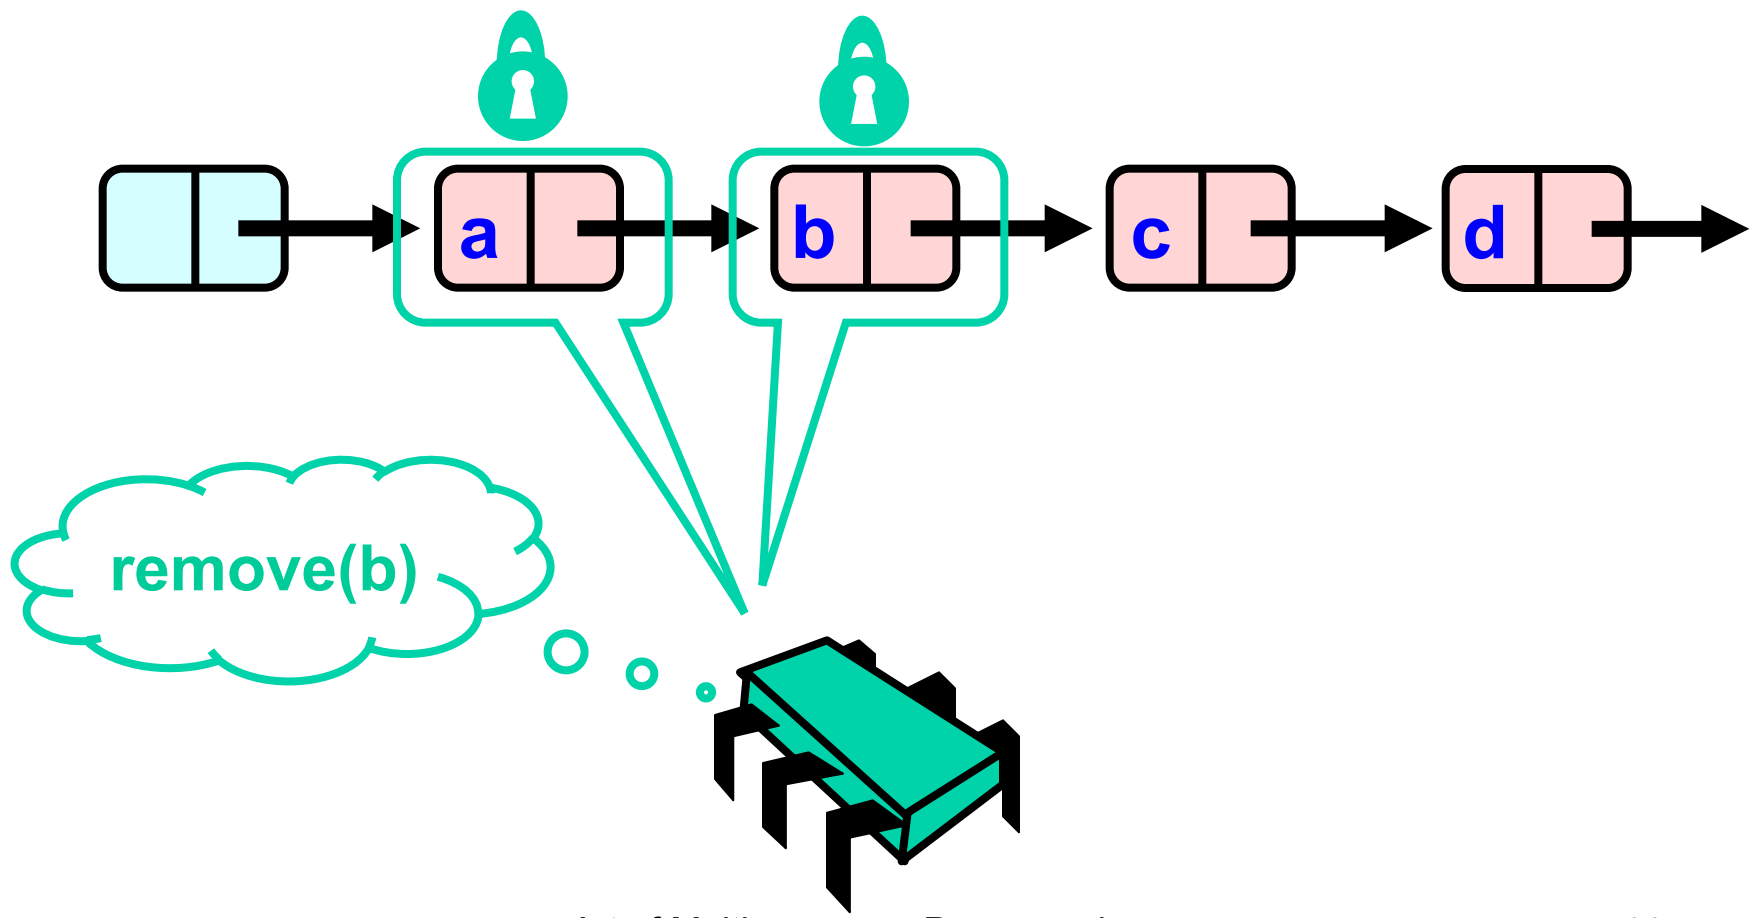
# Removing a Node




# Removing a Node



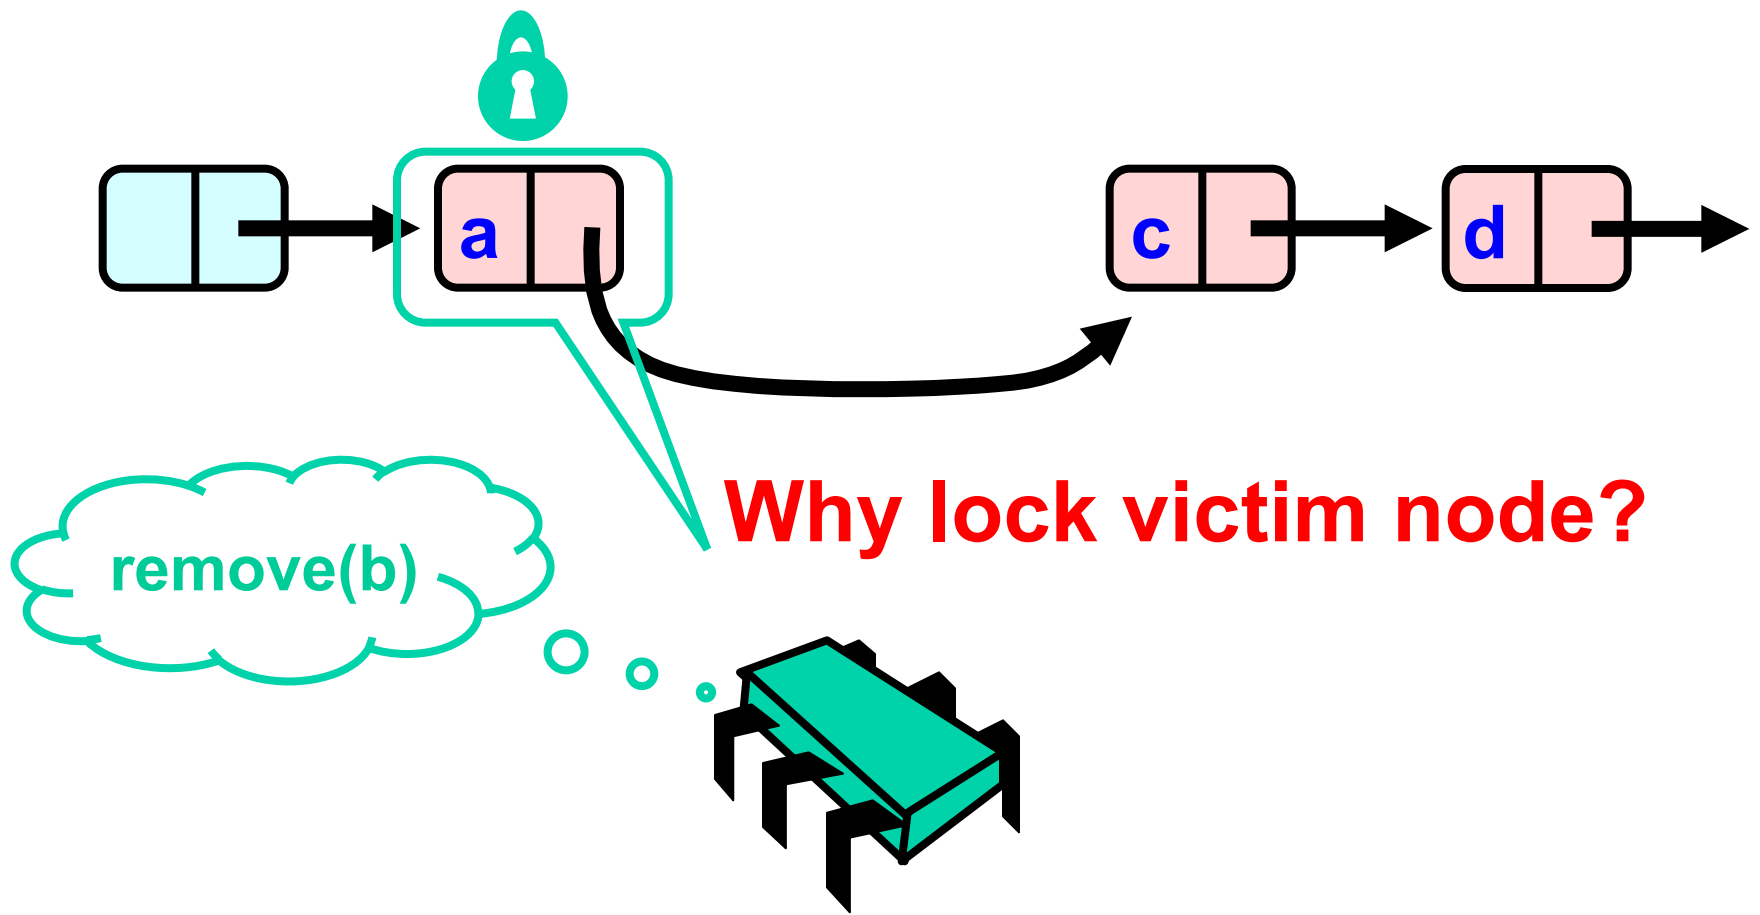
# Removing a Node



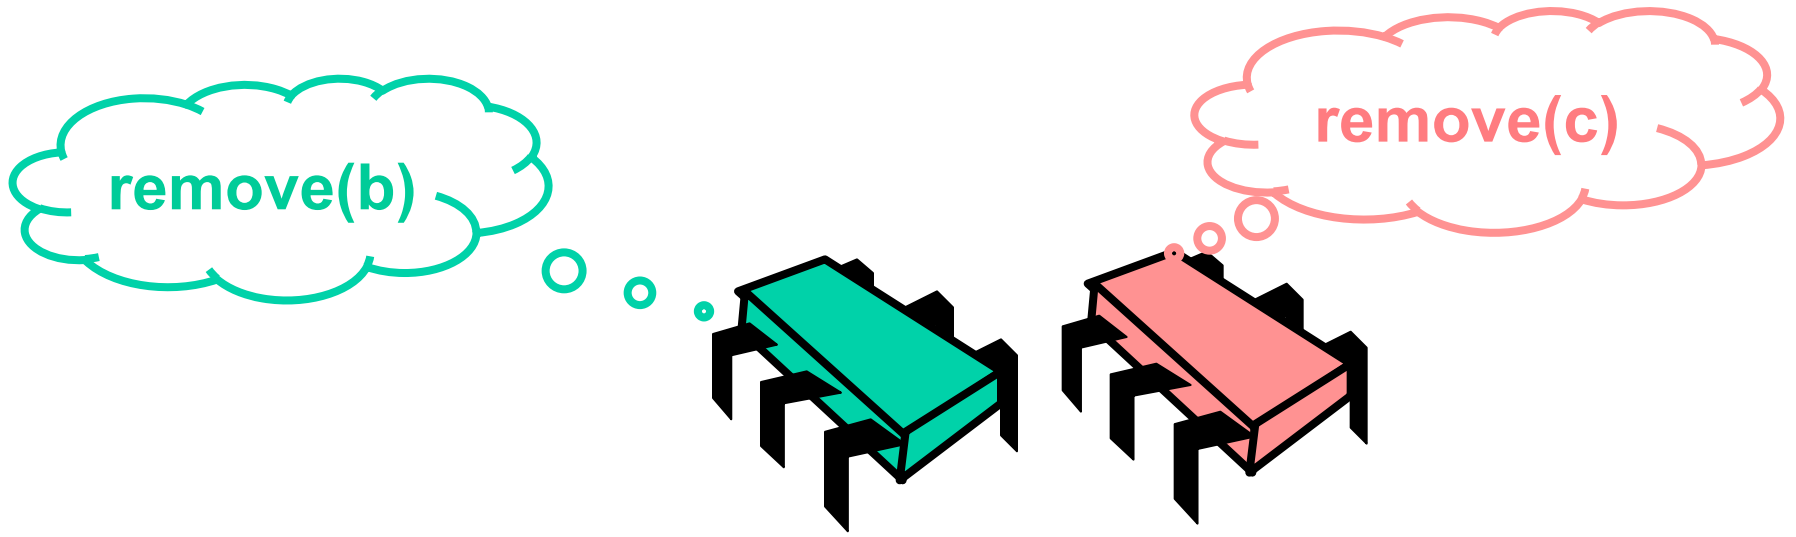
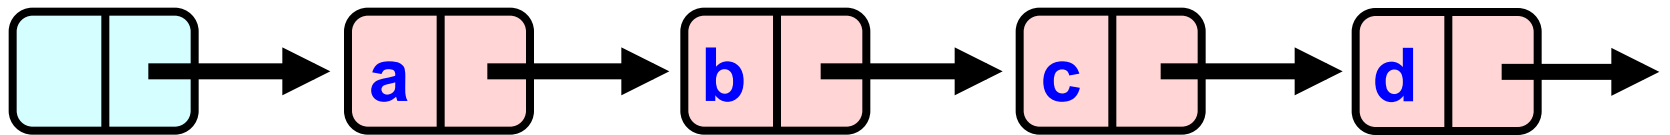
# Removing a Node



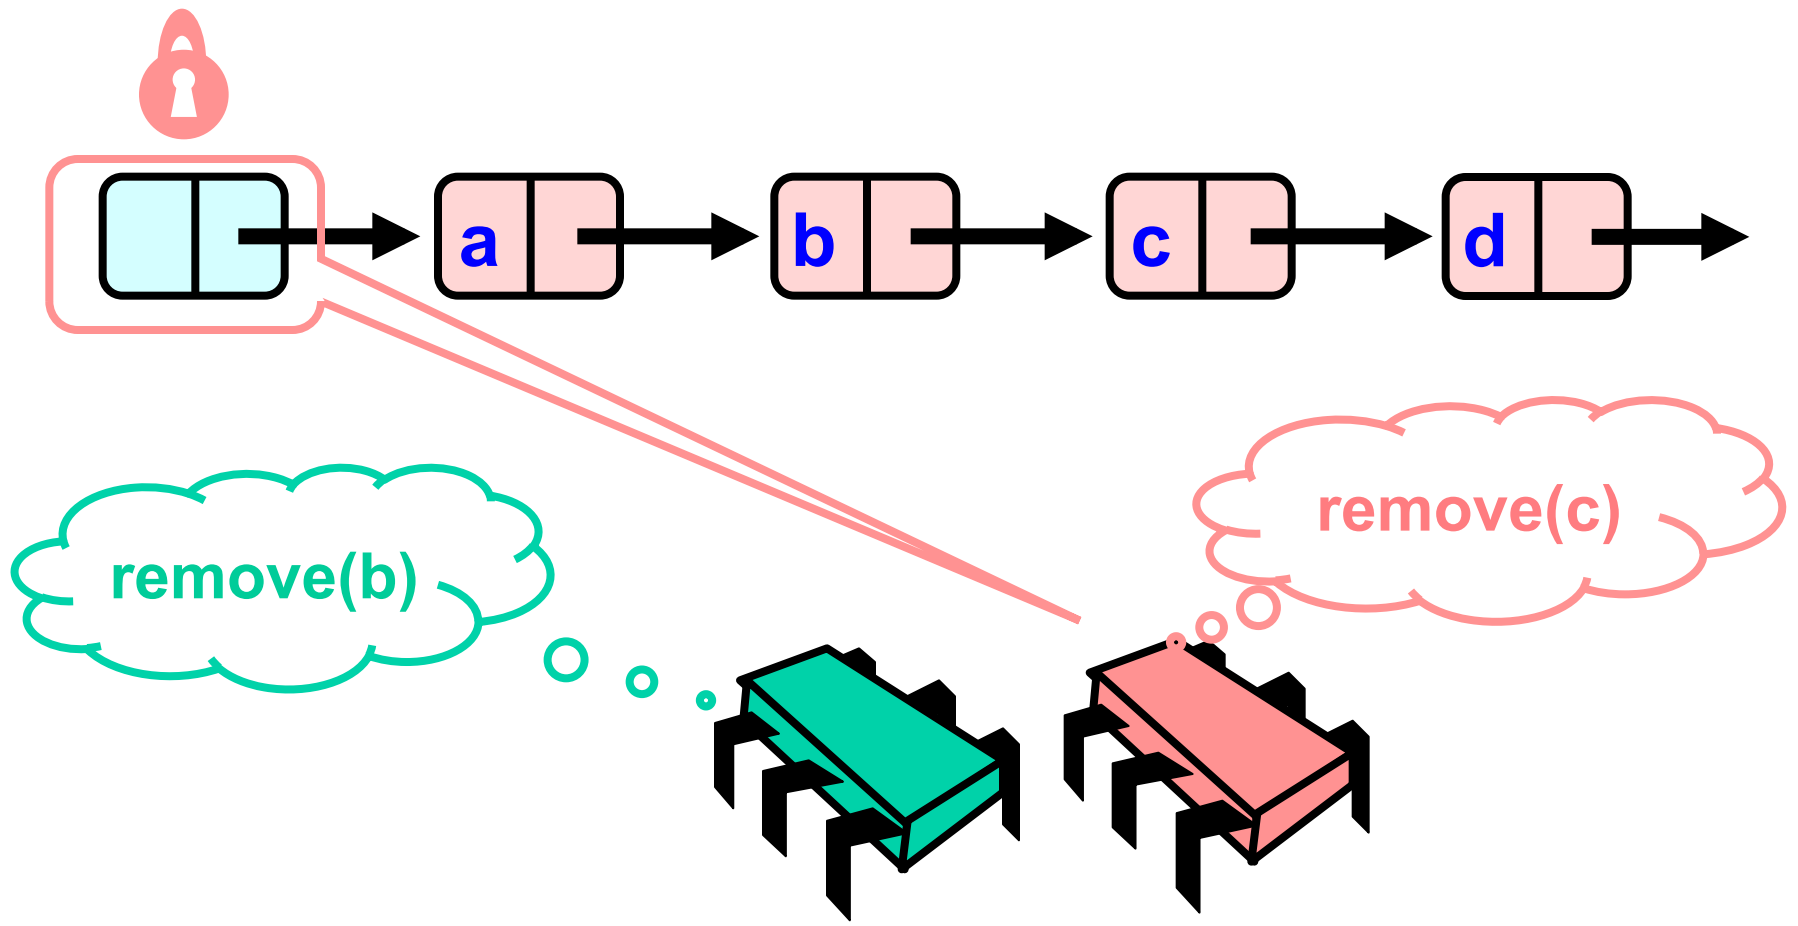
# Removing a Node



# Concurrent Removes

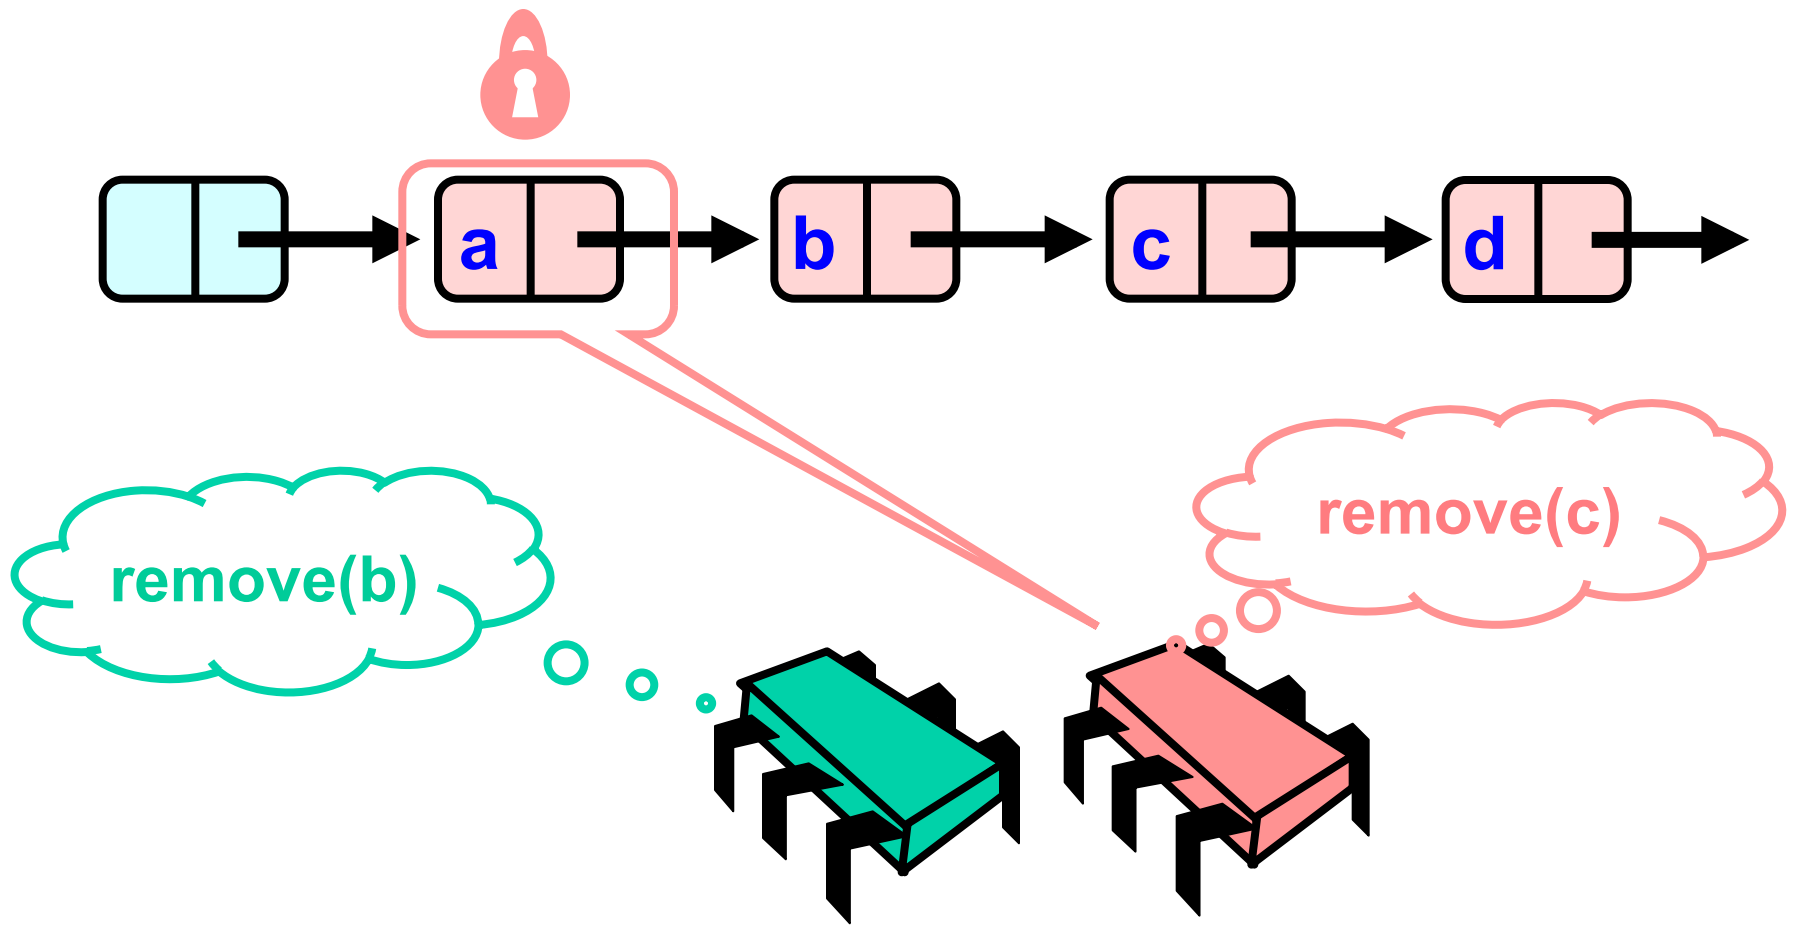


# Concurrent Removes

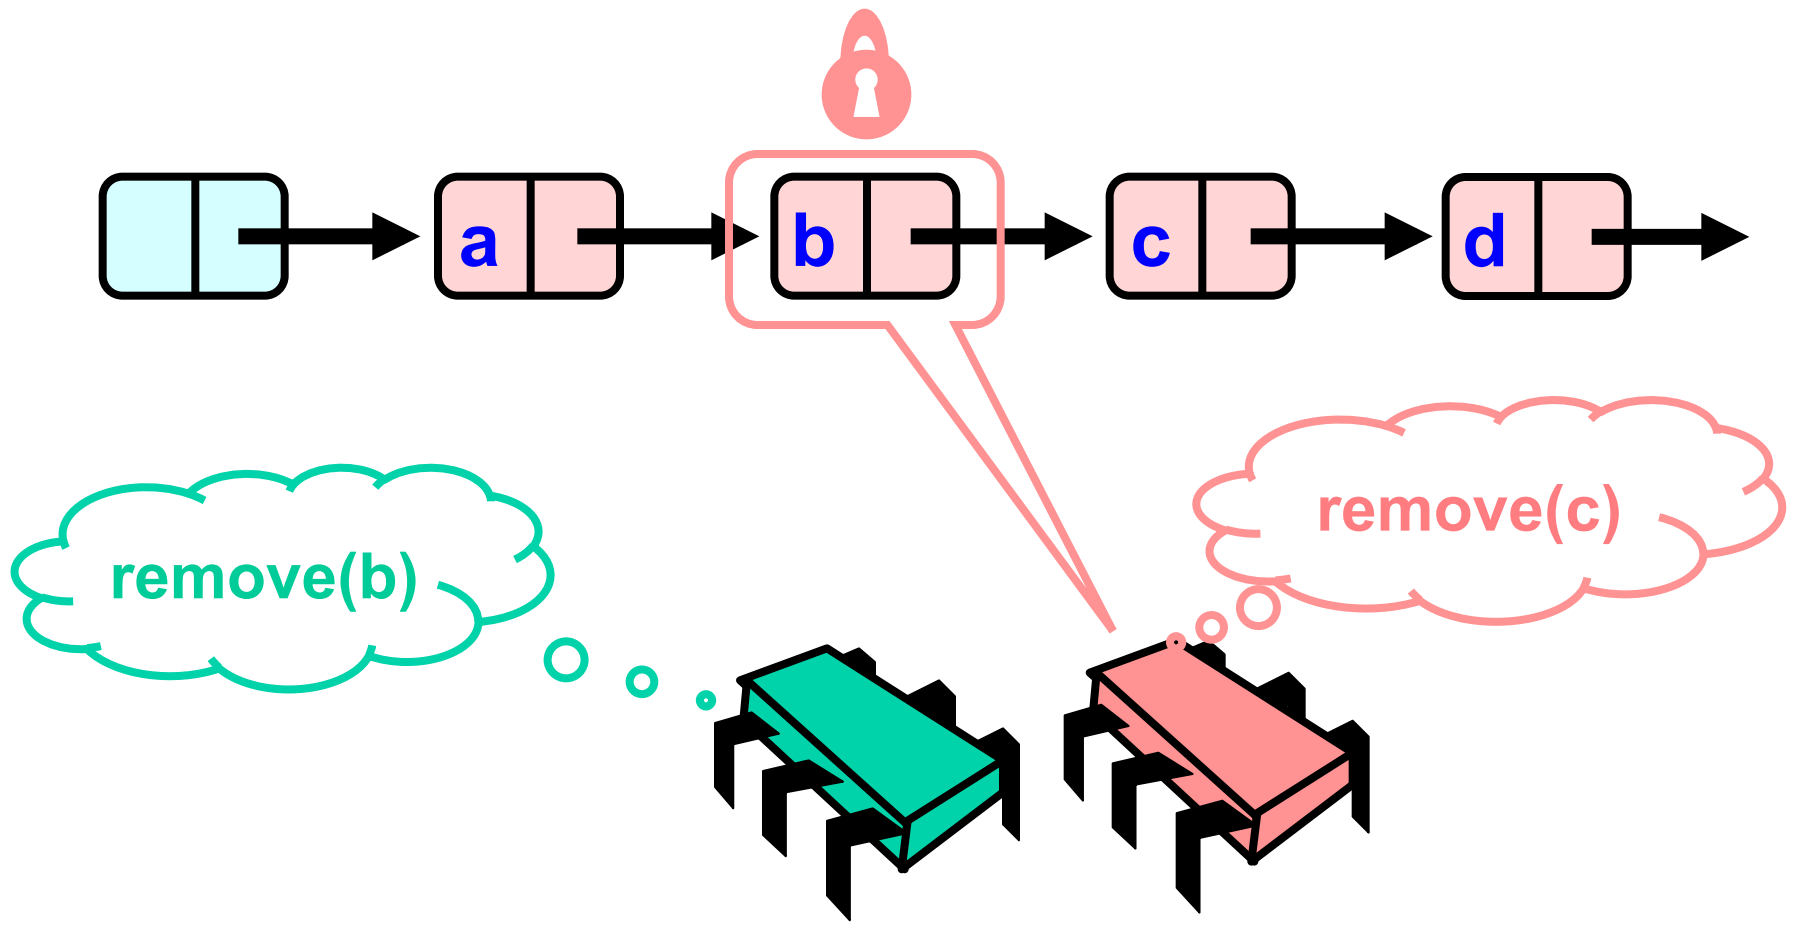




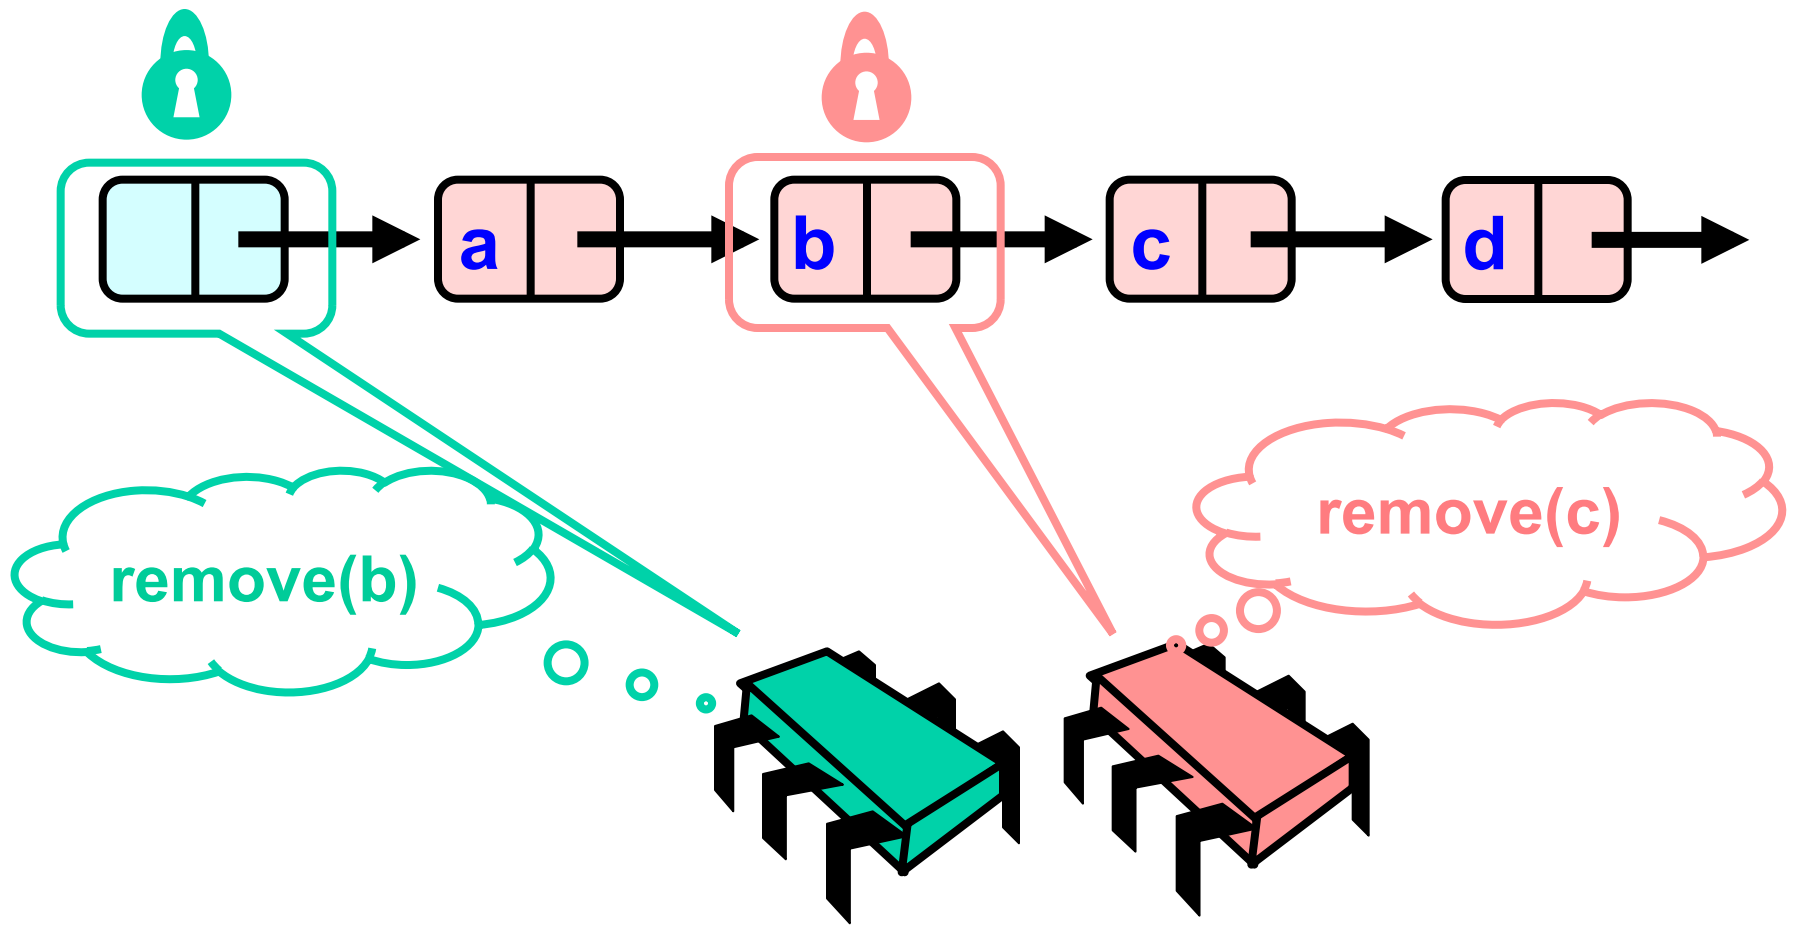
# Concurrent Removes



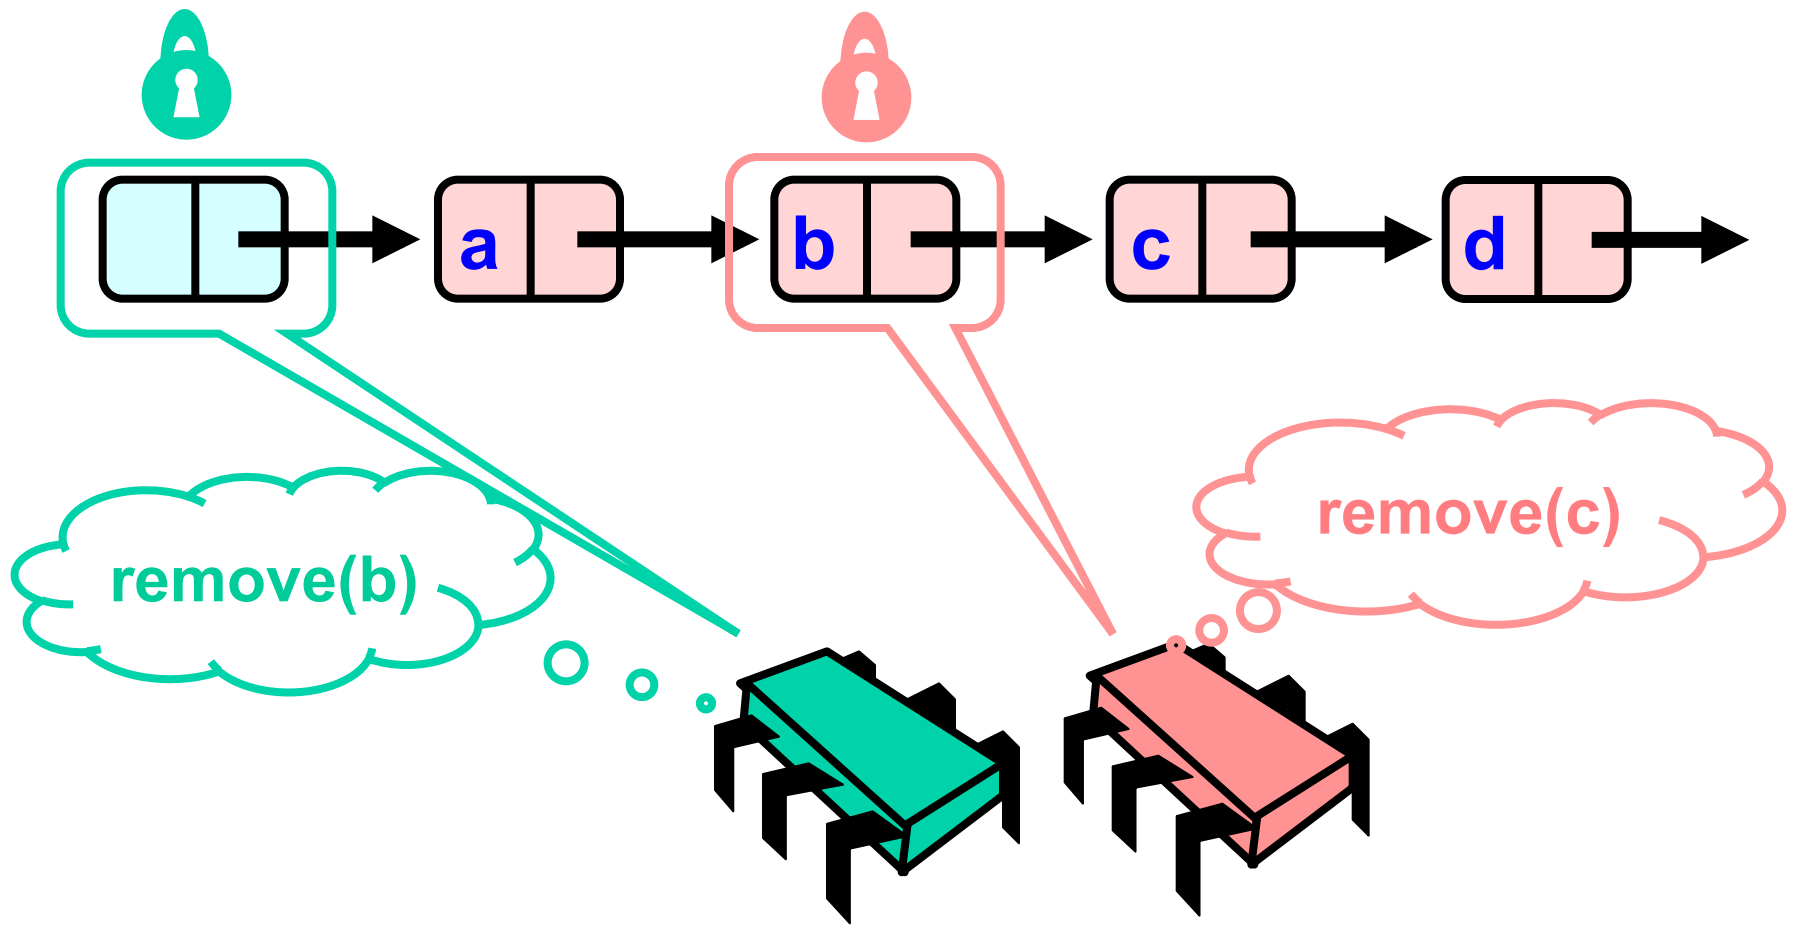
# Concurrent Removes



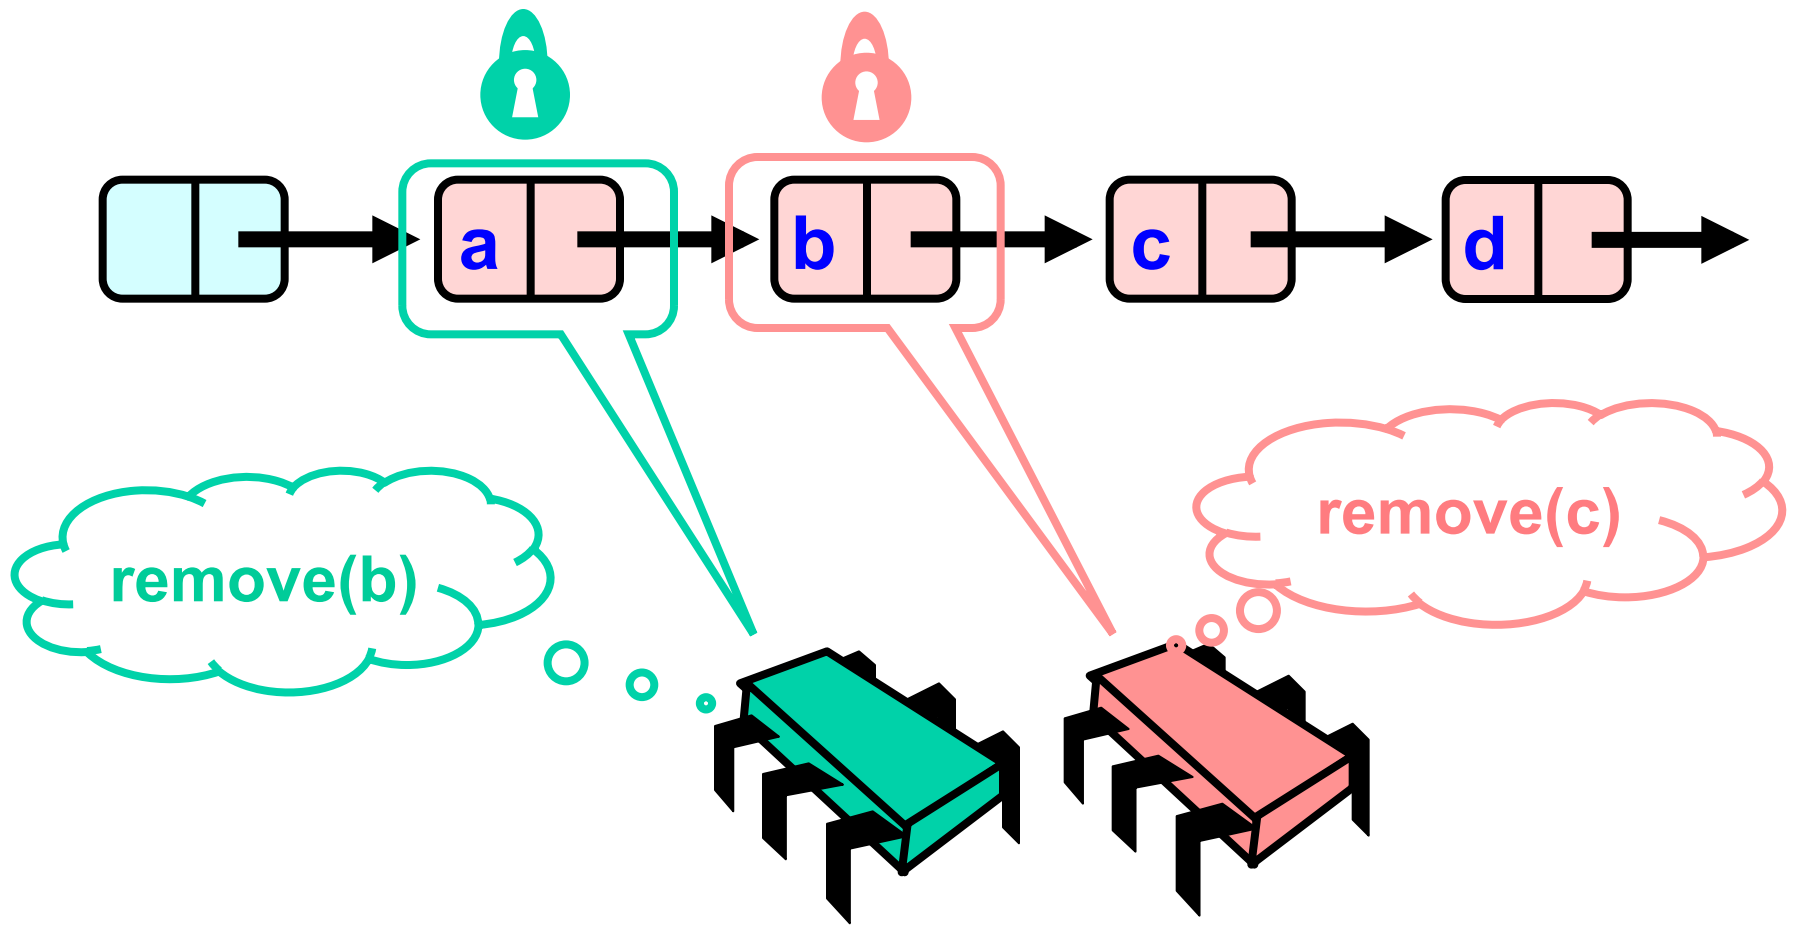
# Concurrent Removes



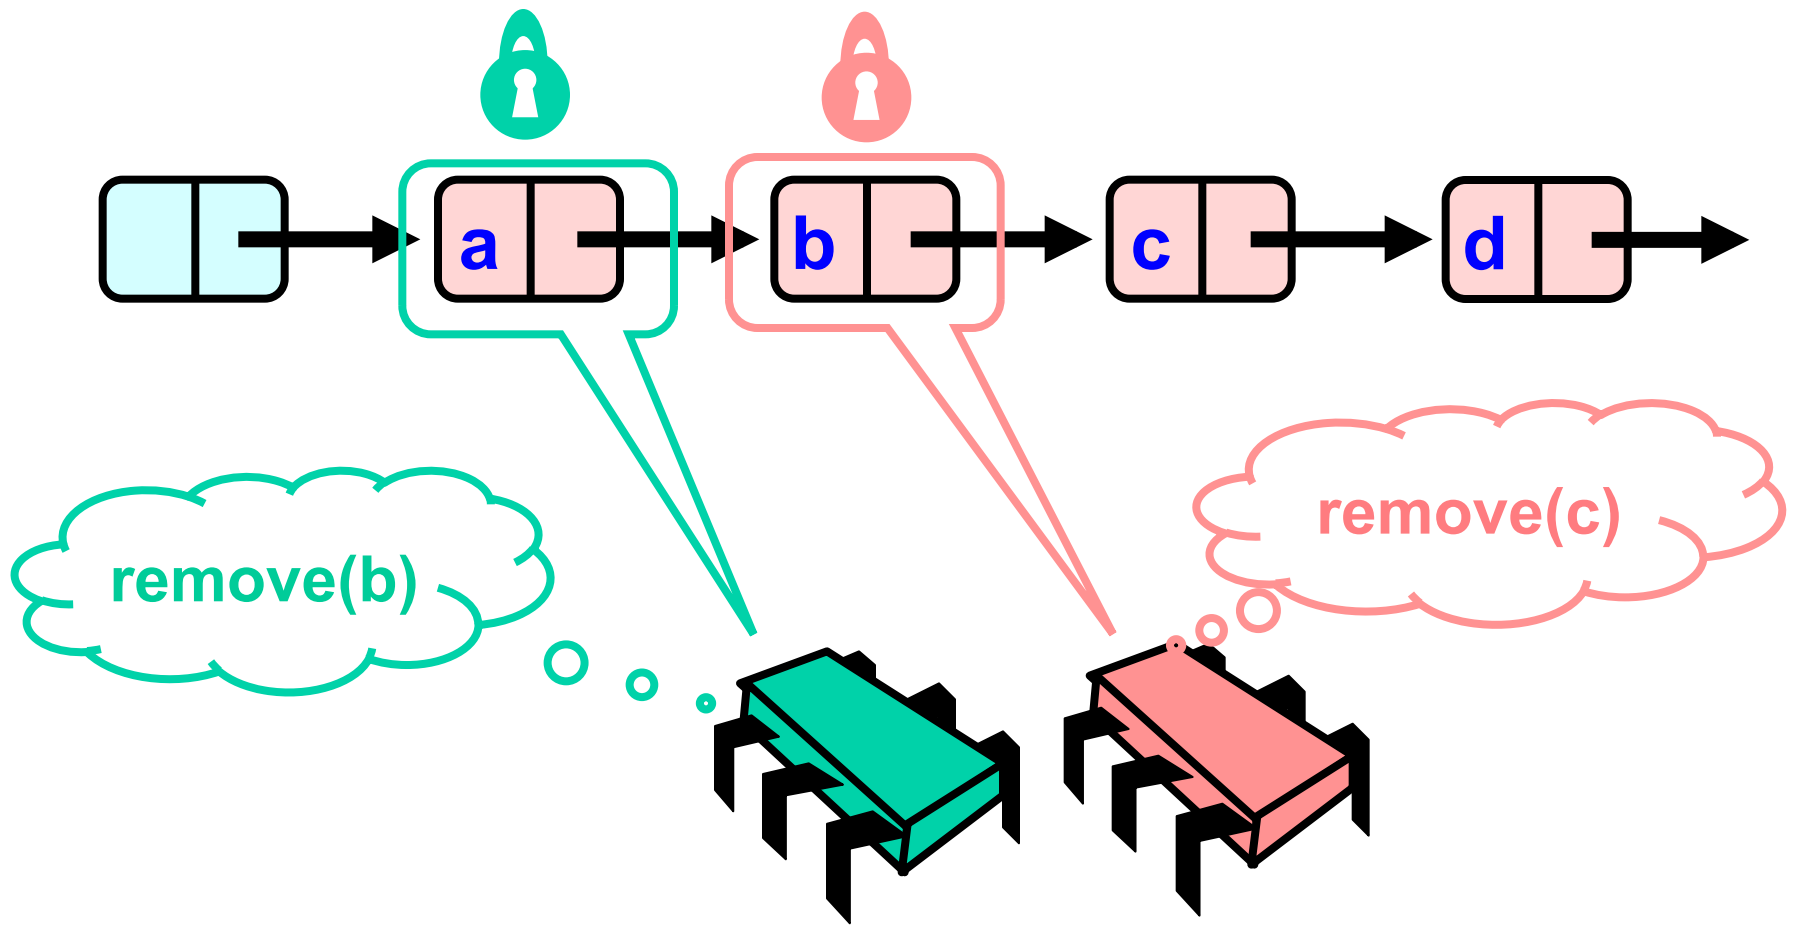
# Concurrent Removes



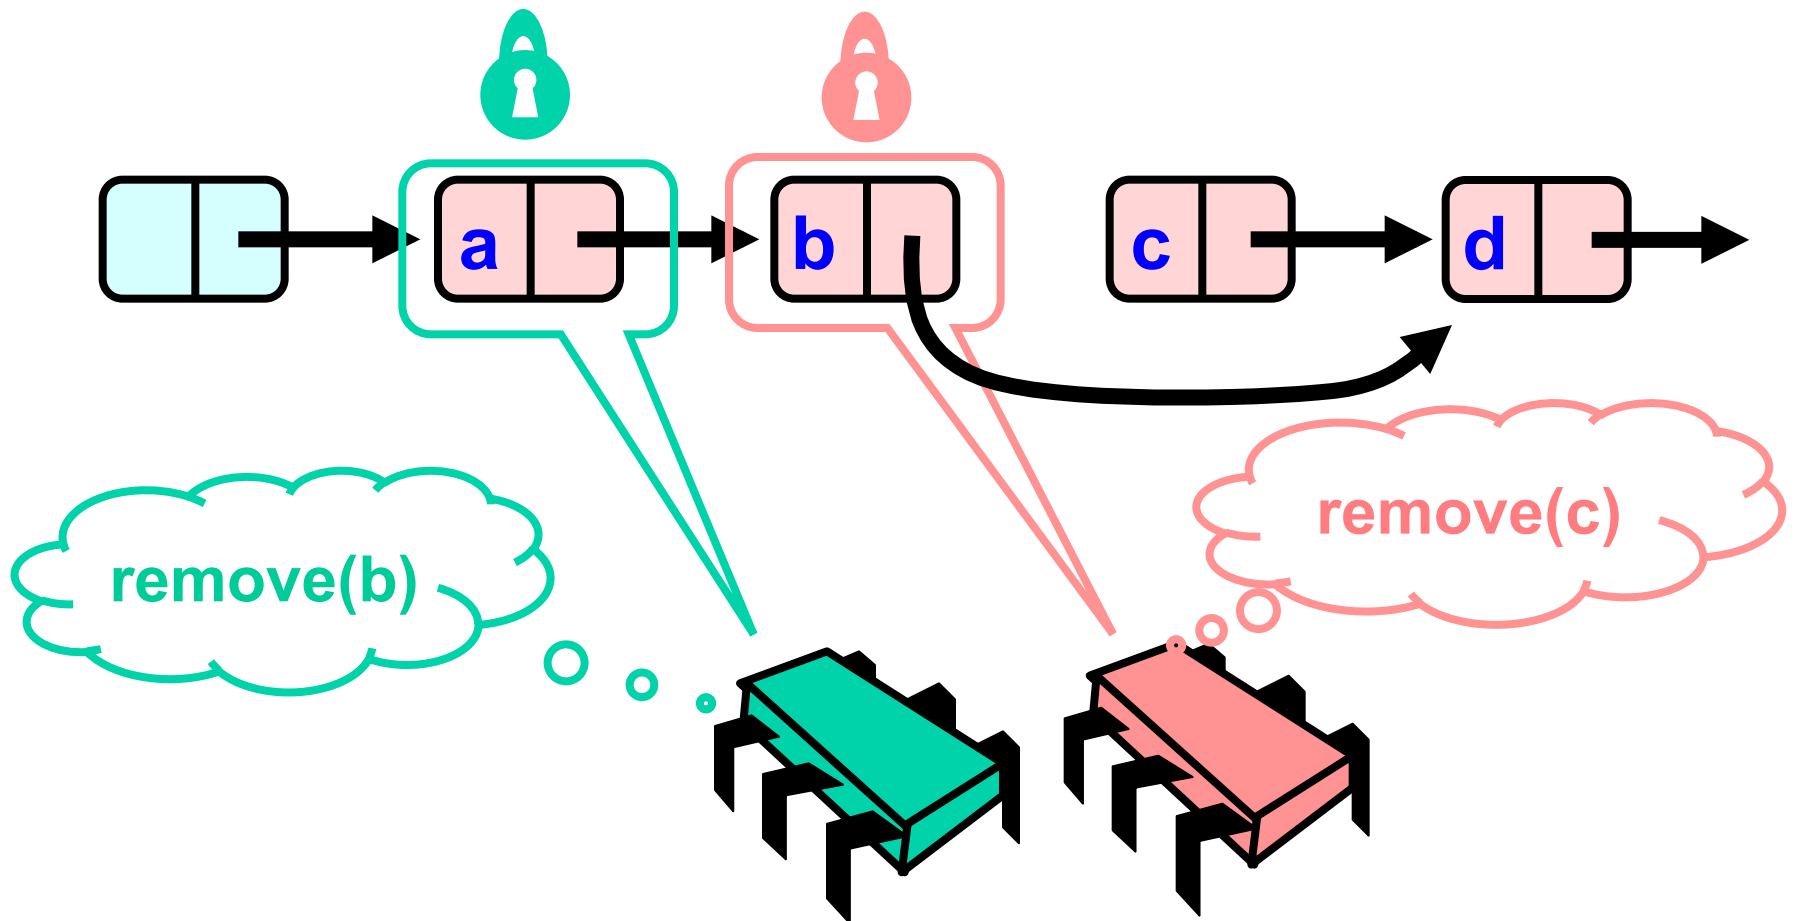
# Concurrent Removes



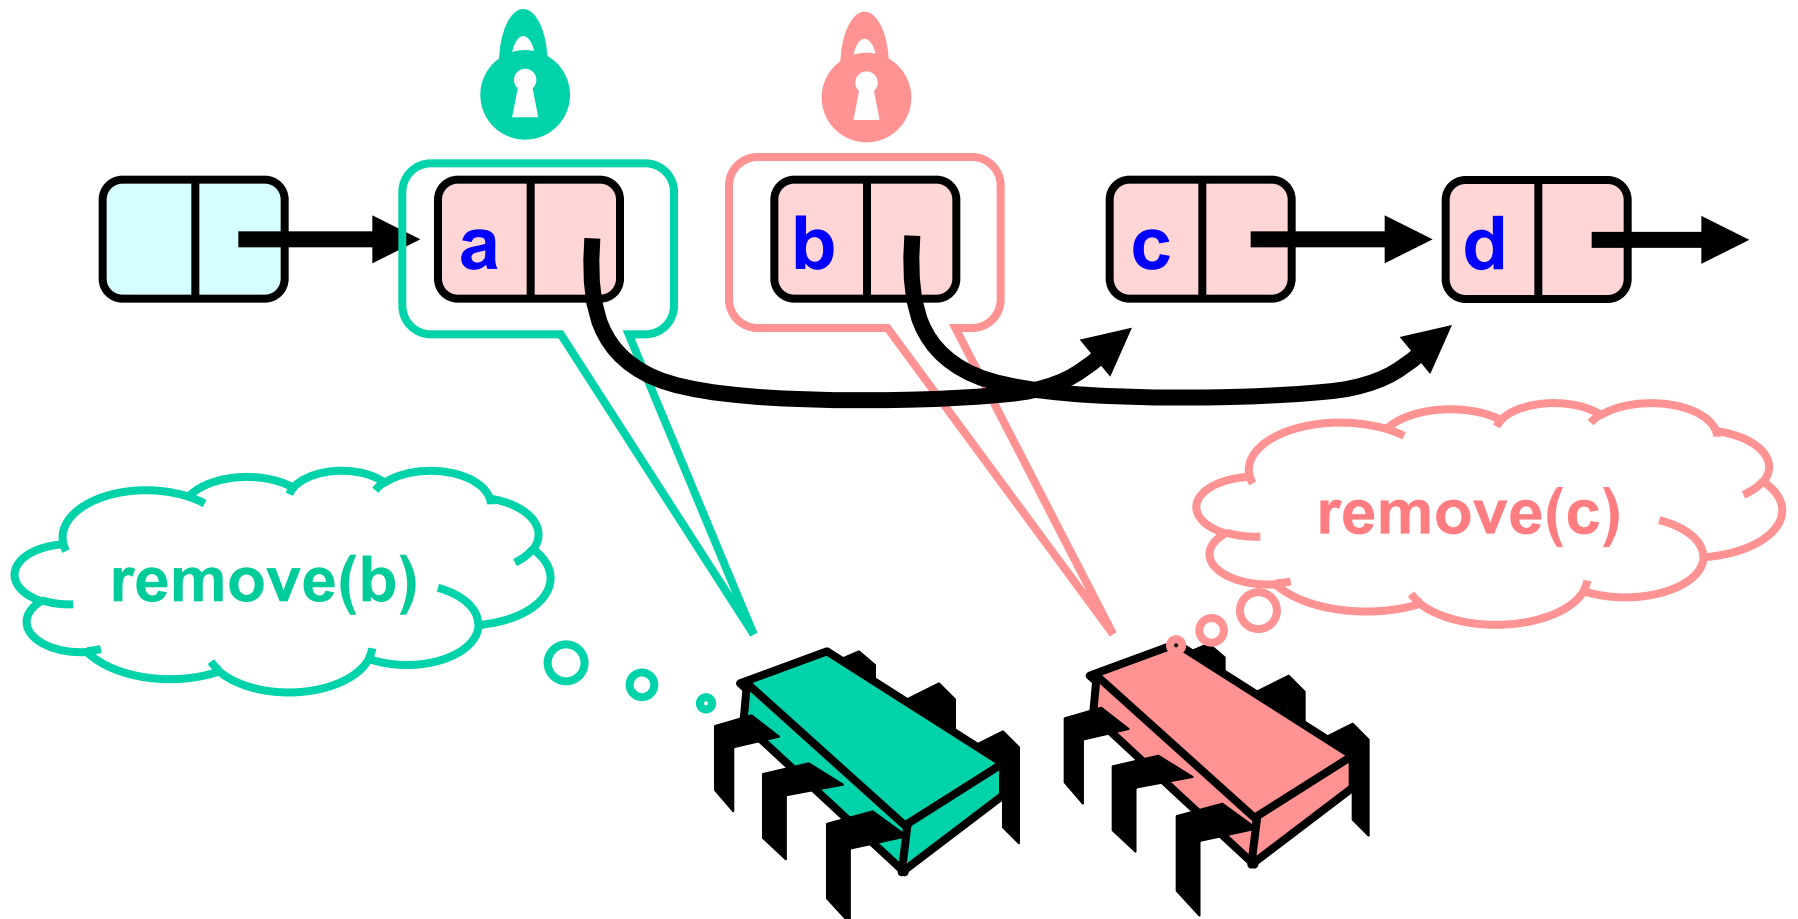
# Concurrent Removes



# Concurrent Removes




# Concurrent Removes






Uh, Oh



Uh, Oh

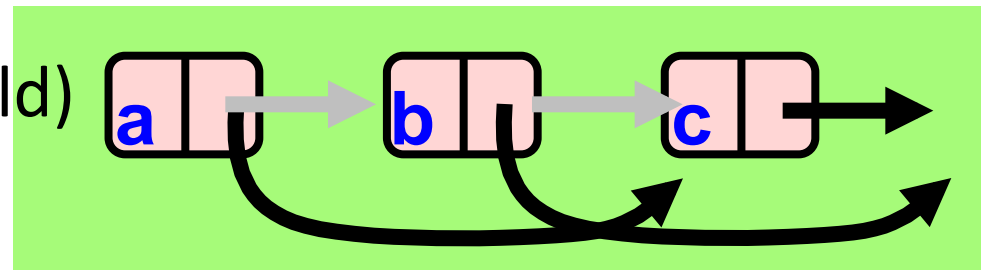
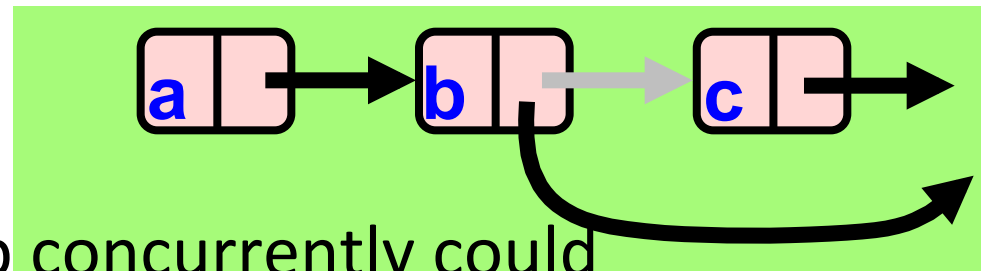
Bad news, **c** not removed



# Problem

- To delete node c
  - Swing node b's next field to d (c's next field)

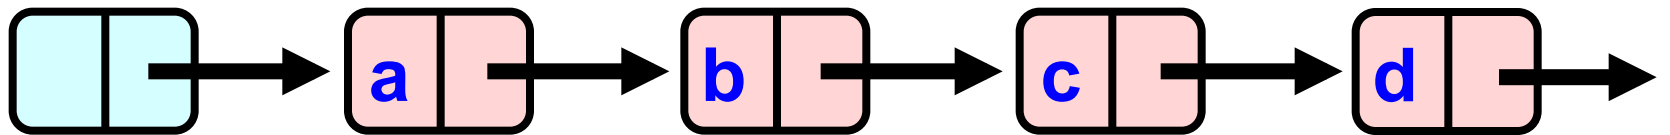
- Problem is,
  - Someone deleting b concurrently could direct a pointer to c (reading b's next field)



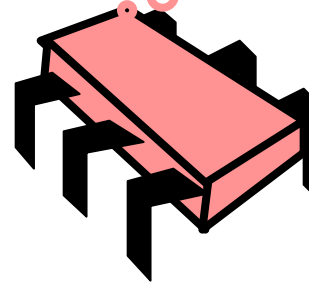
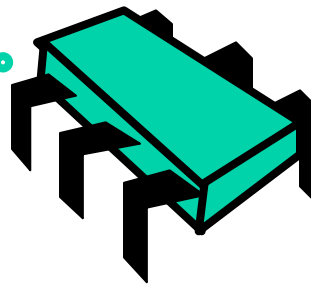
# Insight

- If a node is locked
  - No one can **change** node's *successor*
- If a thread locks
  - Node to be deleted (so its successor don't change)
  - And its predecessor (so you are the only one changing its successor)
  - Then it works

# Removing a Node

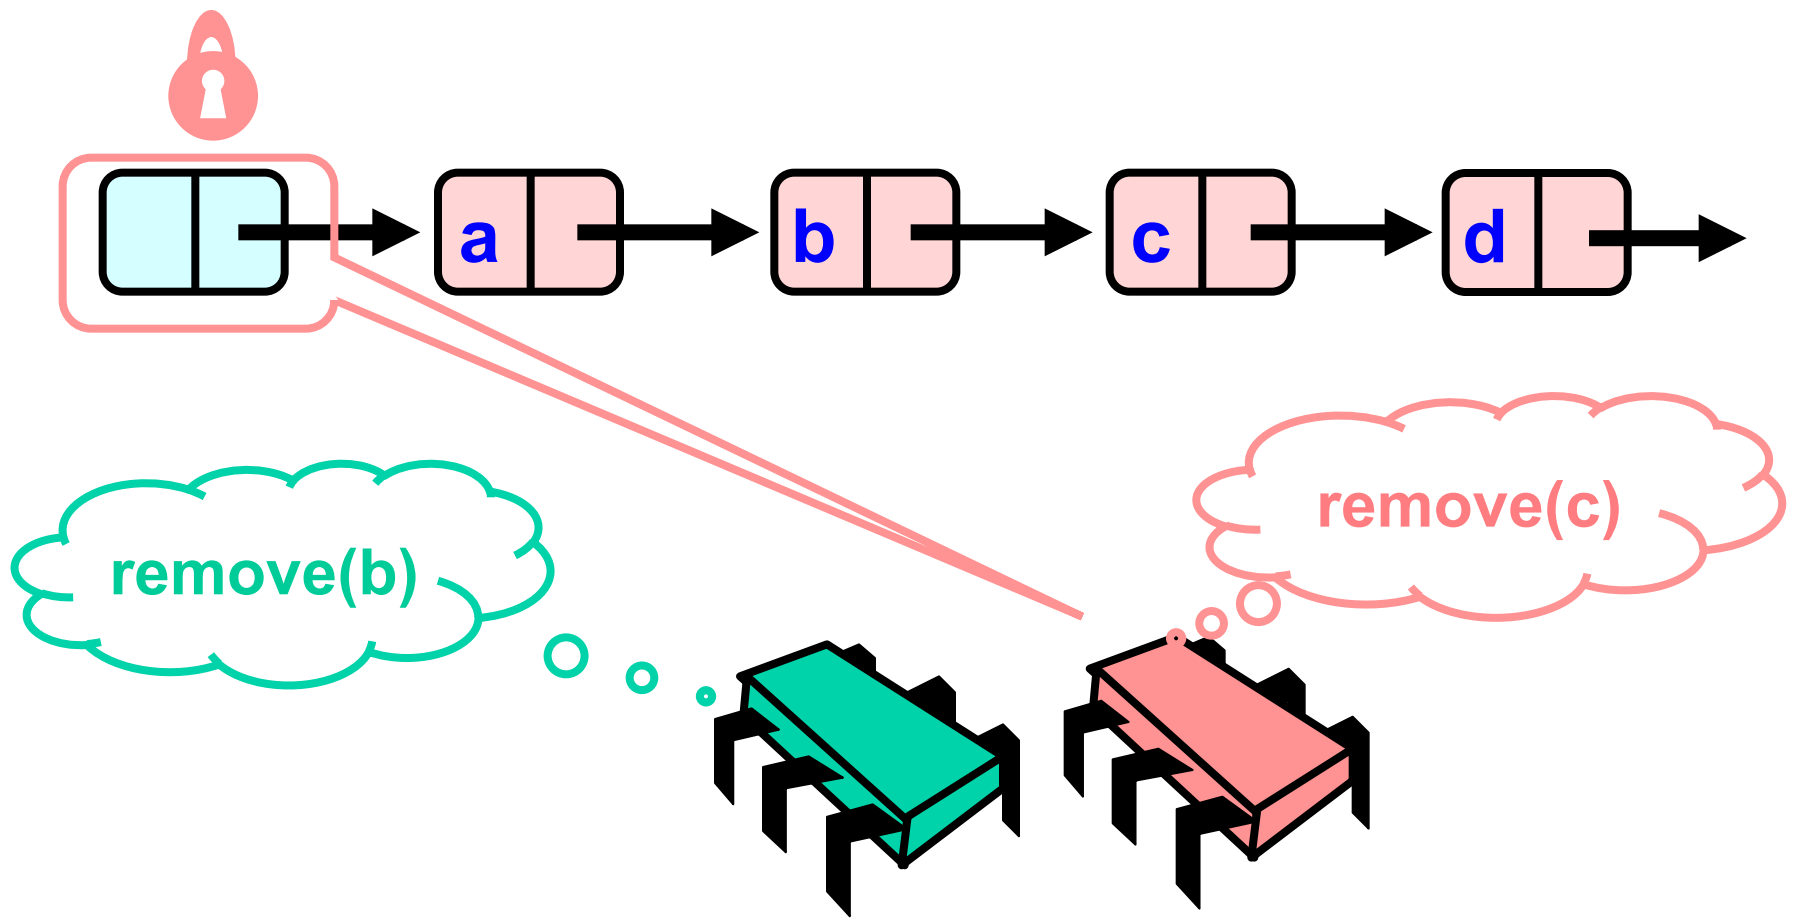


remove(b)

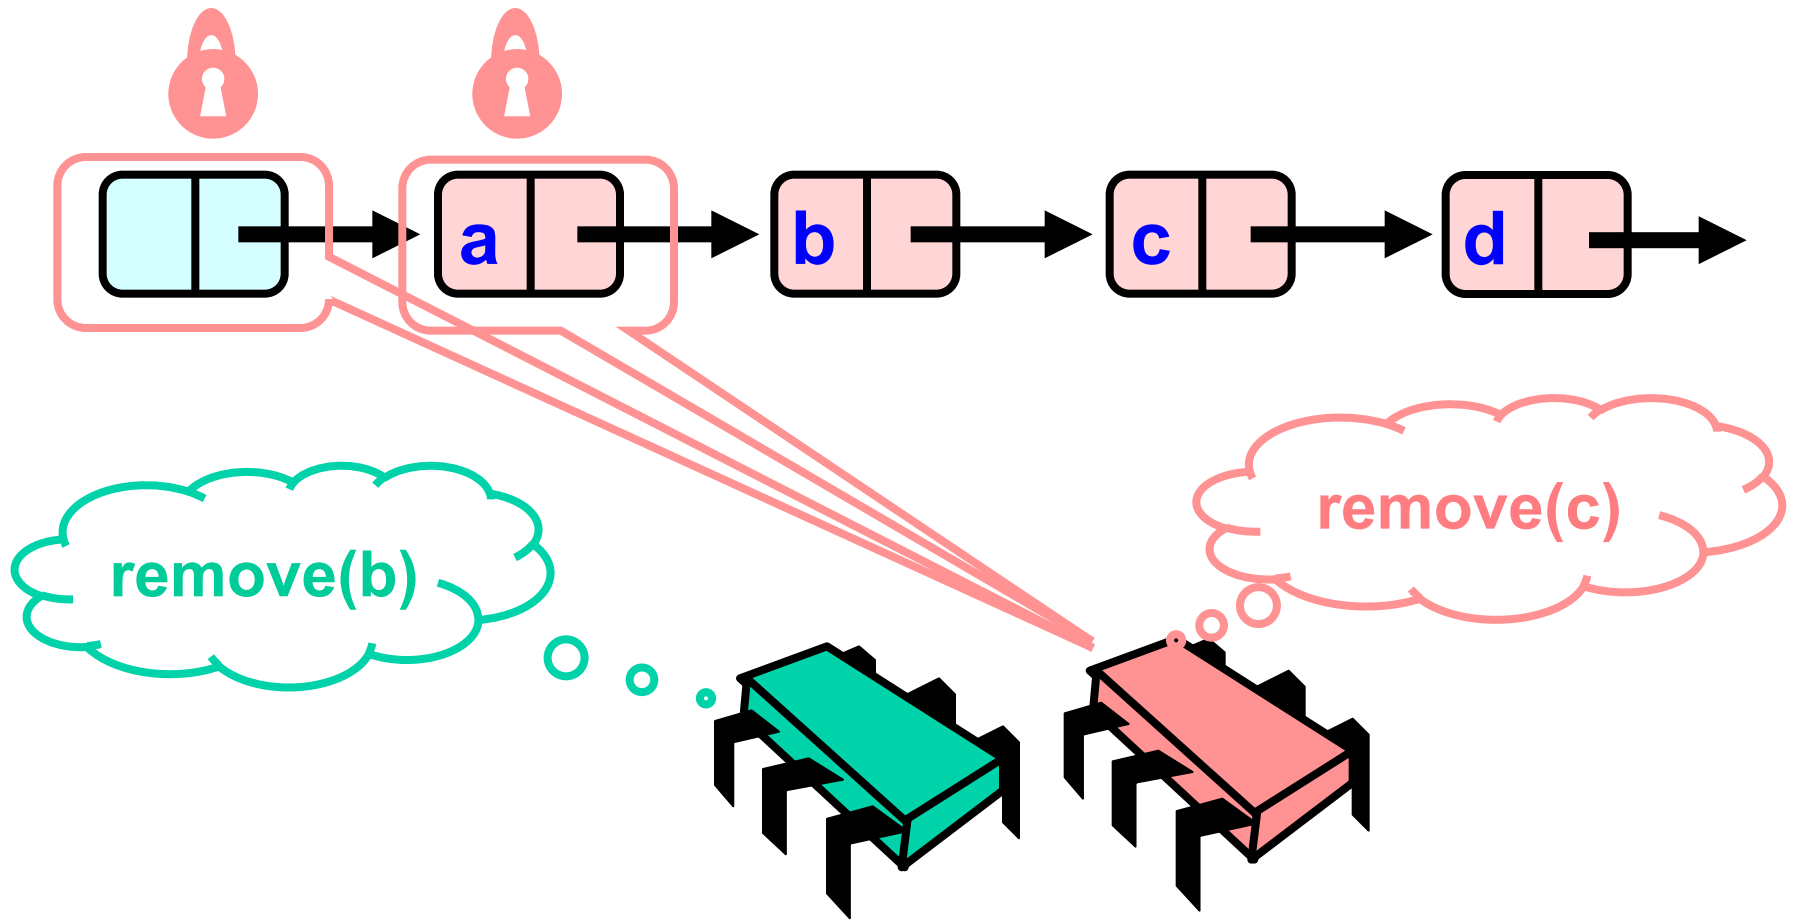


remove(c)

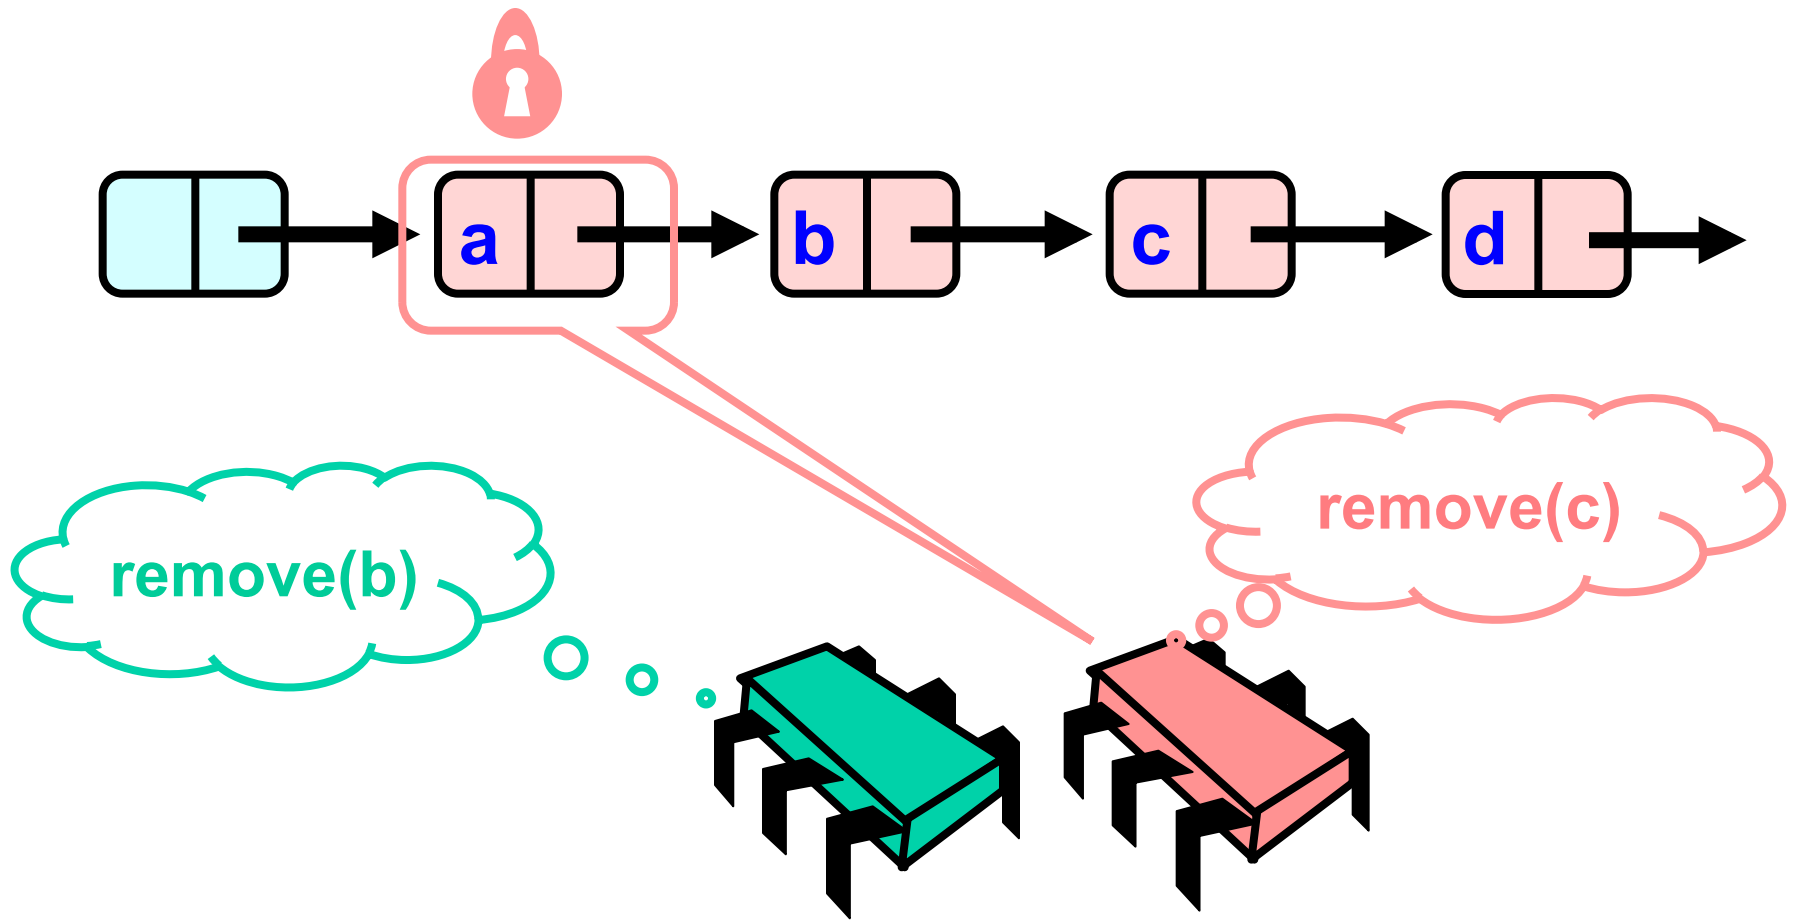
# Removing a Node



# Removing a Node

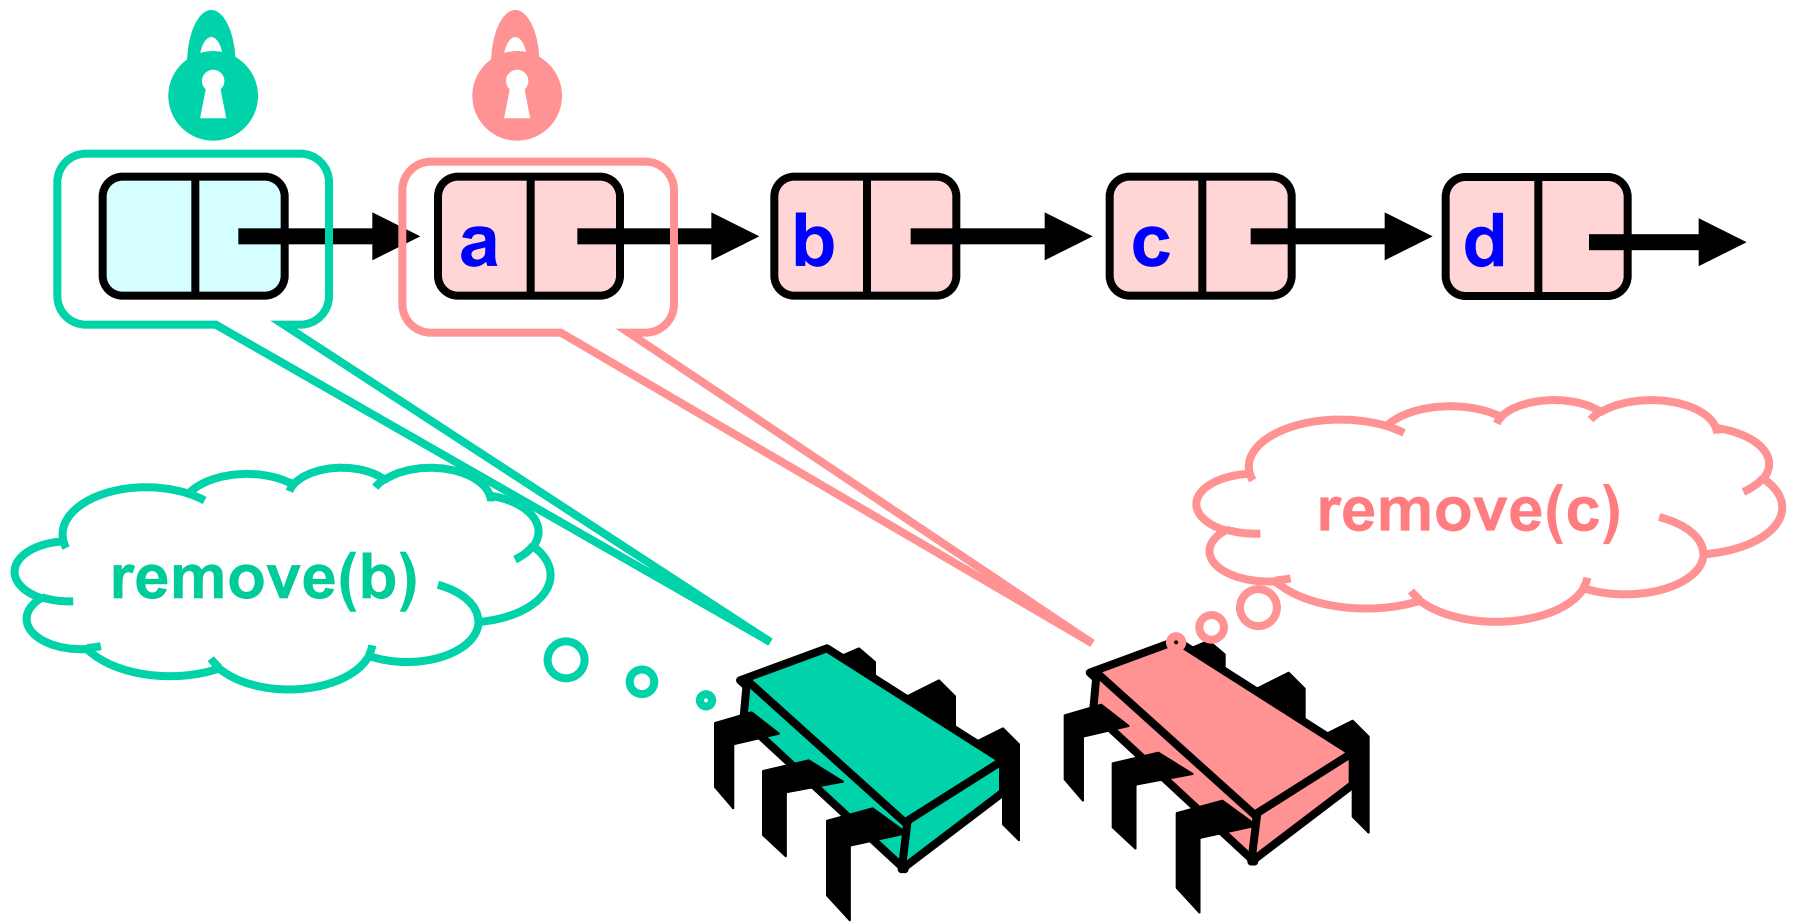


# Removing a Node

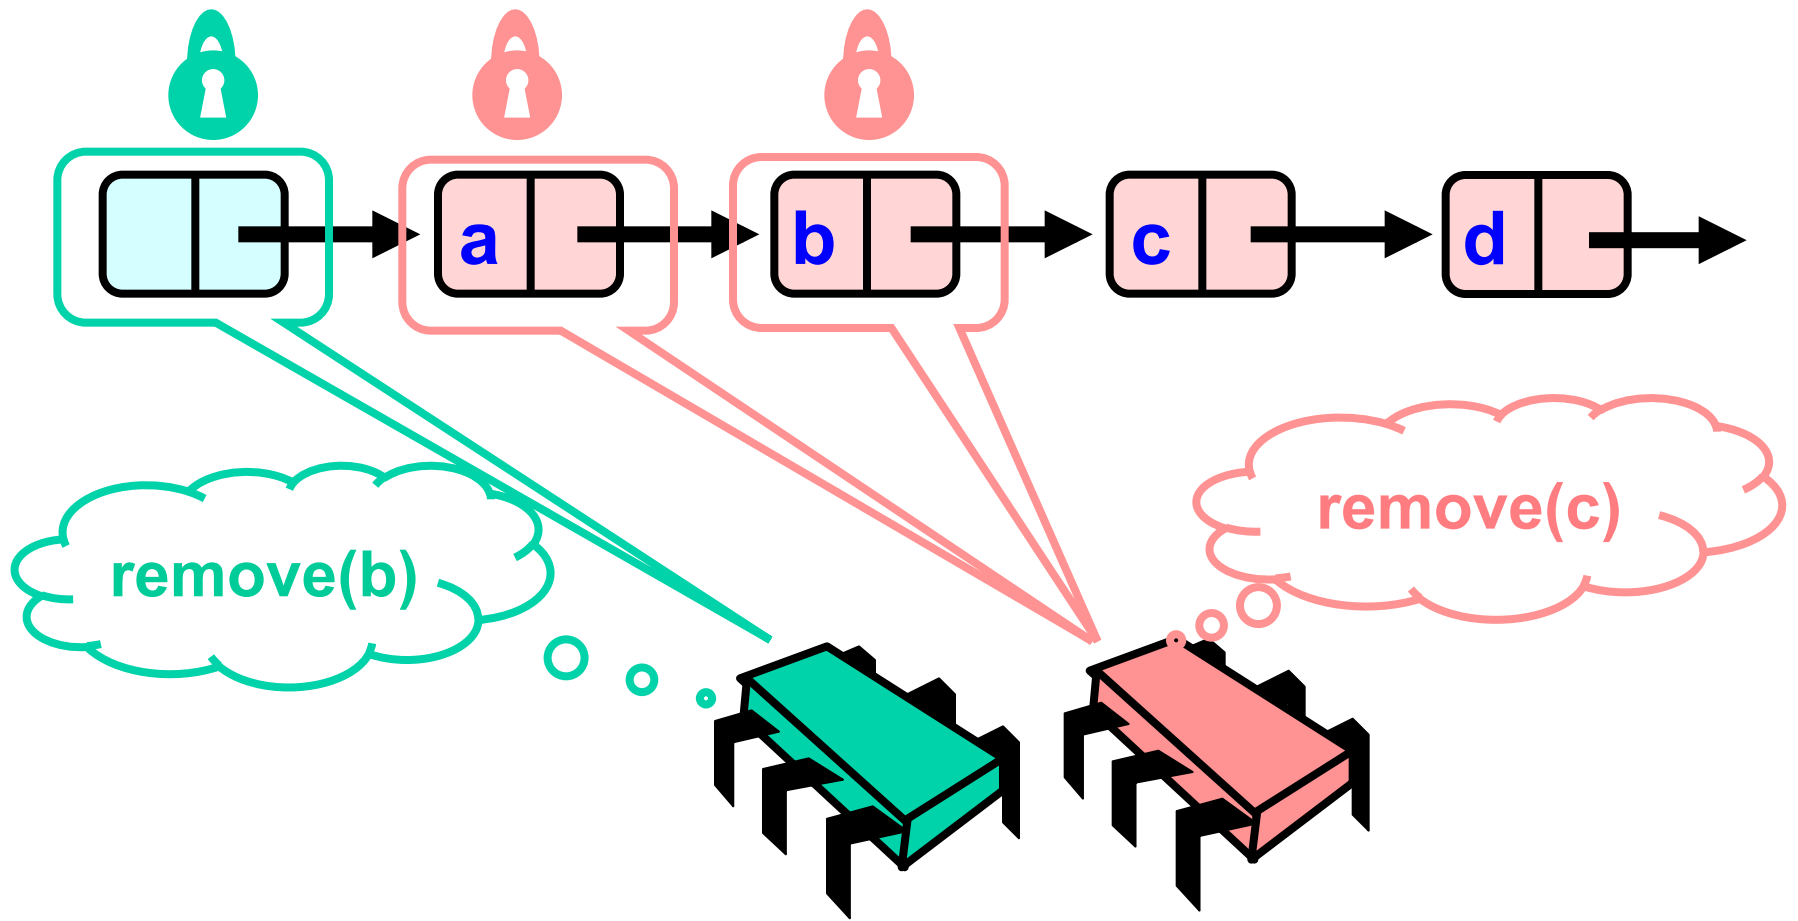




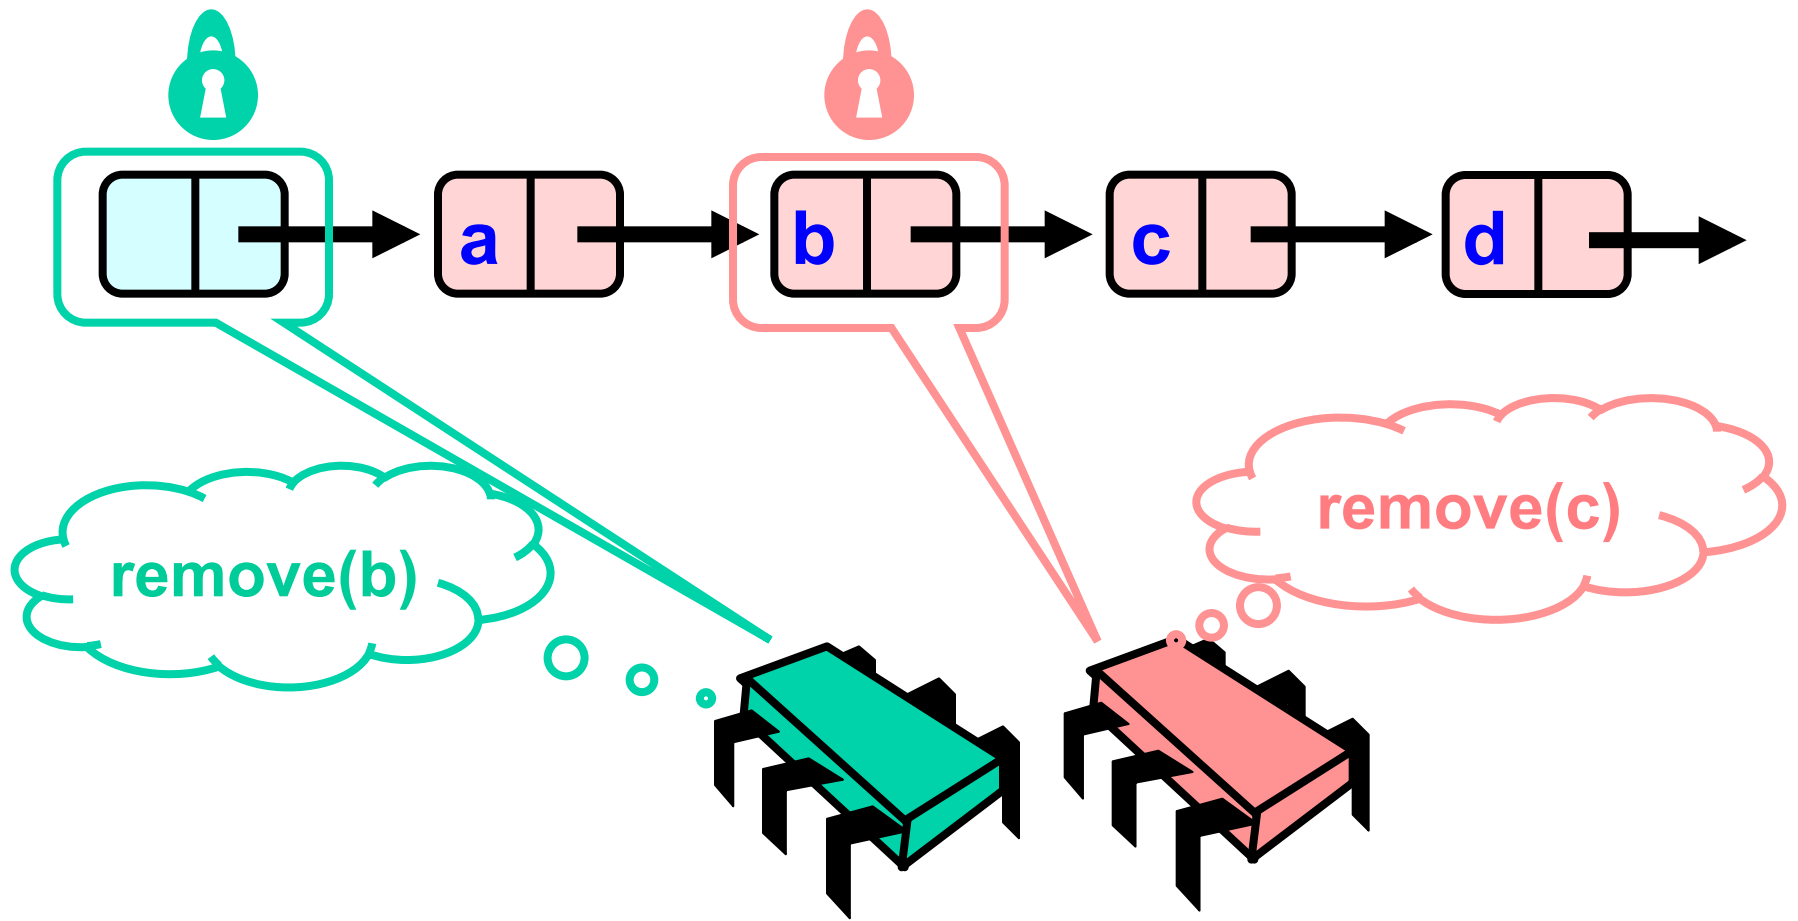
# Removing a Node



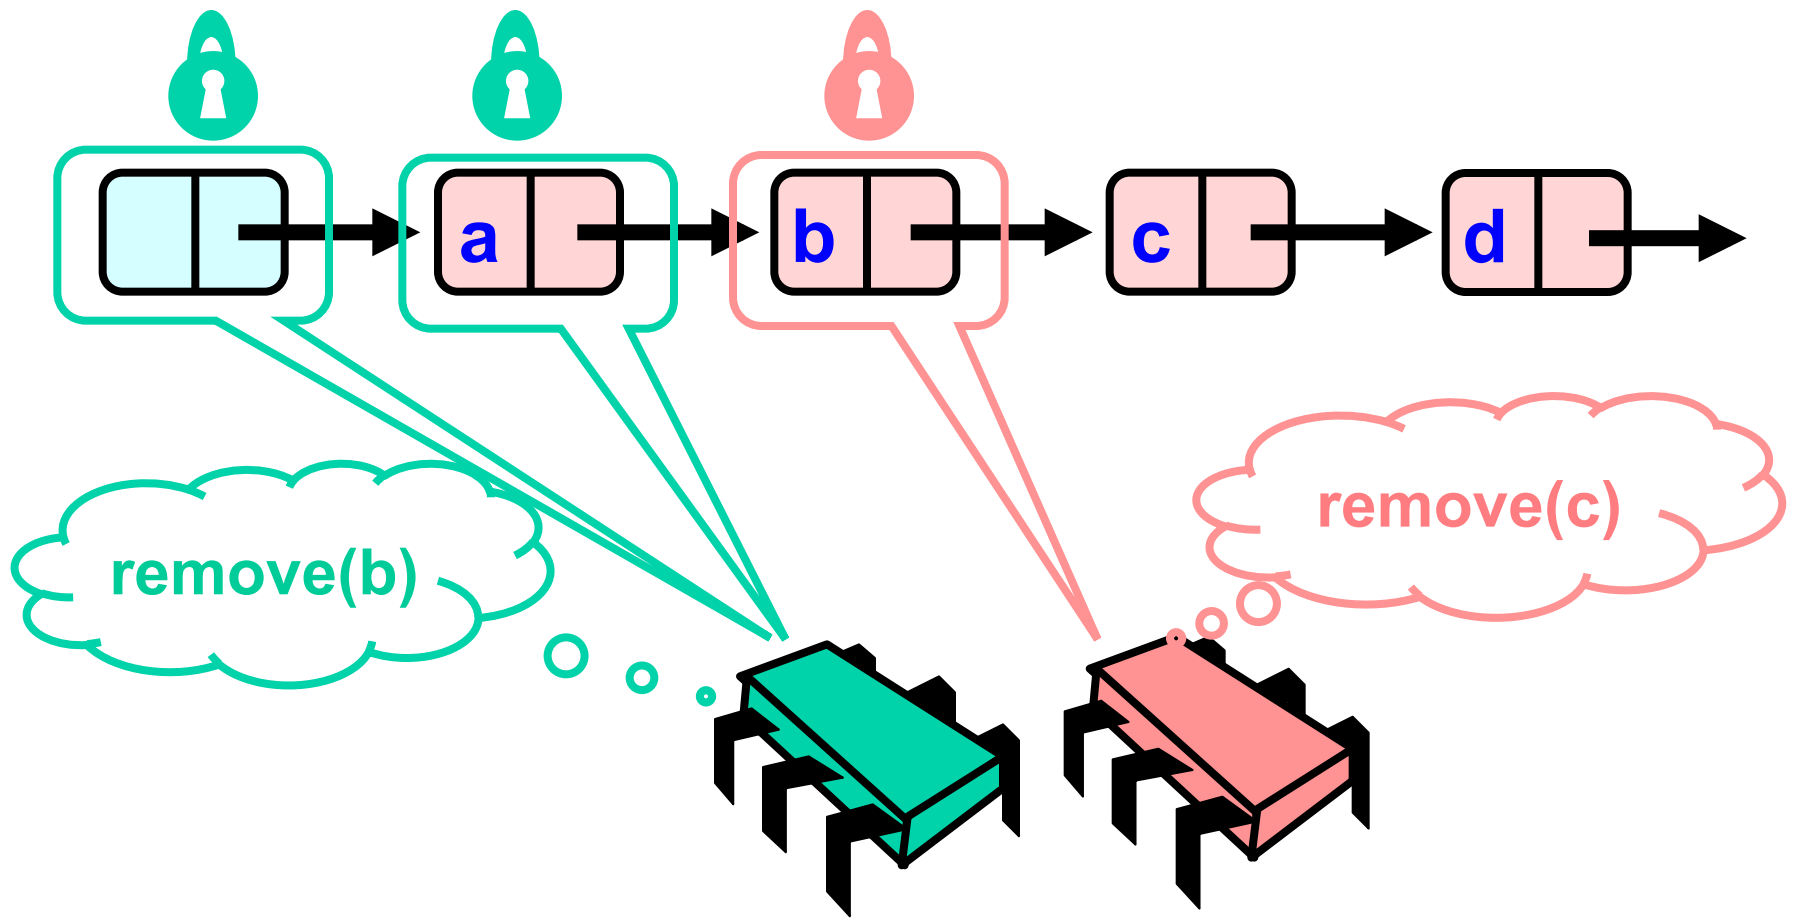
# Removing a Node



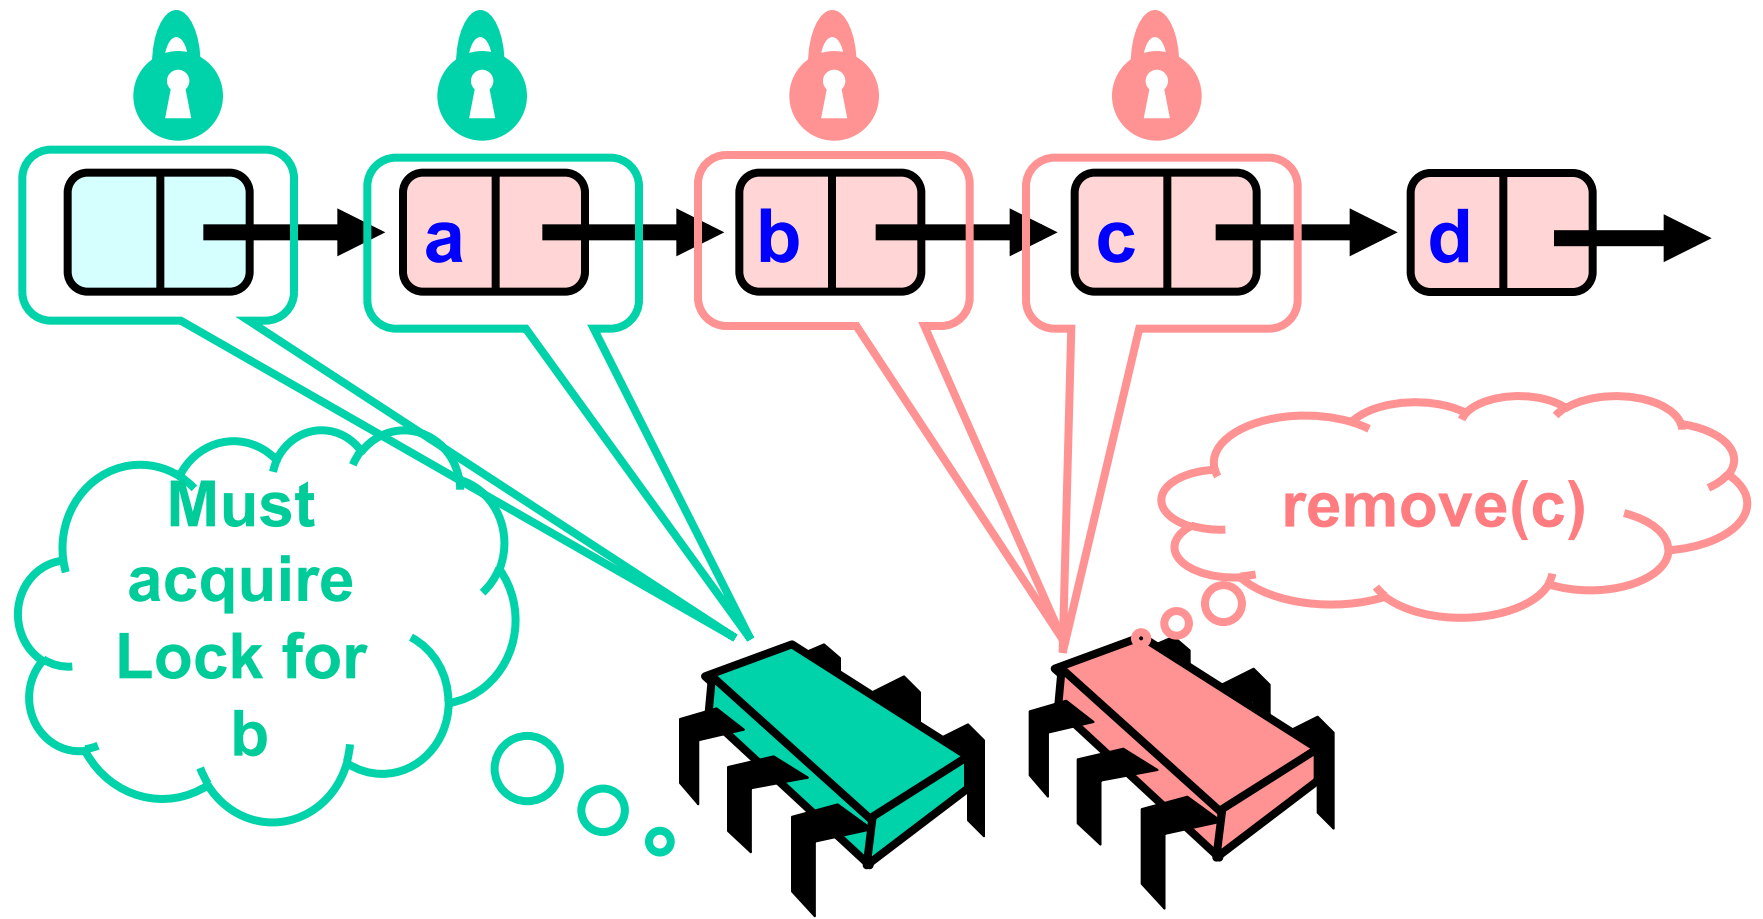
# Removing a Node



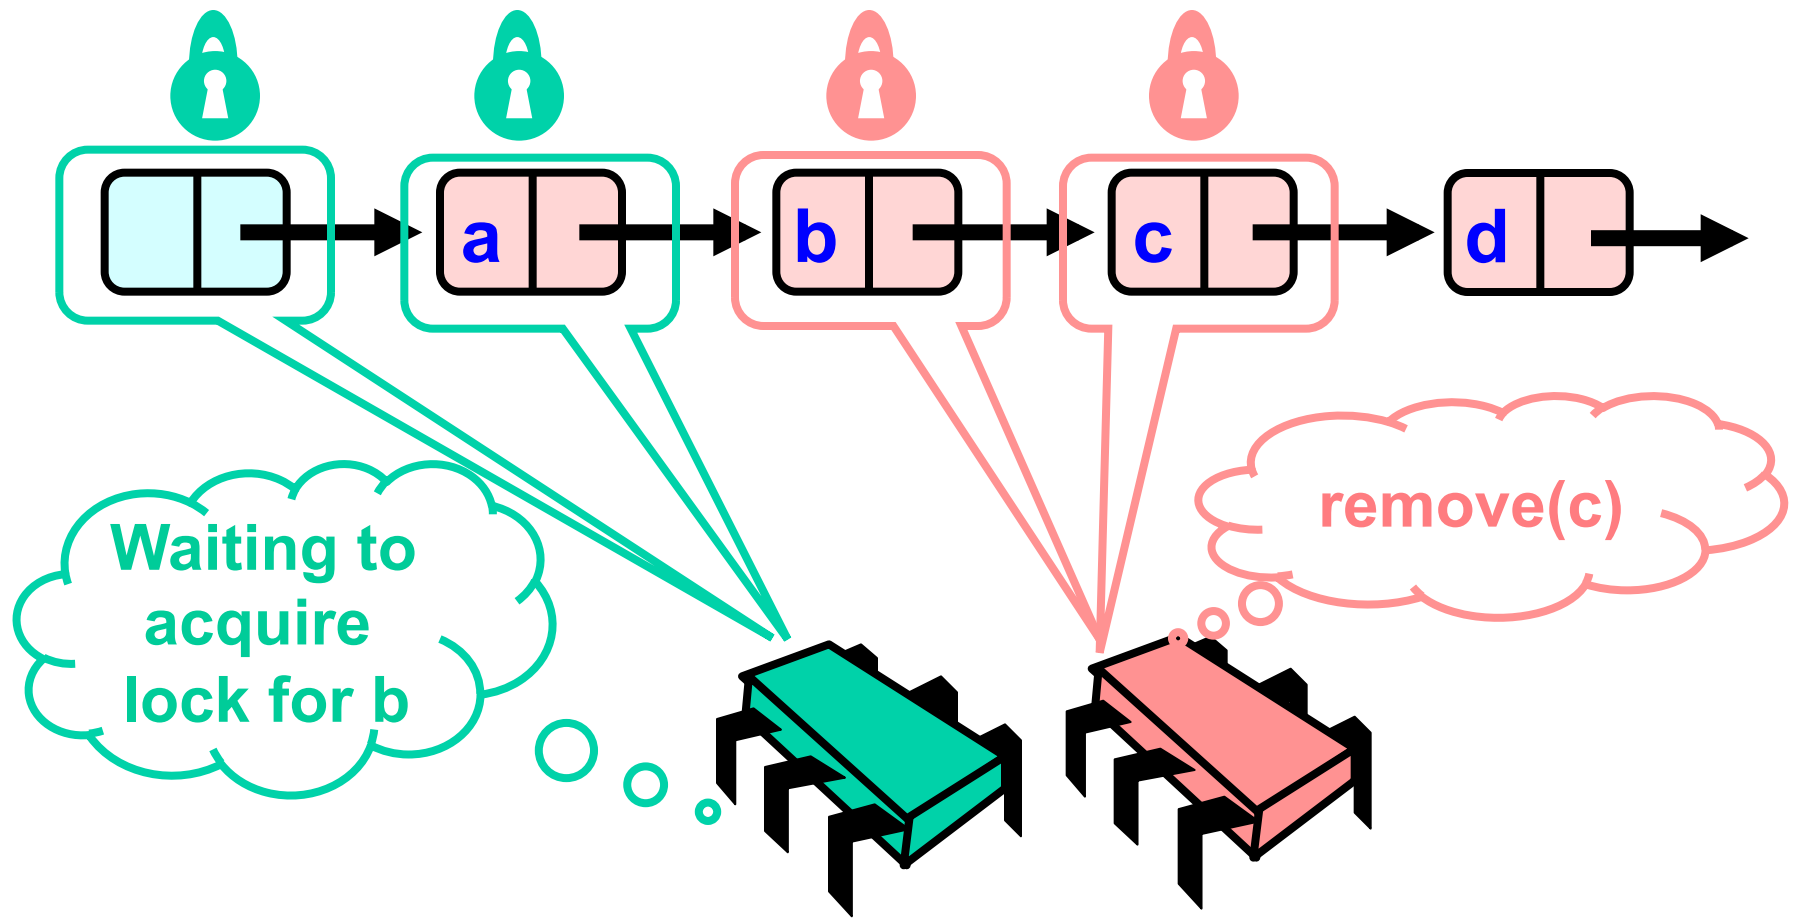
# Removing a Node



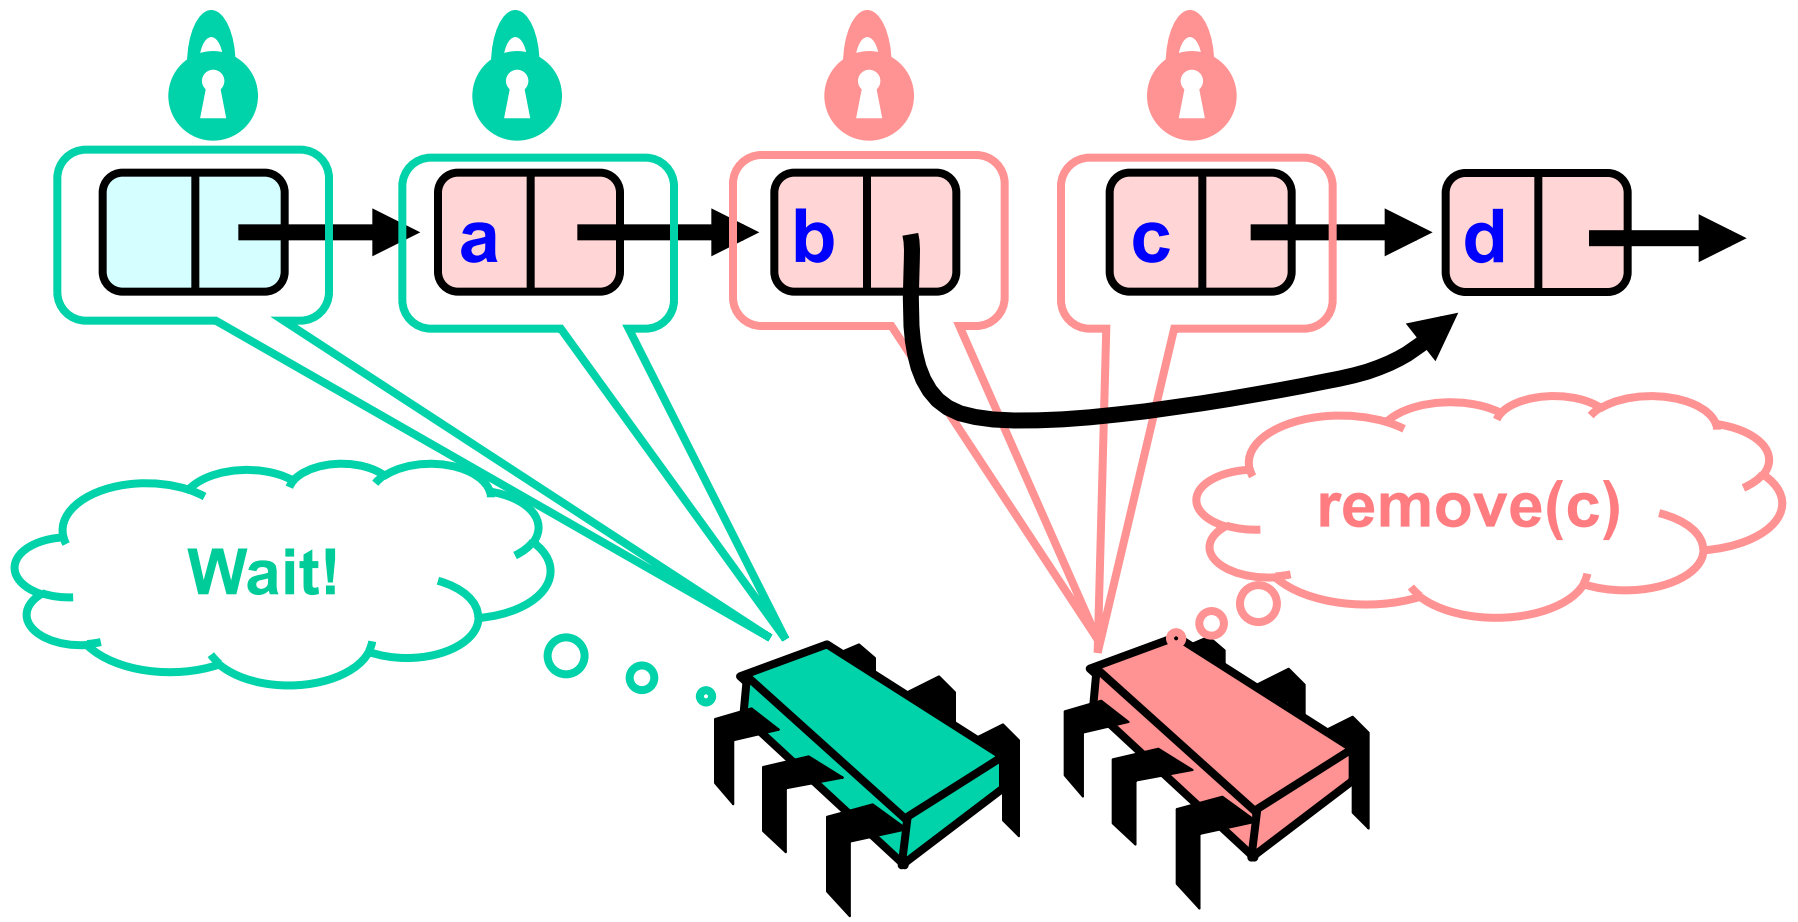
# Removing a Node



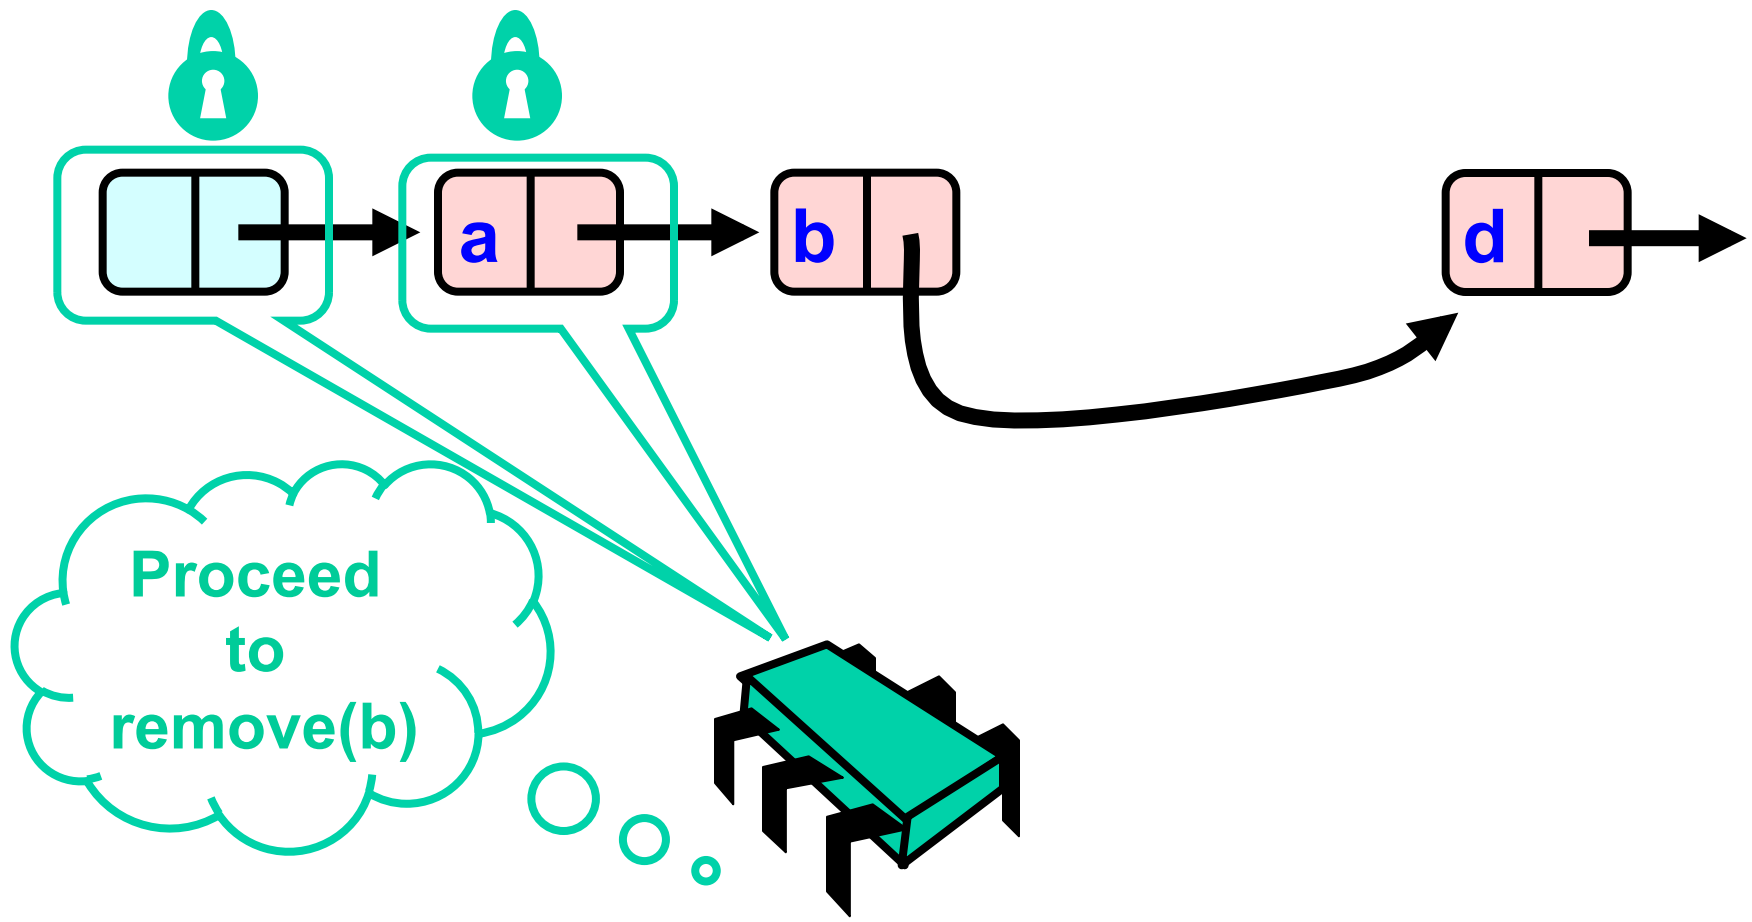
# Removing a Node



# Removing a Node




# Removing a Node

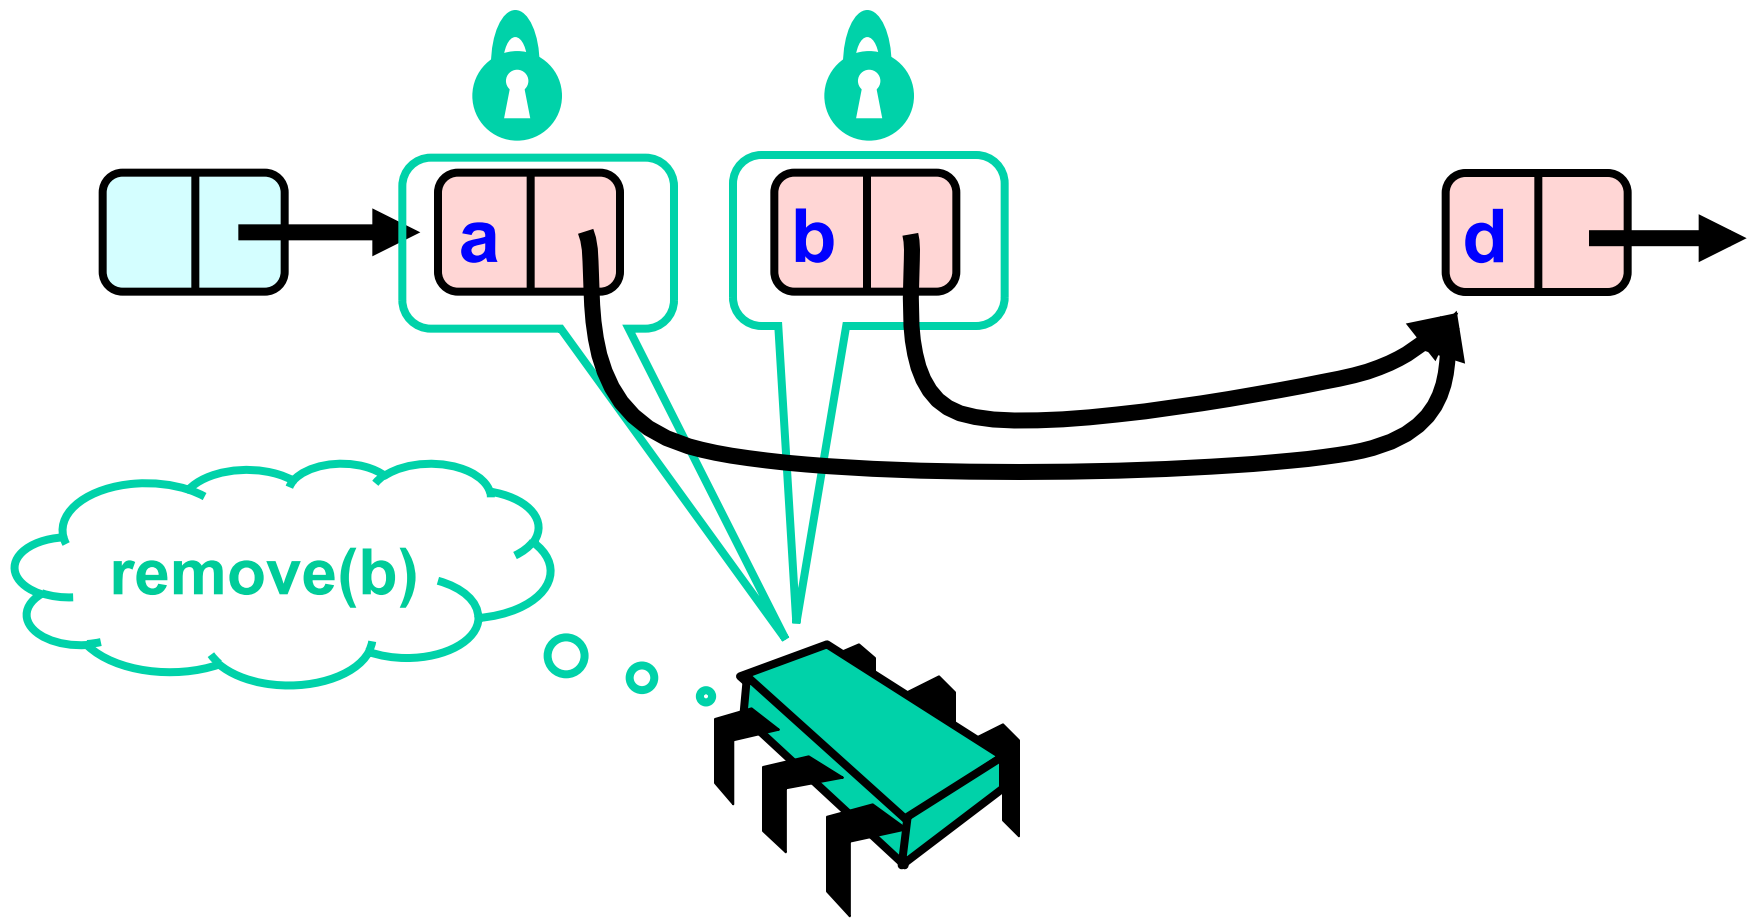




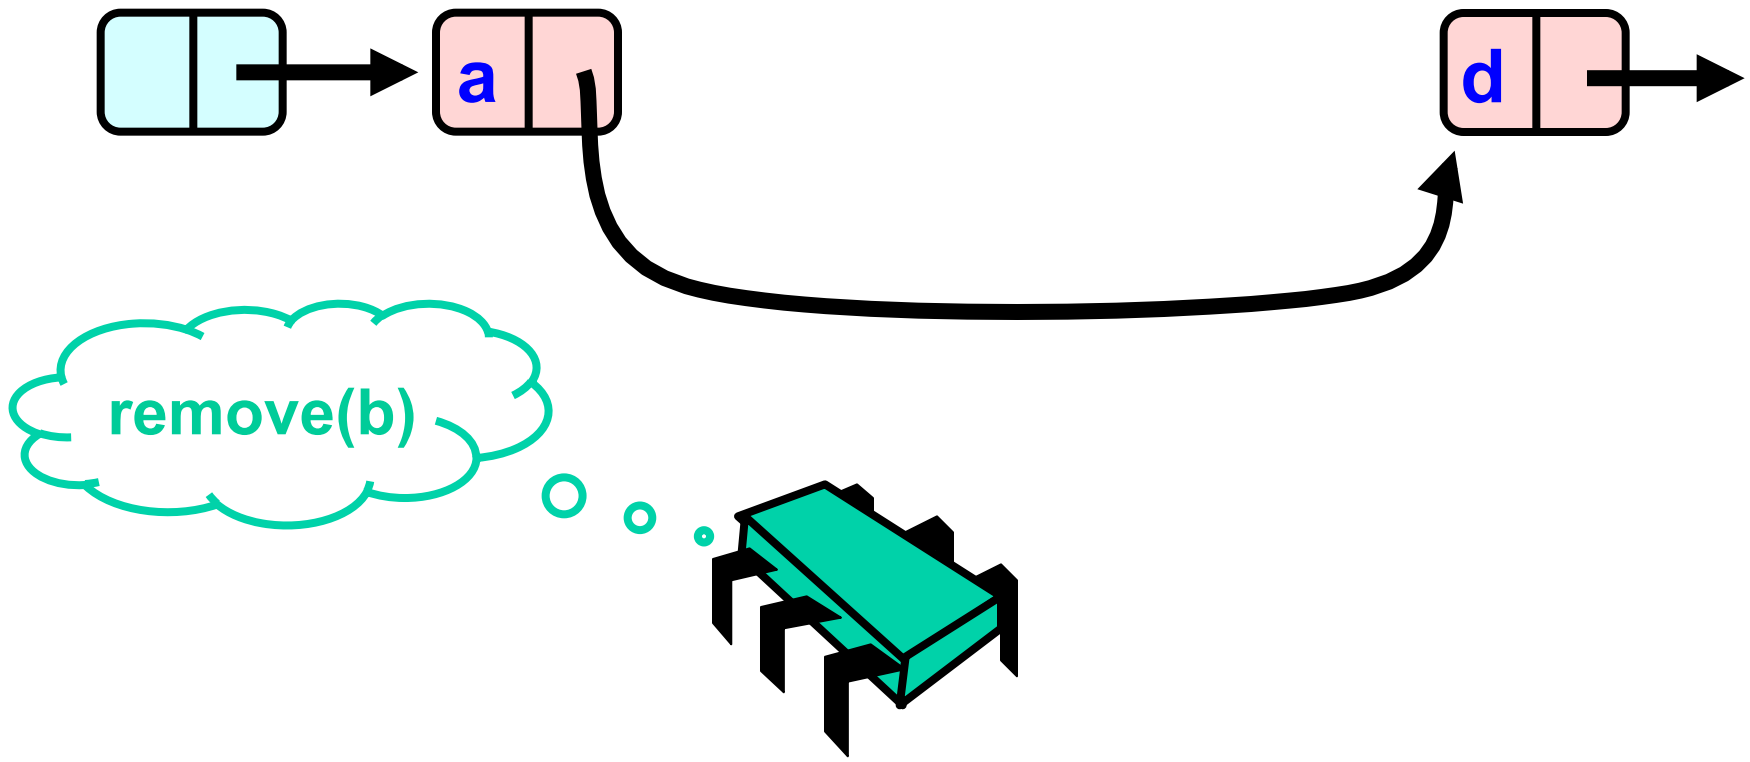
# Removing a Node



# Removing a Node



# Removing a Node



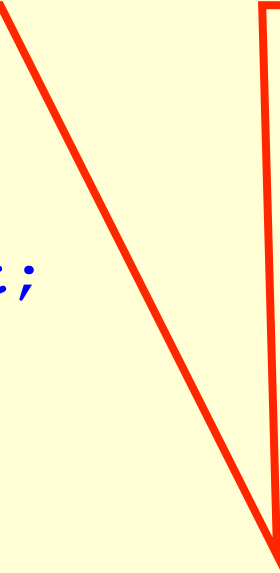
# Remove method

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

# Remove method

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Key used to order node**



# Remove method

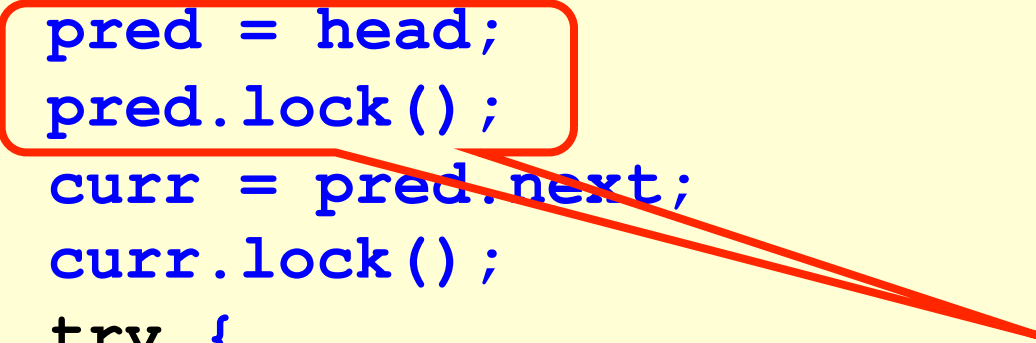
```
public boolean remove(T item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```



**Predecessor and  
current nodes**

# Remove method

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```



**lock pred == head  
before accessing its  
next field**

# Remove method

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**lock the node after head**

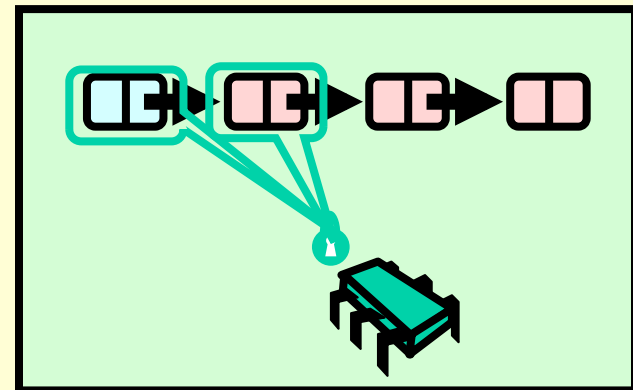




# Remove method

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

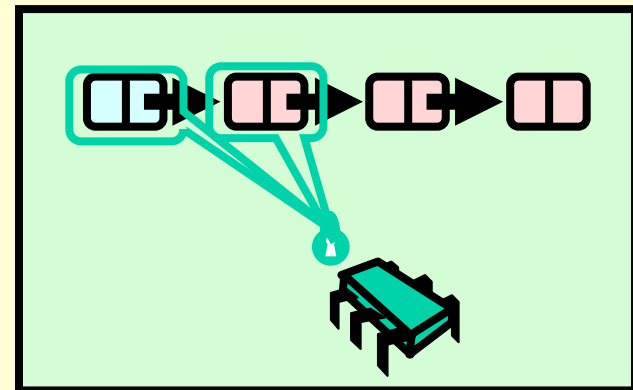
**When enter try,  
we hold locks on  
pred and curr**



# Remove method

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Traverse the rest  
of the list**



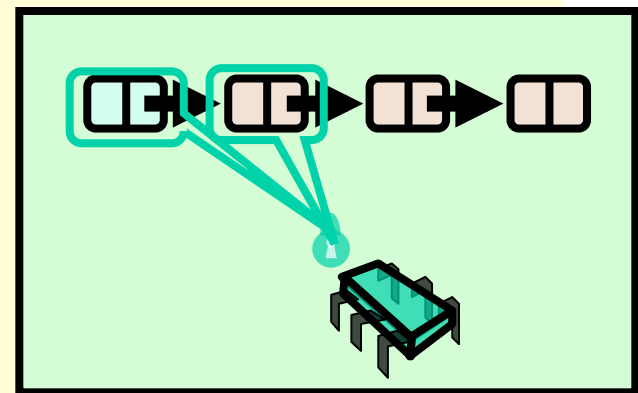
# Remove method

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    Node pred, curr;  
    pred = head;  
    pred.lock();  
    curr = pred.next;  
    curr.lock();  
    try {  
        ...  
    } finally {  
        curr.unlock();  
        pred.unlock();  
    }  
}
```

**Make sure  
locks released**

# Remove: searching (Inside the Try Block)

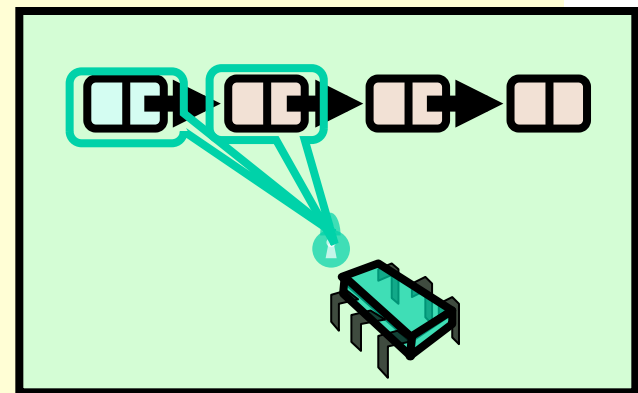
```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```



# Remove: searching (Inside the Try Block)

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

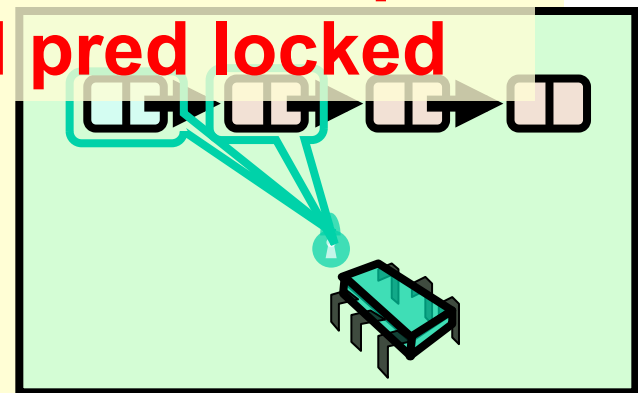
**Search key range**



# Remove: searching (Inside the Try Block)

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

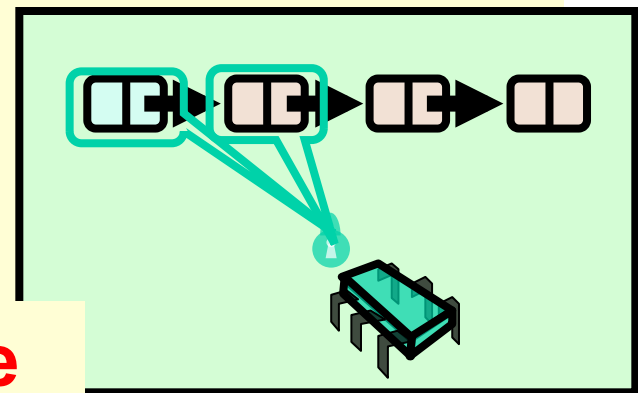
**At start of each loop:  
curr and pred locked**



# Remove: searching (Inside the Try Block)

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

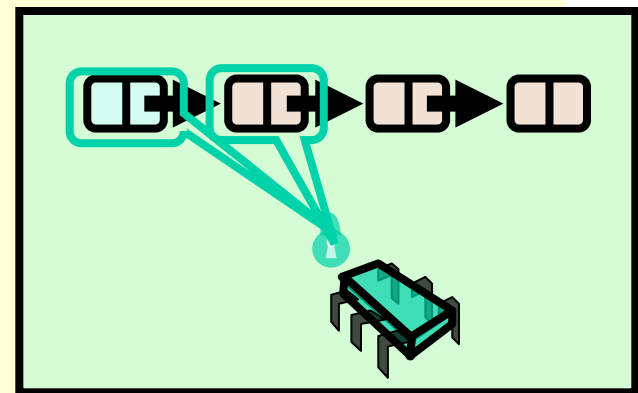
**If item found, remove node**



# Remove: searching (Inside the Try Block)

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Hand-over-hand  
locking again  
otherwise**

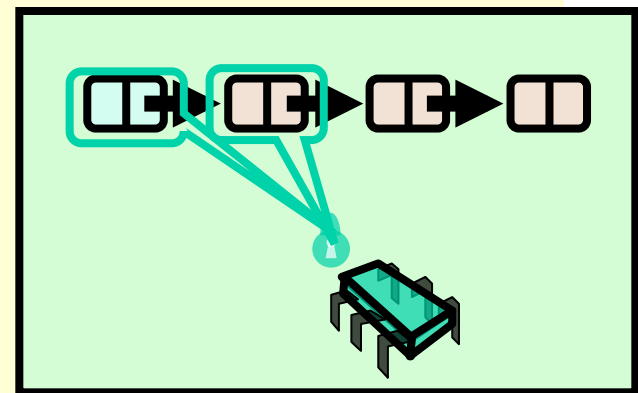




# Remove: searching (Inside the Try Block)

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Lock invariant  
restored**

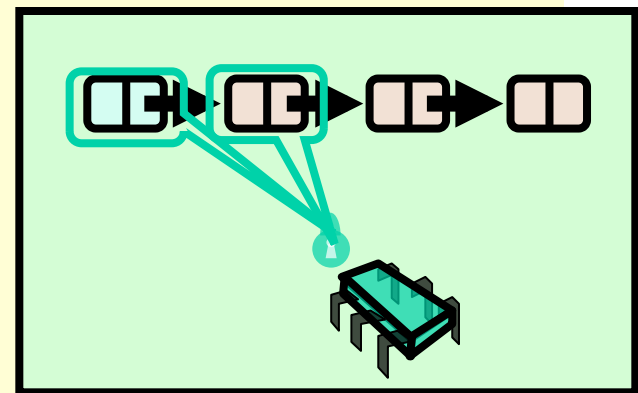


# Remove: searching (Inside the Try Block)

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

**Otherwise, not present**

**return false;**



# Aside: Next Field Must be Volatile!

```
public class Node {  
    public T item;  
    public int key;  
    public volatile Node next;  
}
```

Since we are no longer holding a lock when we read the "next" field, it needs to be volatile to avoid race conditions (more on that in future lecture).

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

- pred reachable from head
- curr is pred.next
- So curr.item is in the set

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Linearization point if  
item is present**

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Node locked, so no other  
thread can remove it ....**

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}
```

**return false;**

**Item not present**

# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

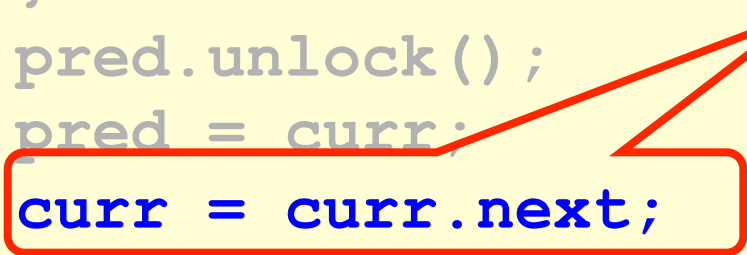
- pred reachable from head
- curr is pred.next
- pred.key < key
- key < curr.key



# Why remove() is linearizable

```
while (curr.key <= key) {  
    if (item == curr.item) {  
        pred.next = curr.next;  
        return true;  
    }  
    pred.unlock();  
    pred = curr;  
    curr = curr.next;  
    curr.lock();  
}  
return false;
```

**Linearization point:  
the most recent  
read before return**

A red callout box with a pointer indicating the linearization point. The box is rectangular with a red border and contains the text 'curr = curr.next;'. A red line extends from the top-right corner of the box, pointing towards the right side of the slide.

# Adding Nodes

- To add node e
  - Must lock predecessor
  - Must lock successor
- Neither can be deleted
  - (Is successor lock actually required?)

# Abstraction Map

- $S(\text{head}) =$ 
  - $\{ x \mid \text{there exists } a \text{ such that}$ 
    - $a \text{ reachable from head and}$
    - $a.\text{item} = x$
  - $\}$

# Representation Invariant

- Easy to check that
  - tail always reachable from head
  - Nodes sorted, no duplicates

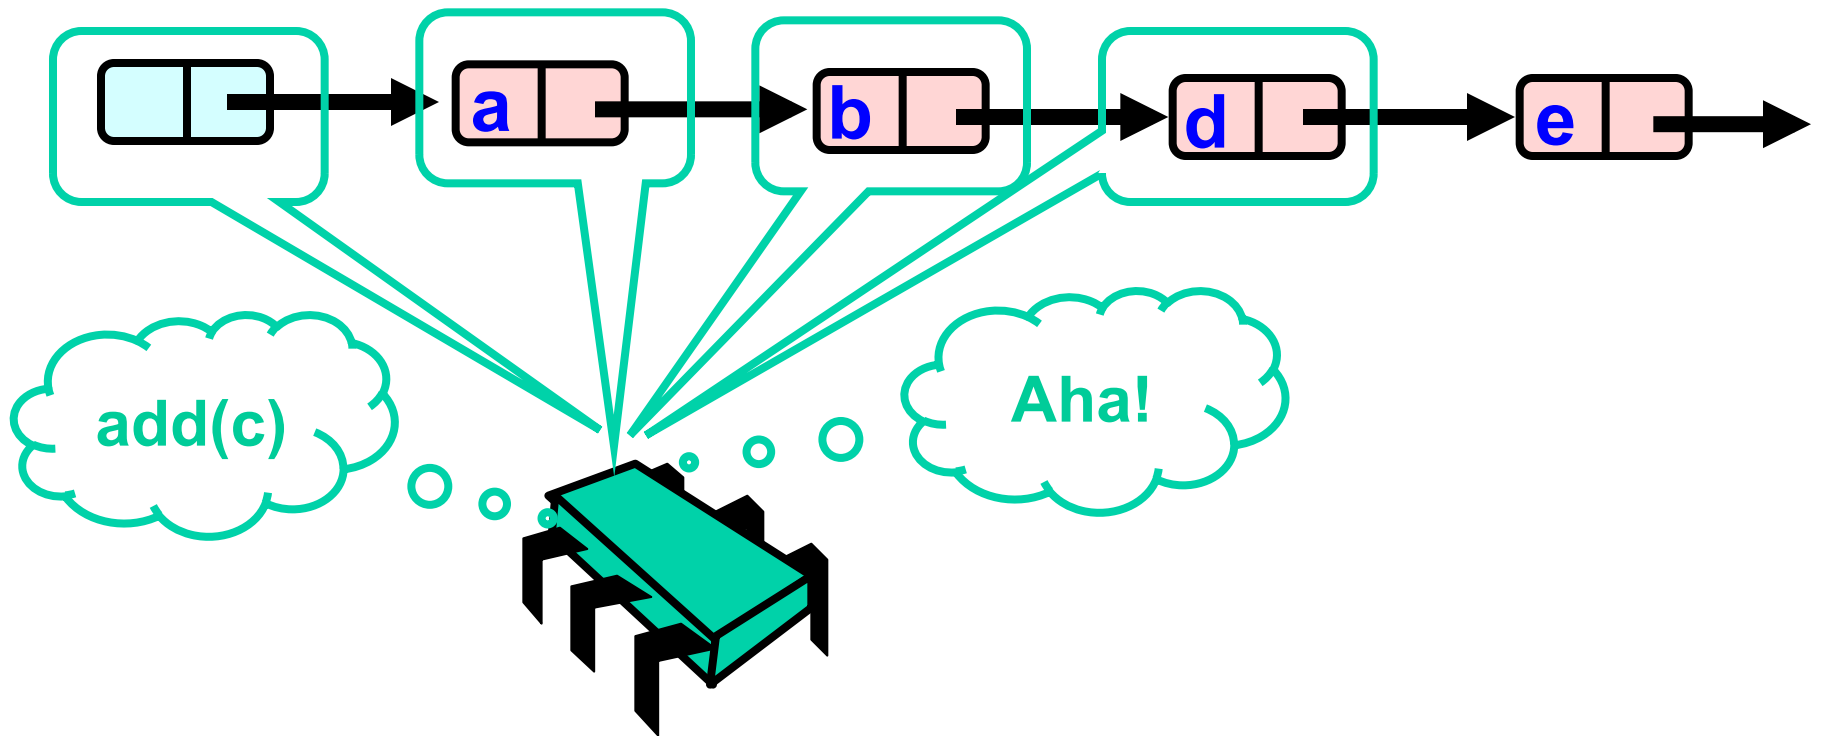
# Drawbacks

- Better than coarse-grained lock
  - Threads can traverse in parallel
- Still not ideal
  - Long chain of acquire/release
  - Inefficient

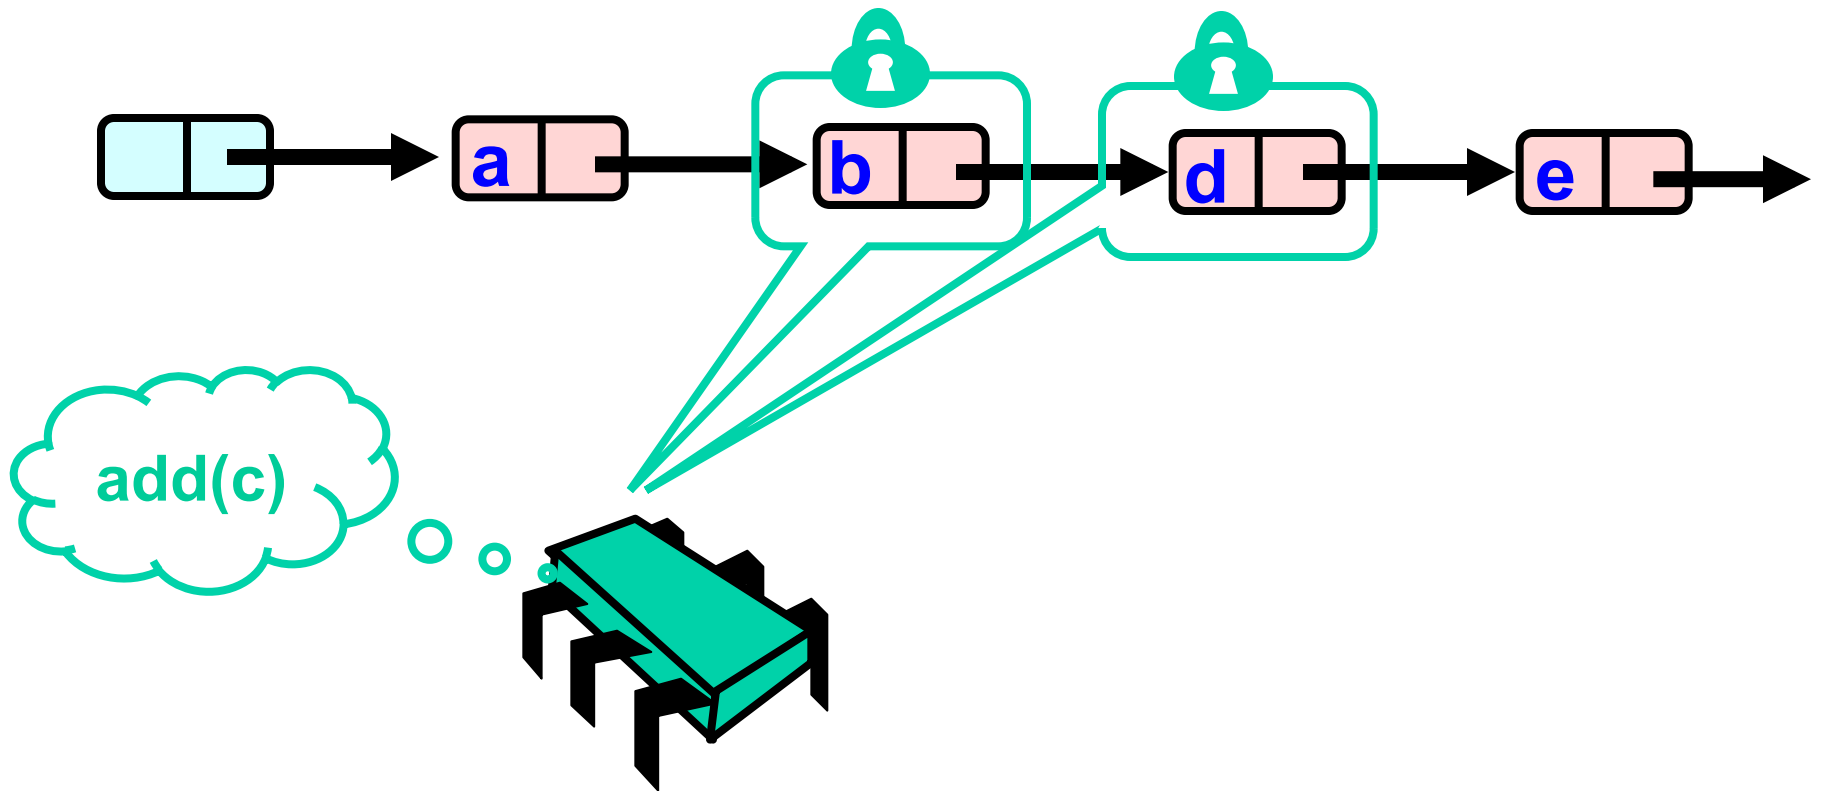
# Optimistic Synchronization

- Find nodes without locking
- Lock nodes
- Check that everything is OK

# Optimistic: Traverse without Locking

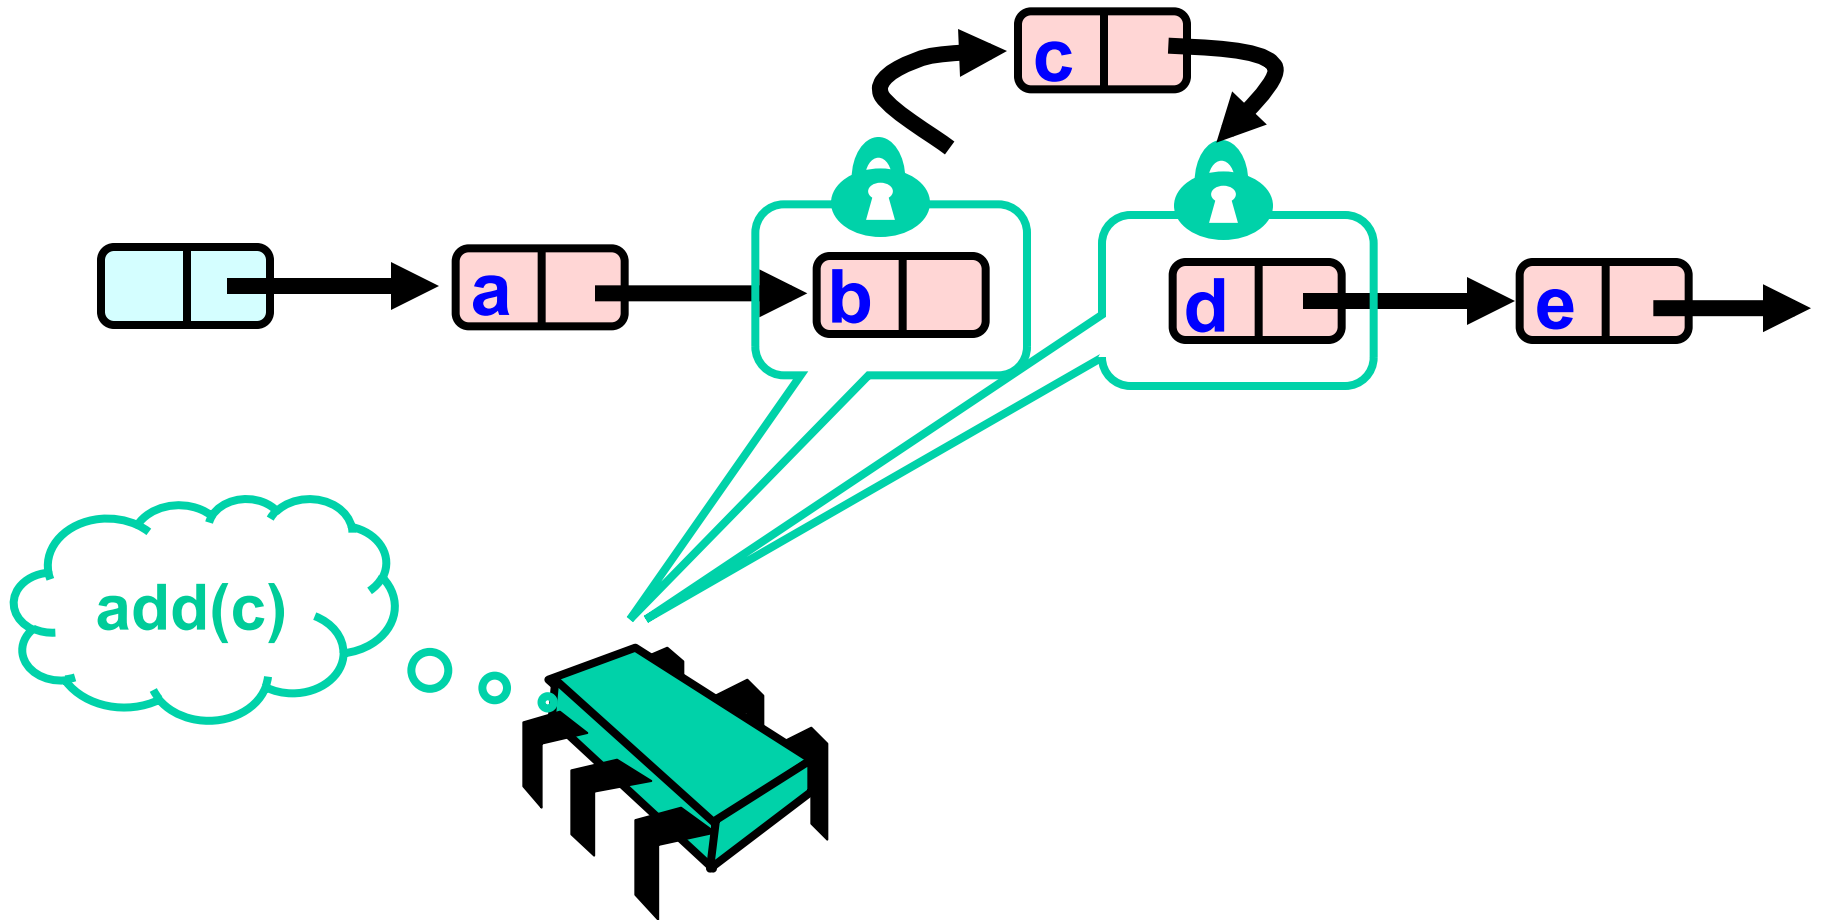


# Optimistic: Lock and Load

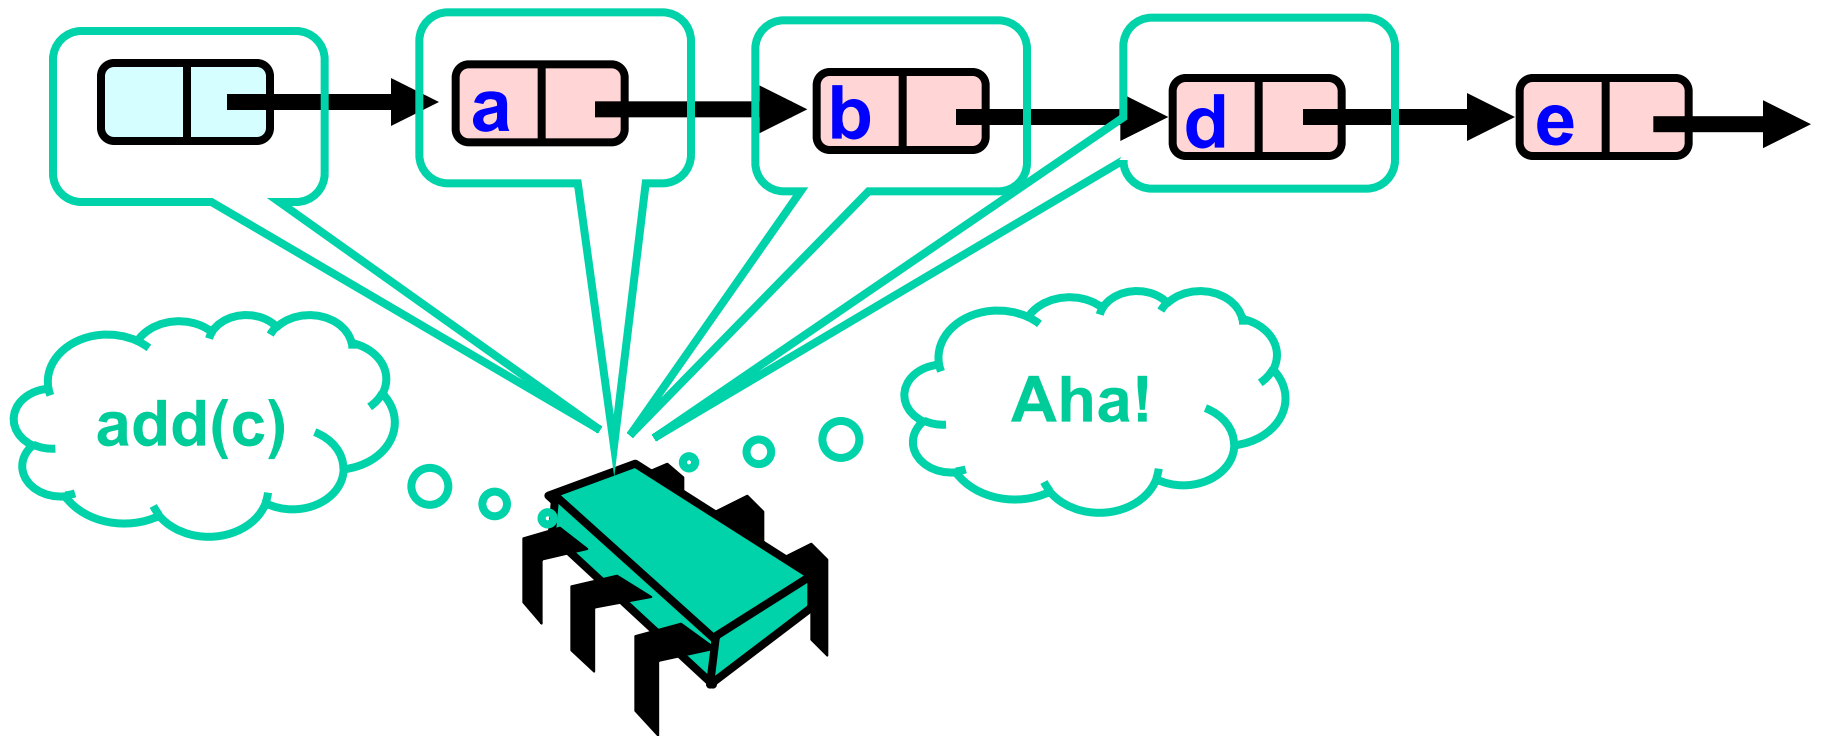




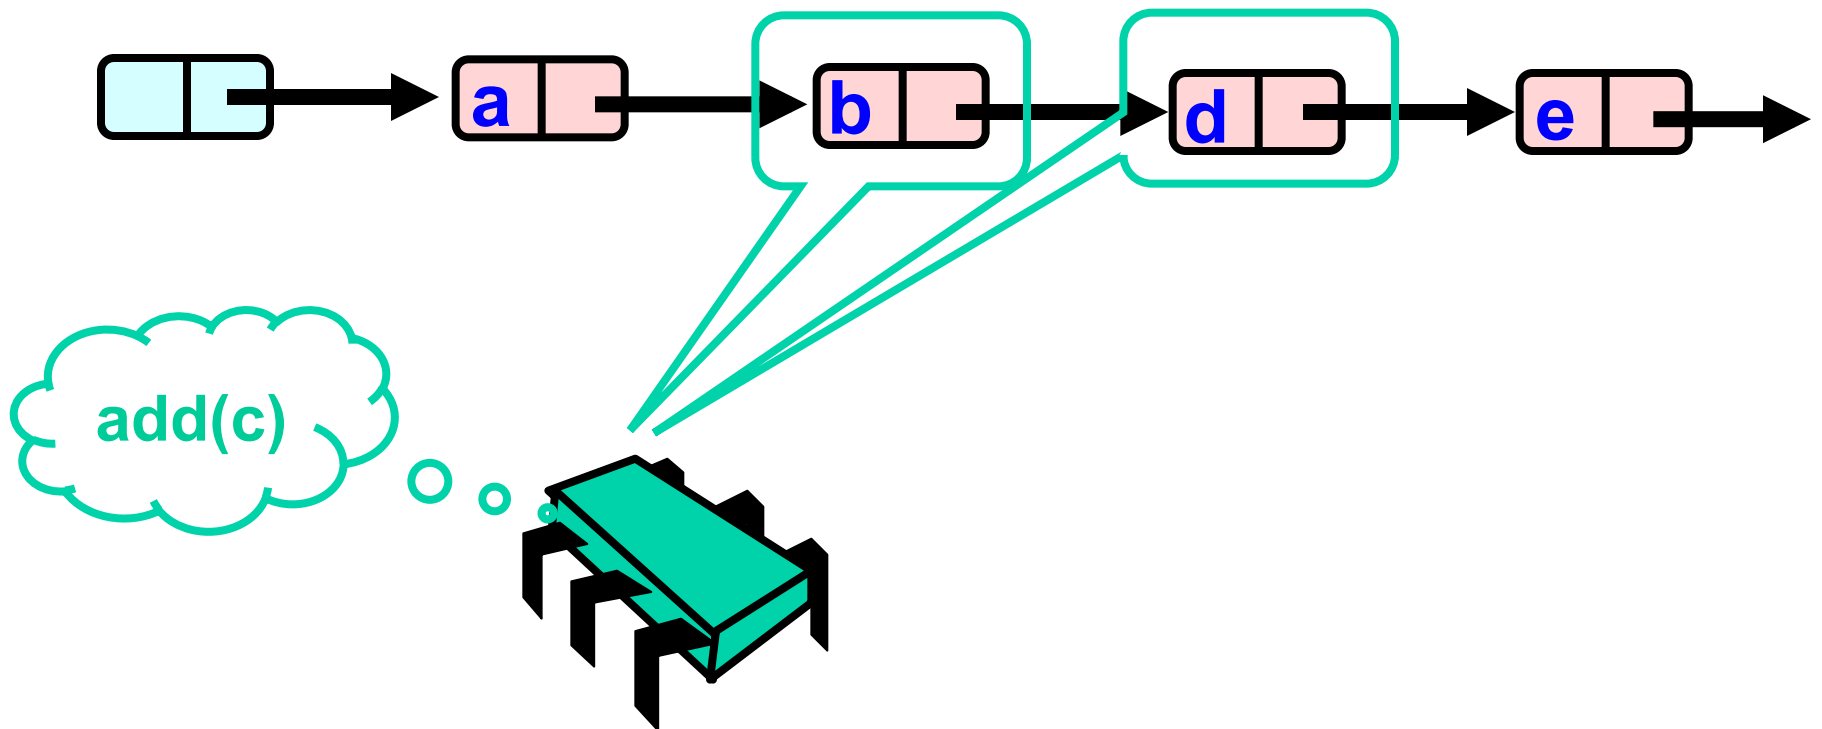
# Optimistic: Lock and Load



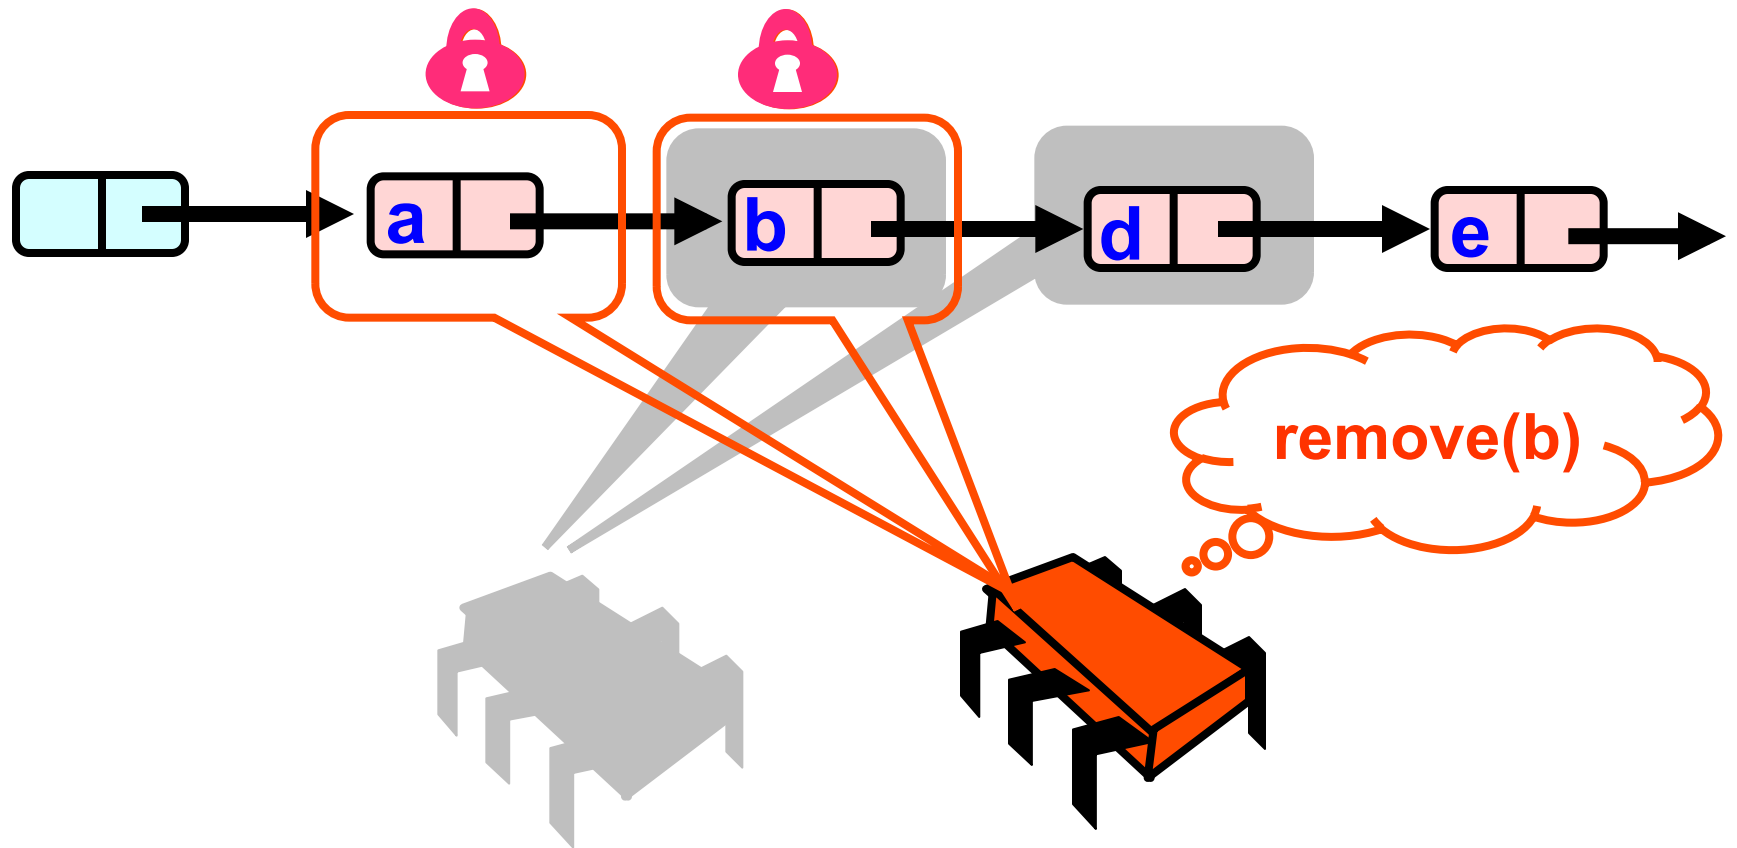
# What could go wrong?



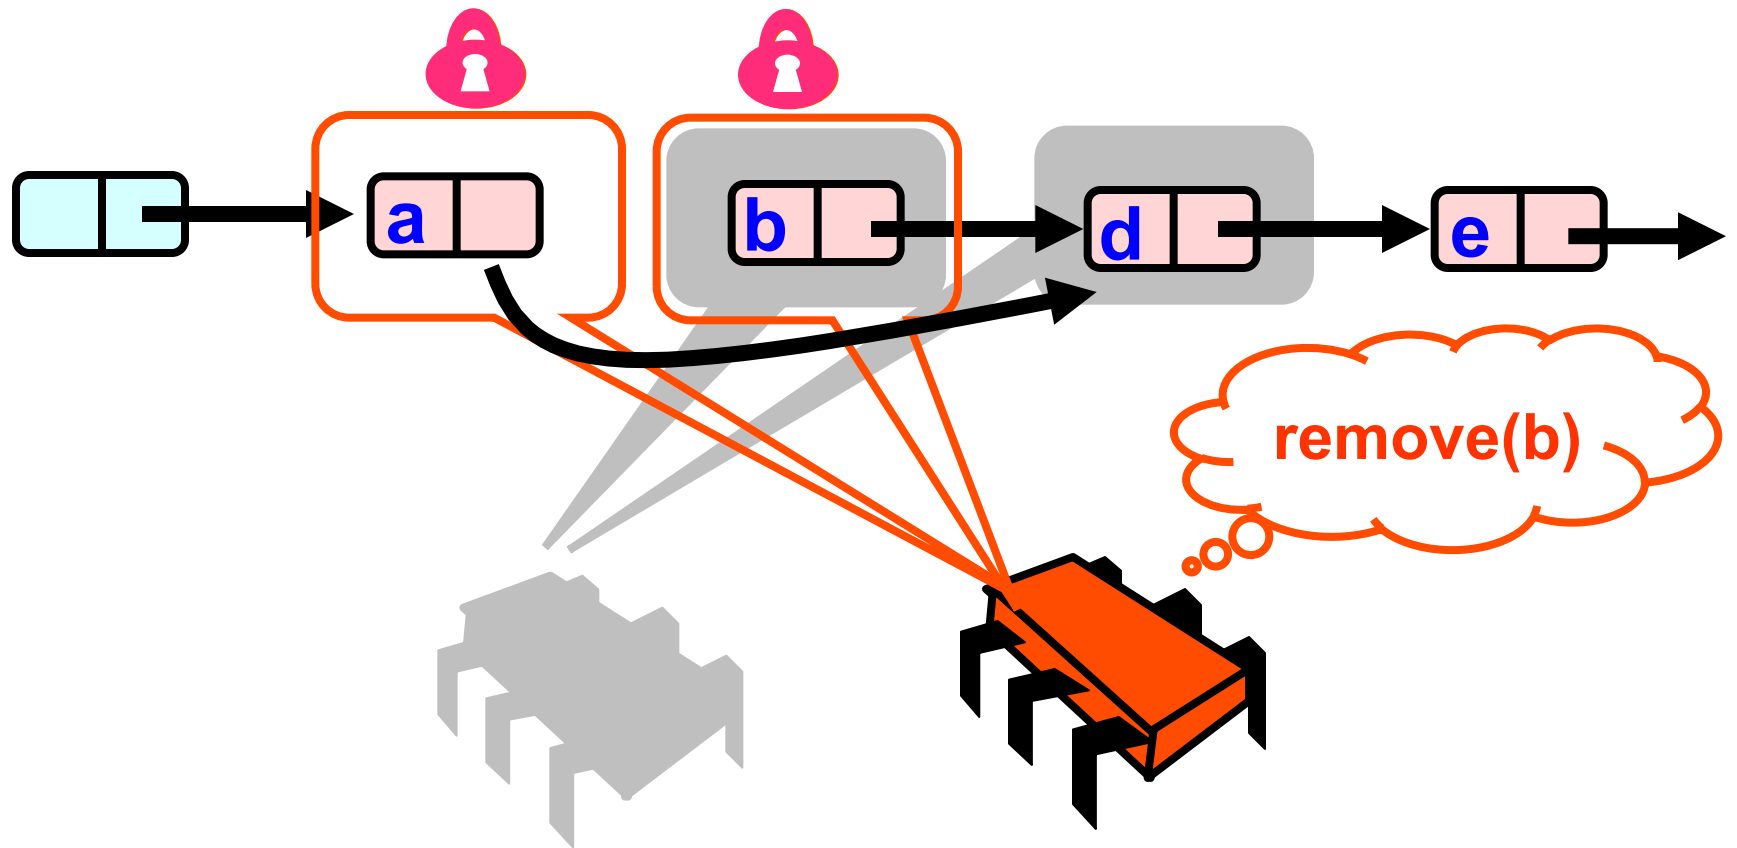
# What could go wrong?



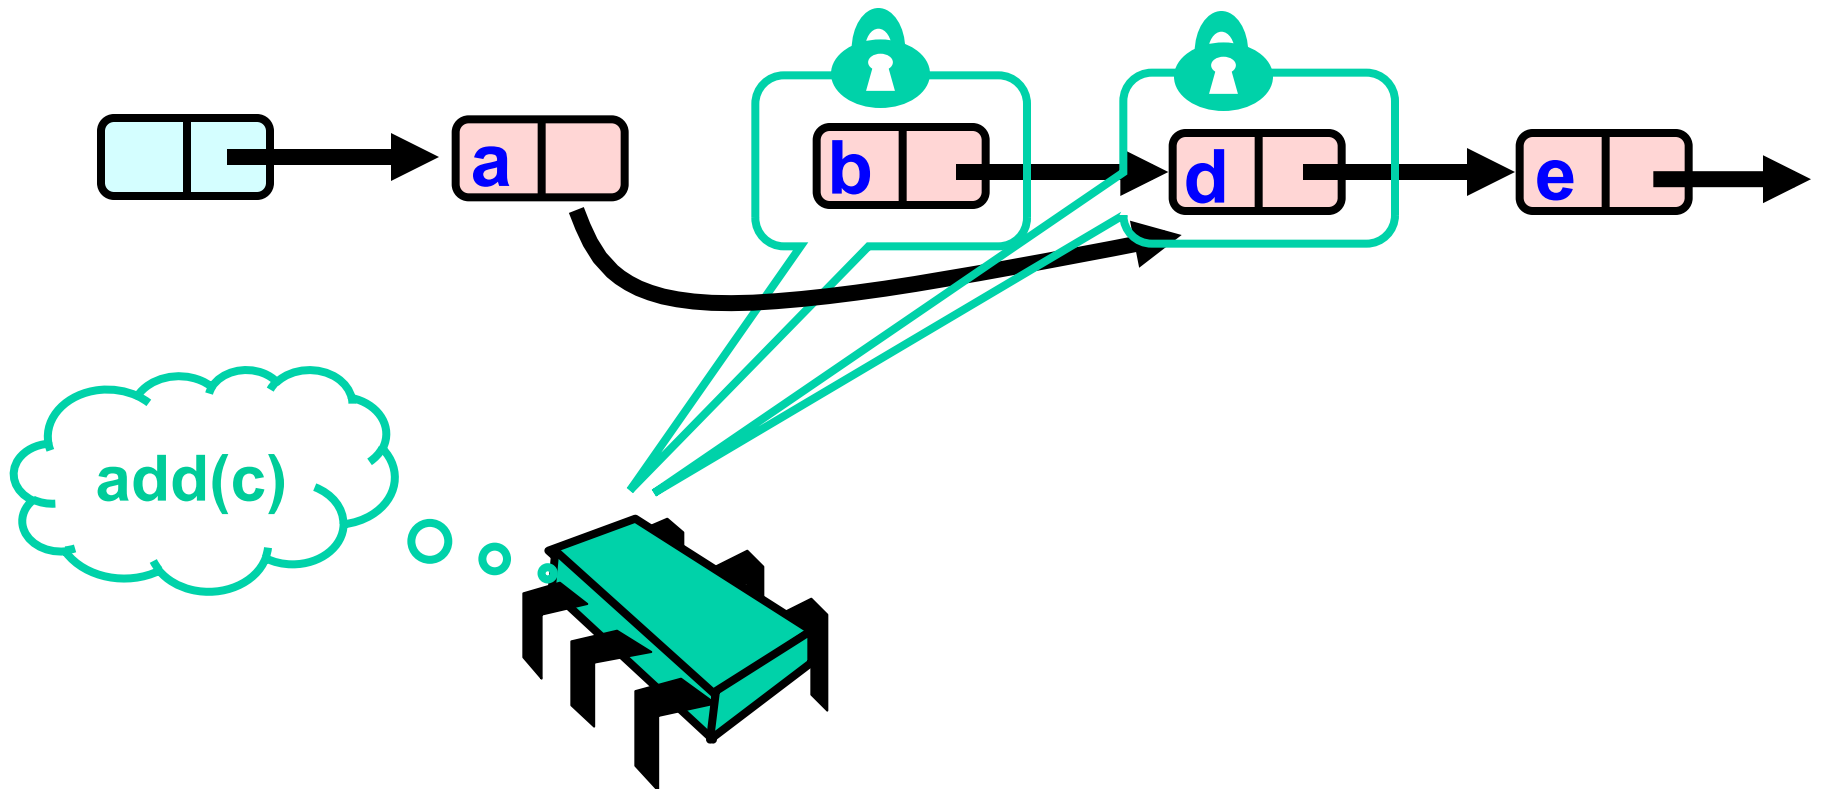
# What could go wrong?



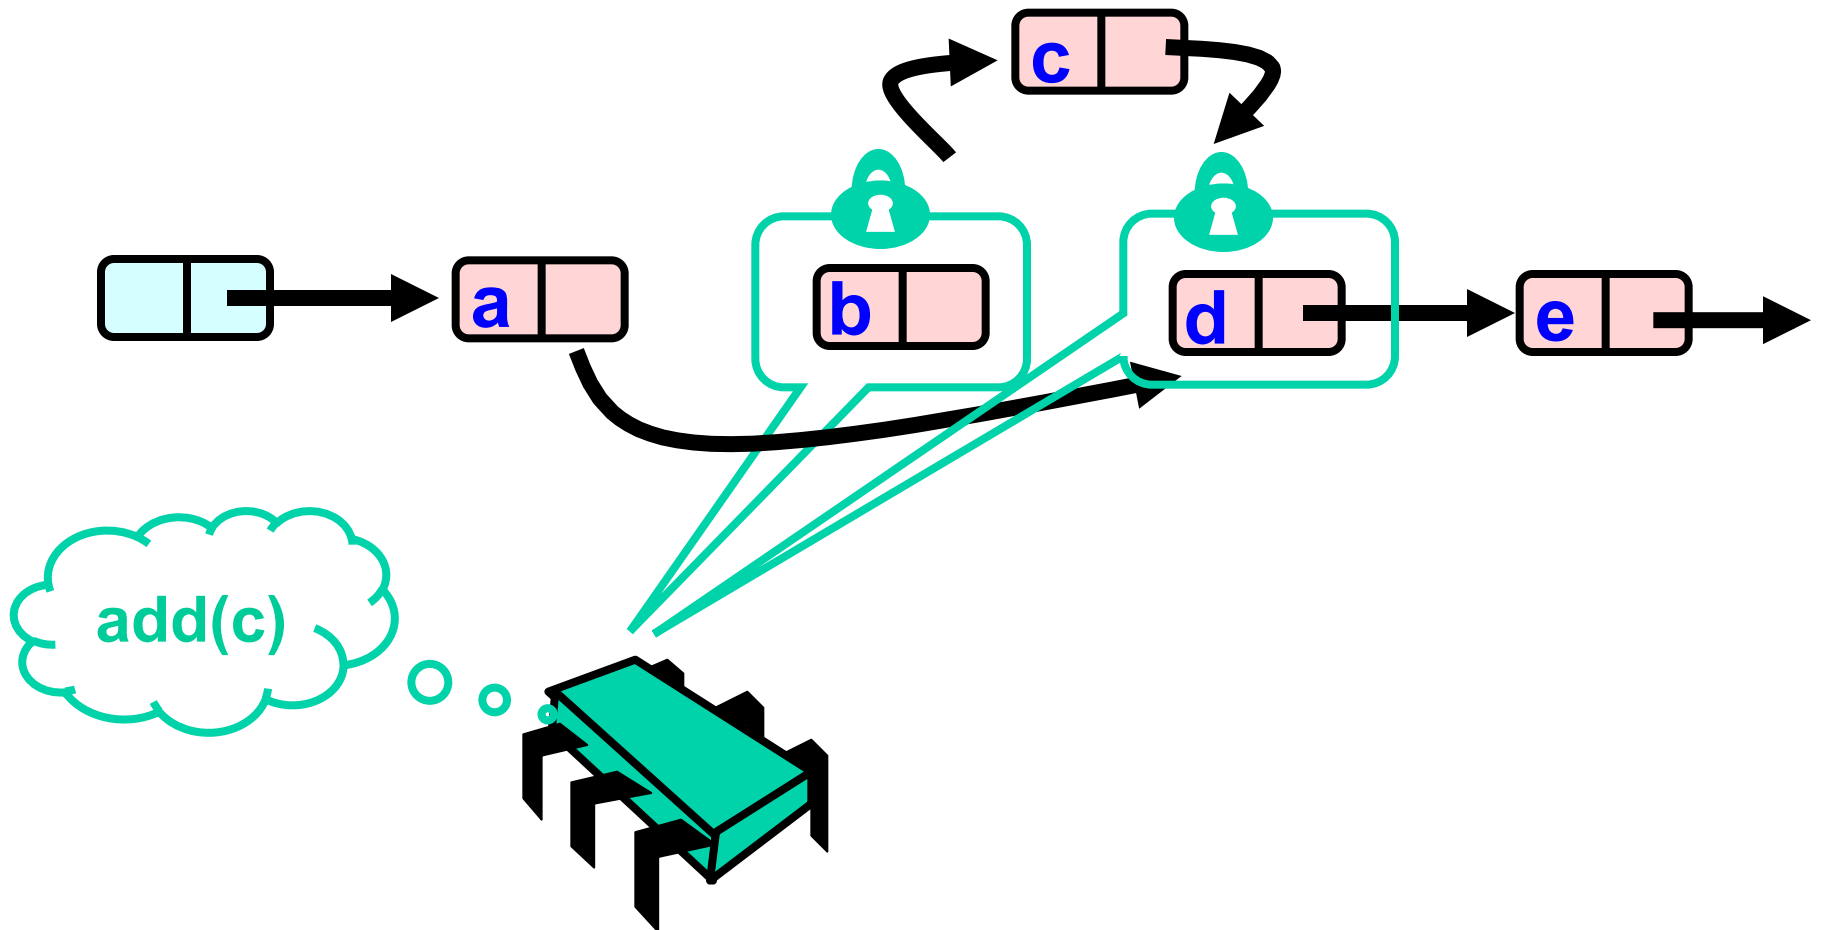
# What could go wrong?



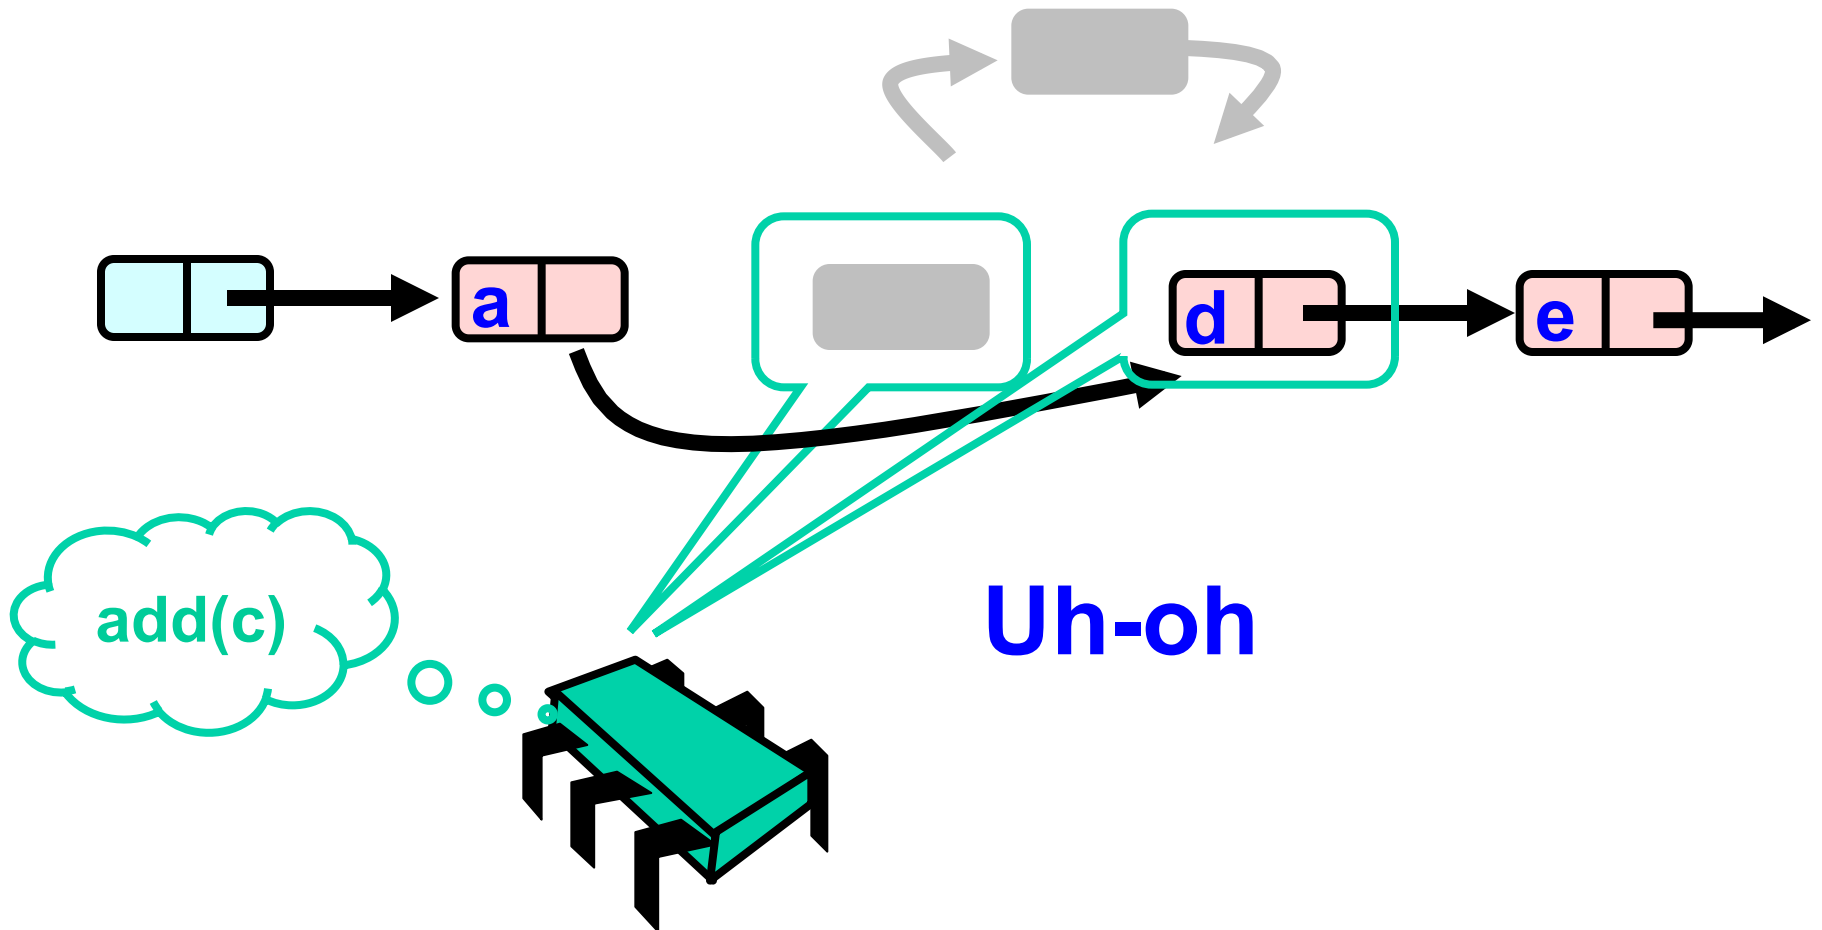
# What could go wrong?



# What could go wrong?

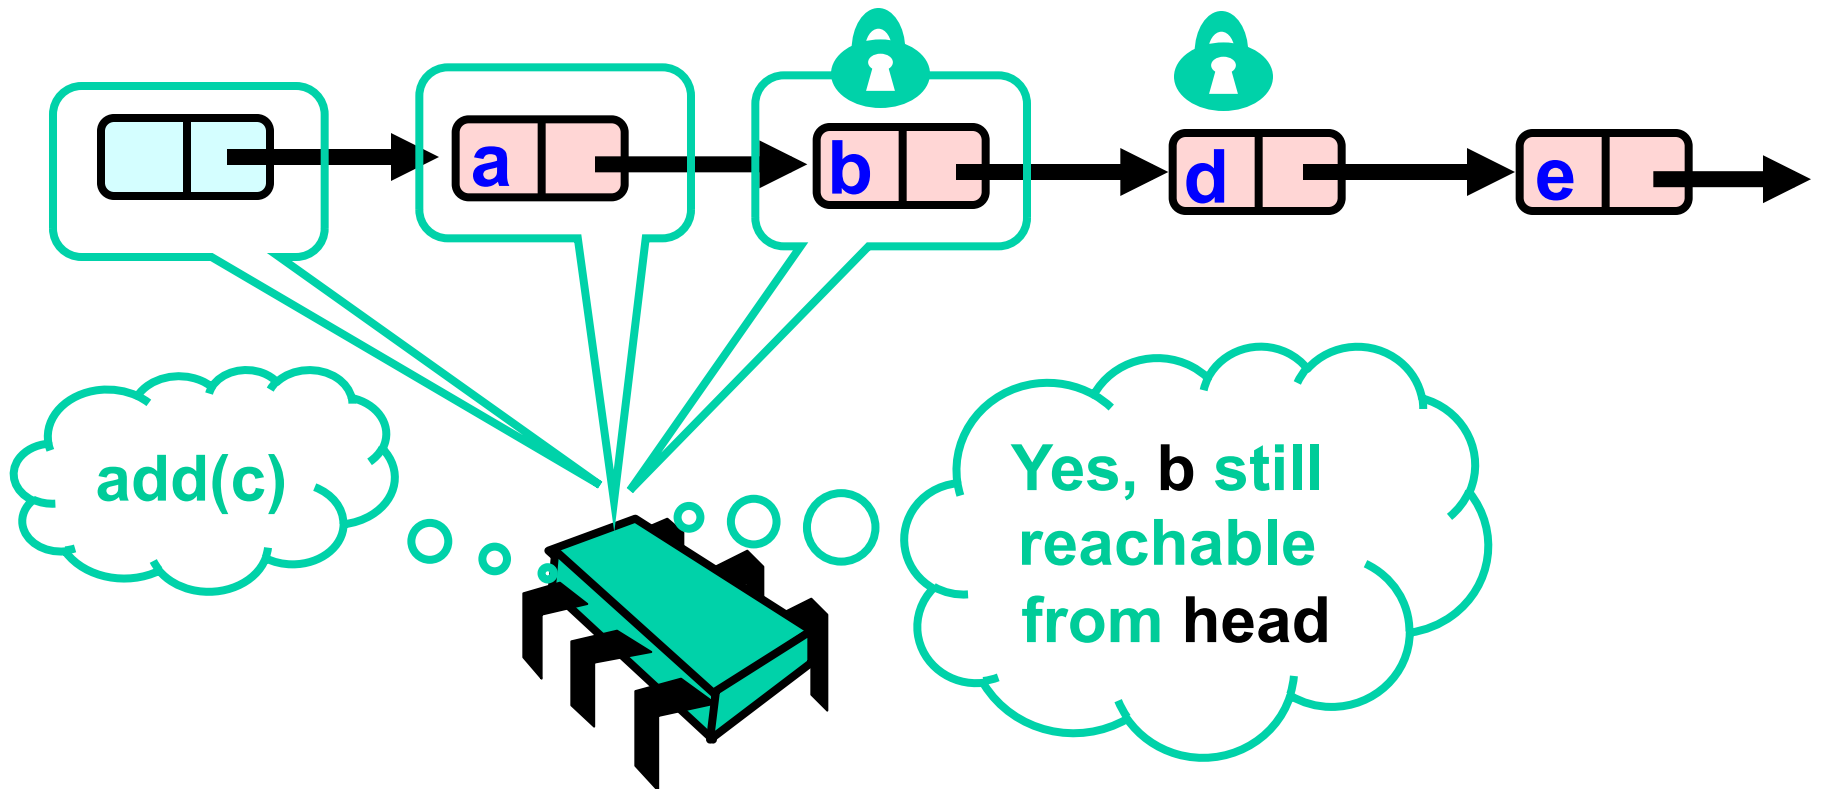


# What could go wrong?

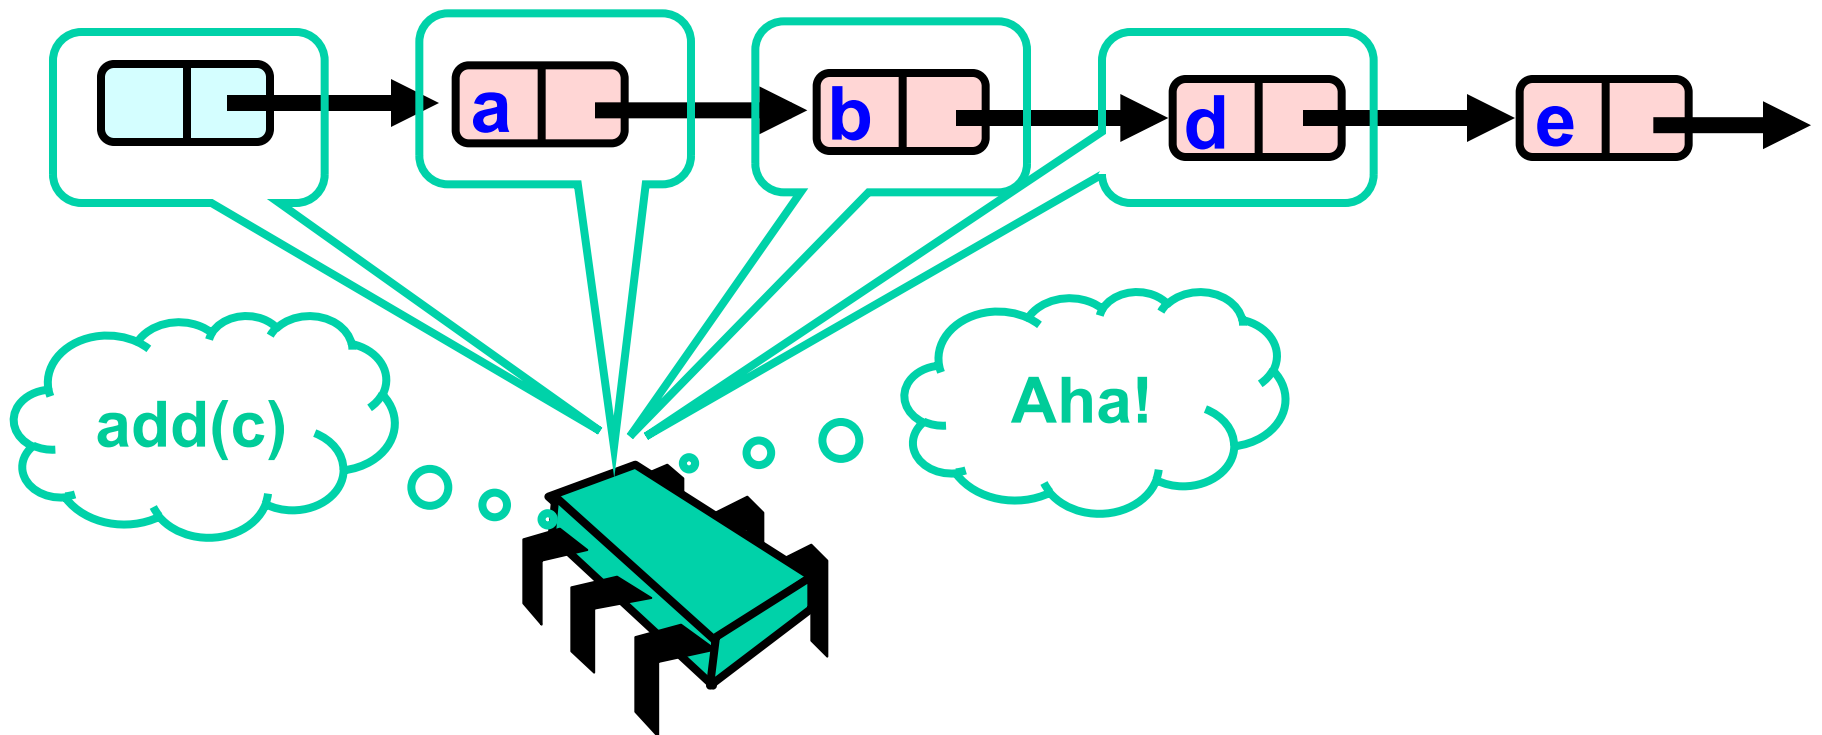




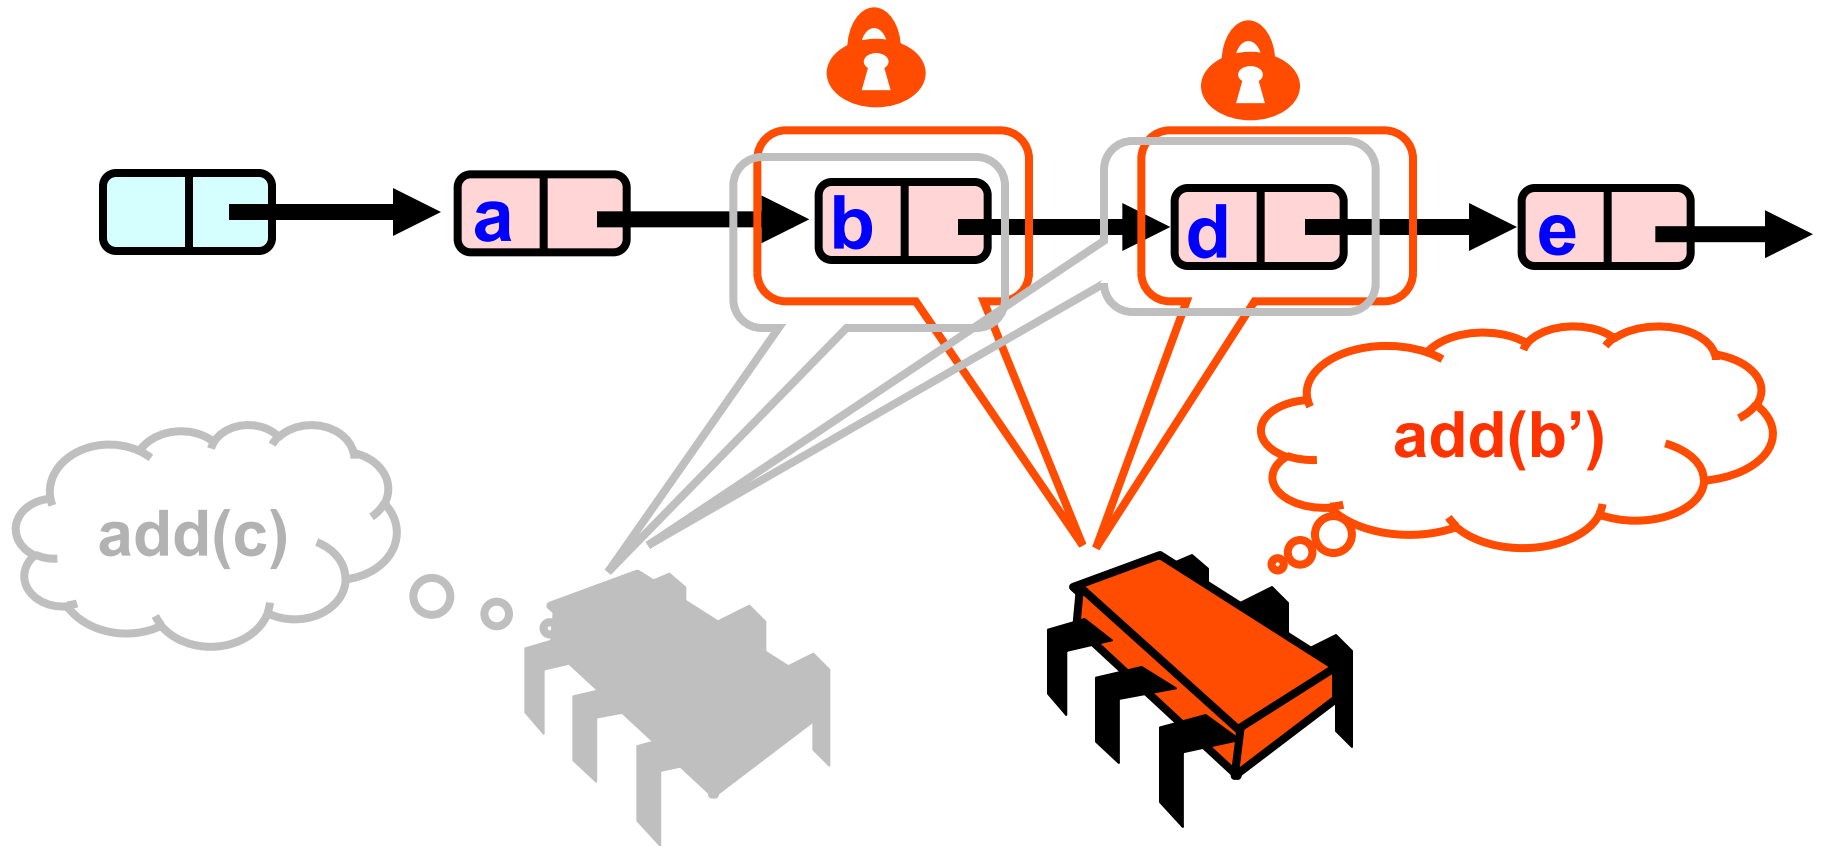
# Validate – Part 1



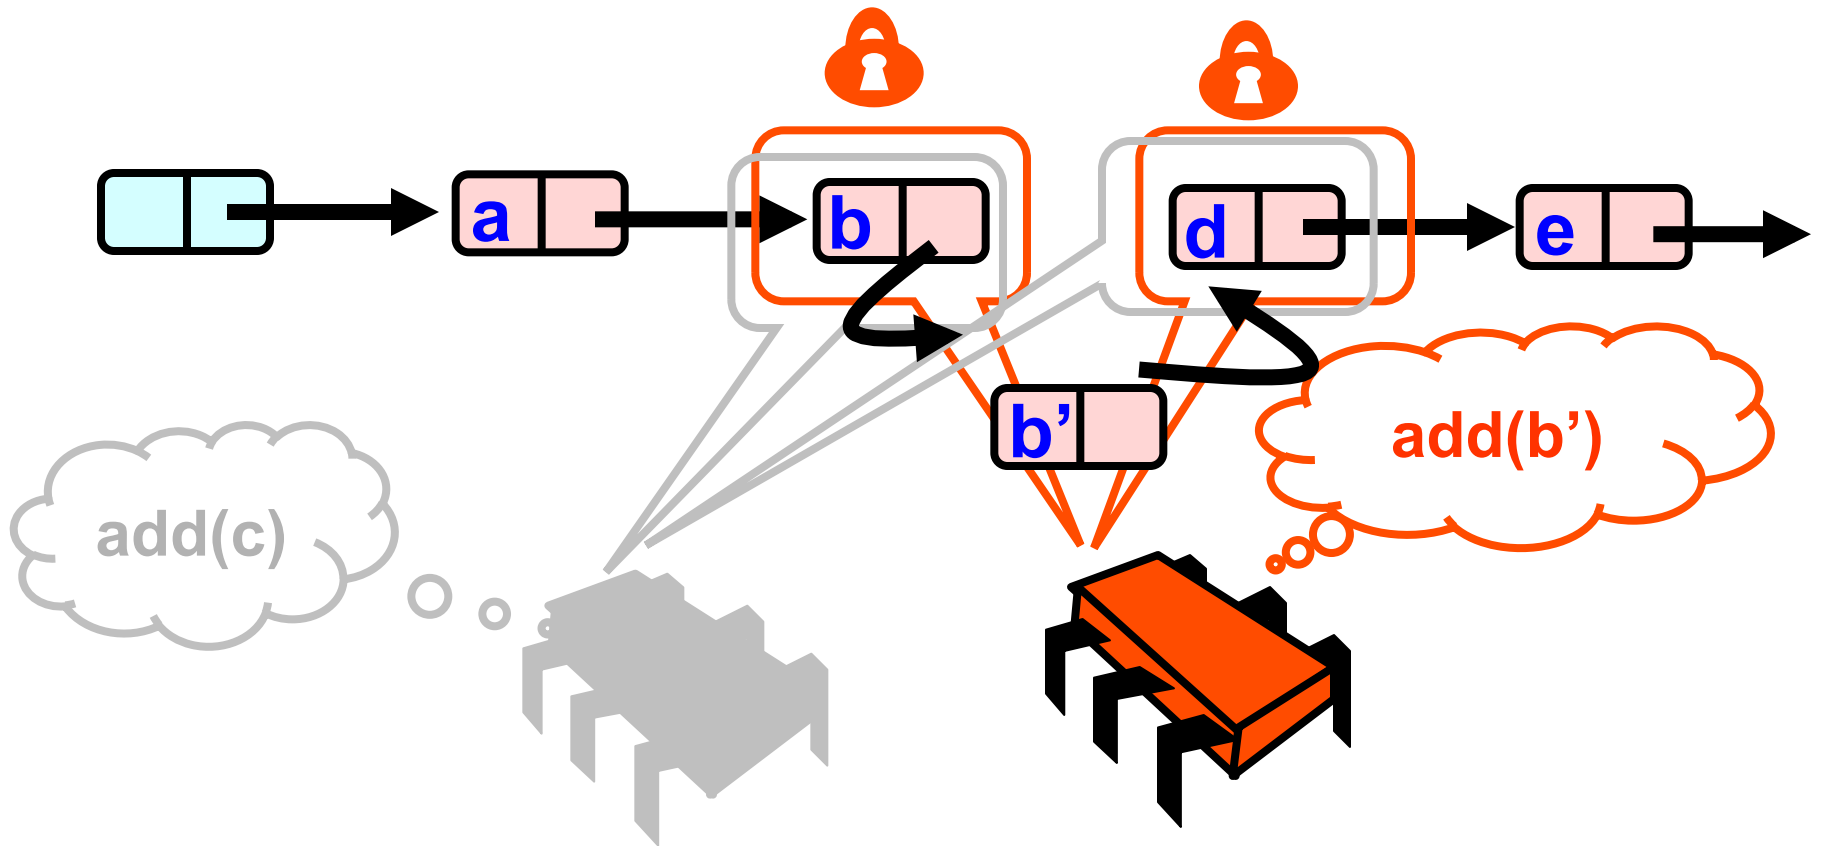
# What Else Could Go Wrong?



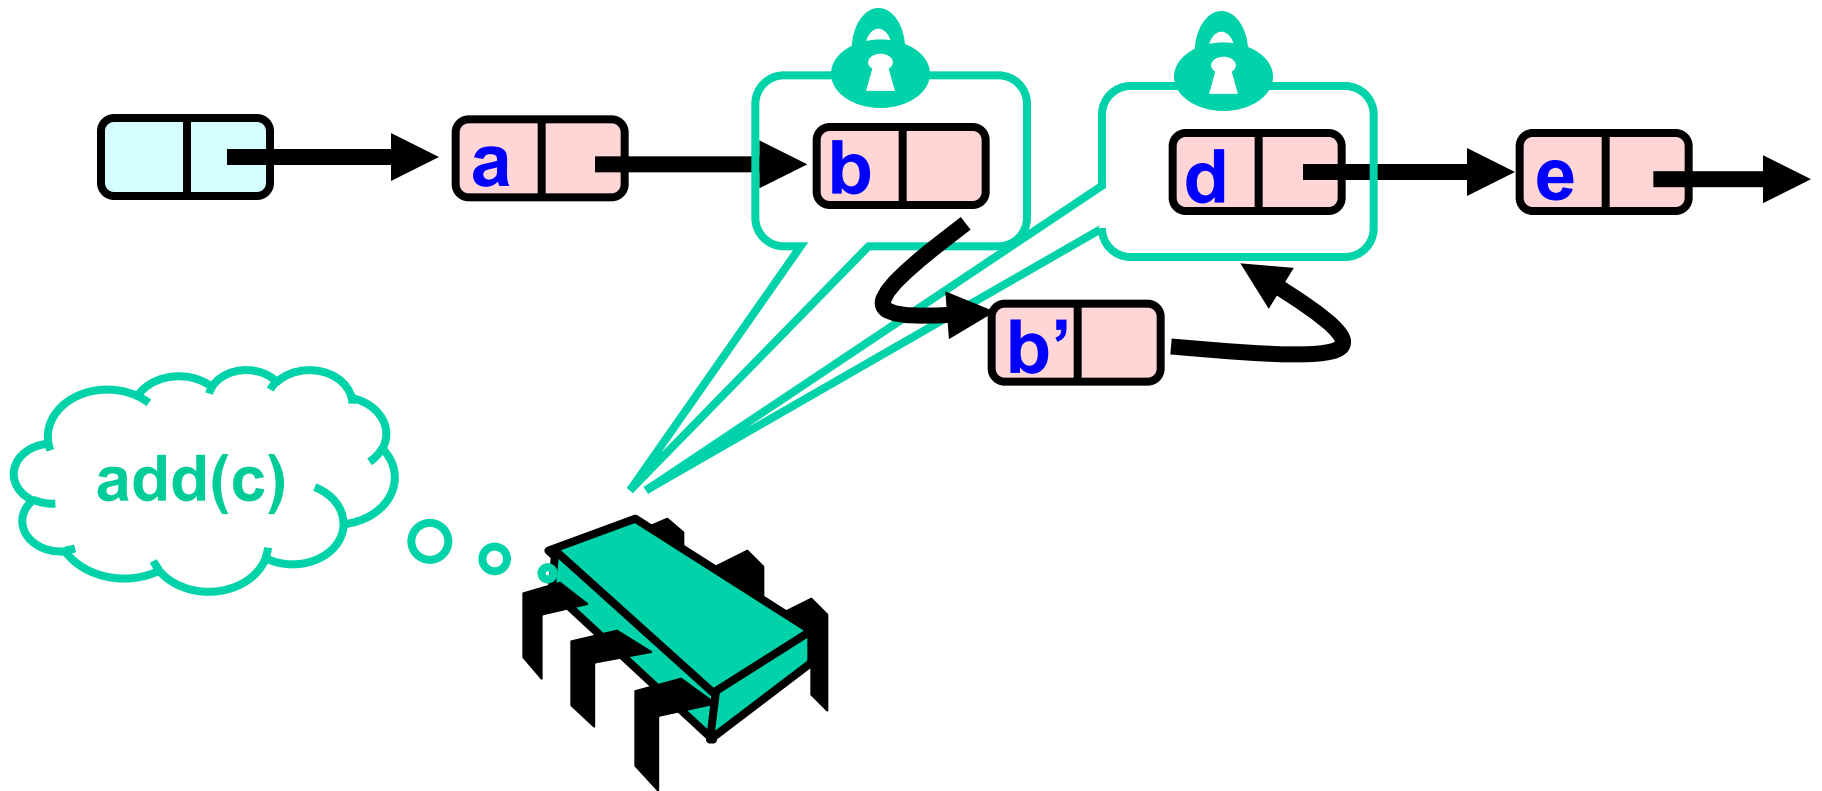
# What Else Could Go Wrong?



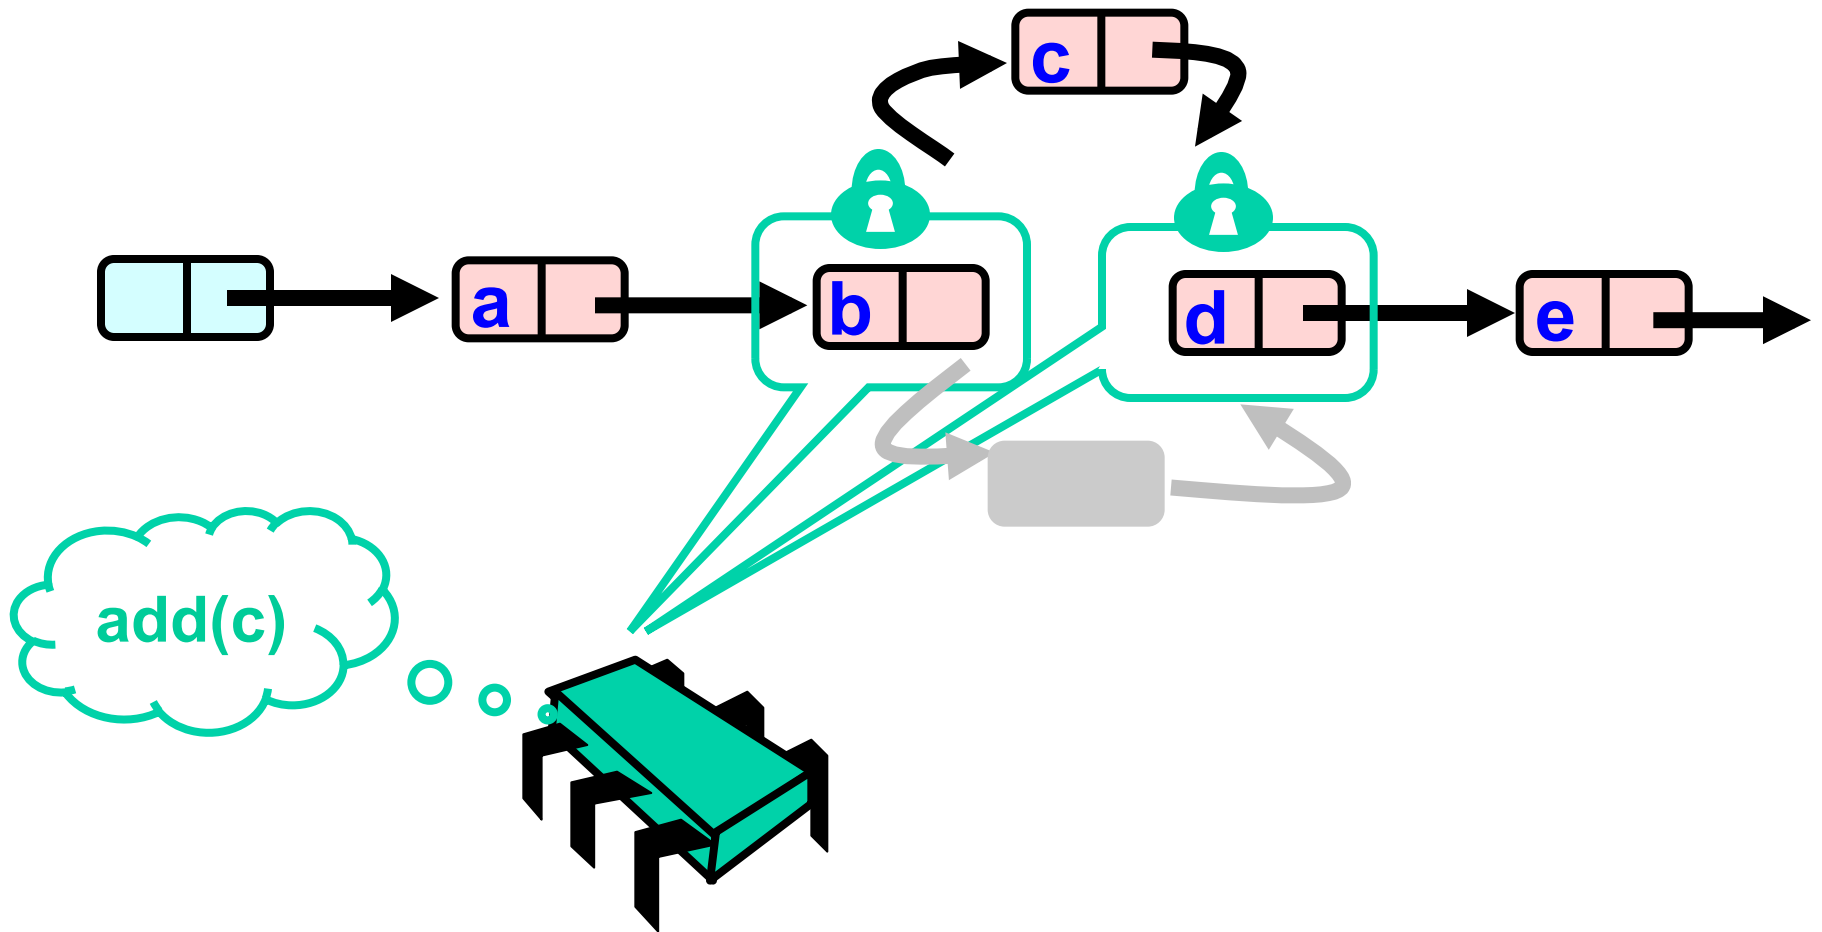
# What Else Could Go Wrong?



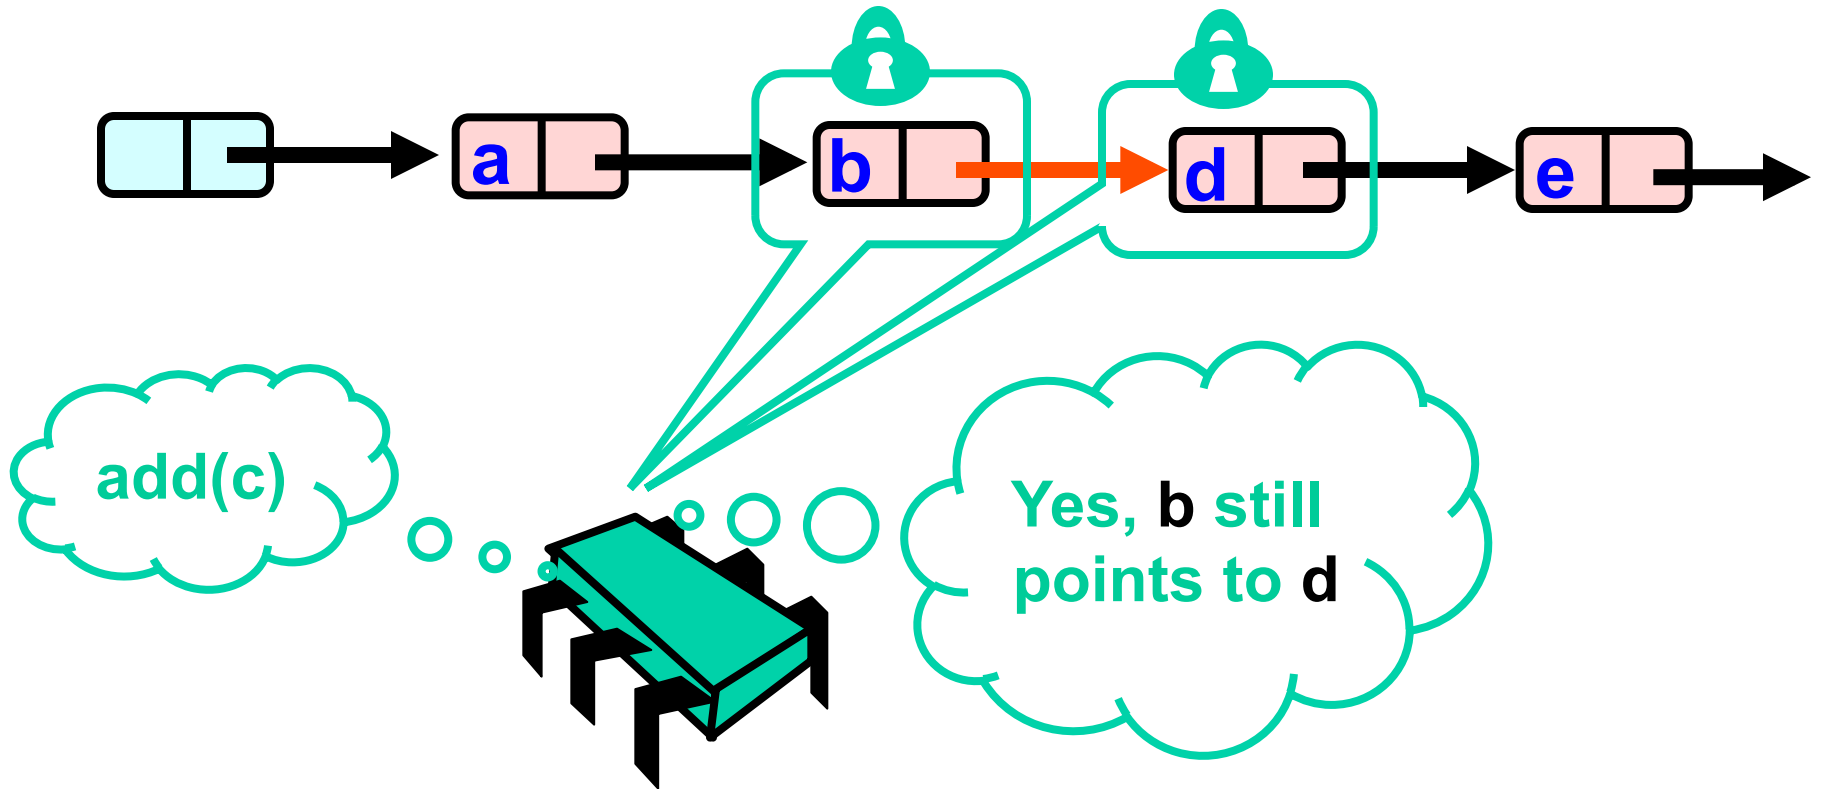
# What Else Could Go Wrong?



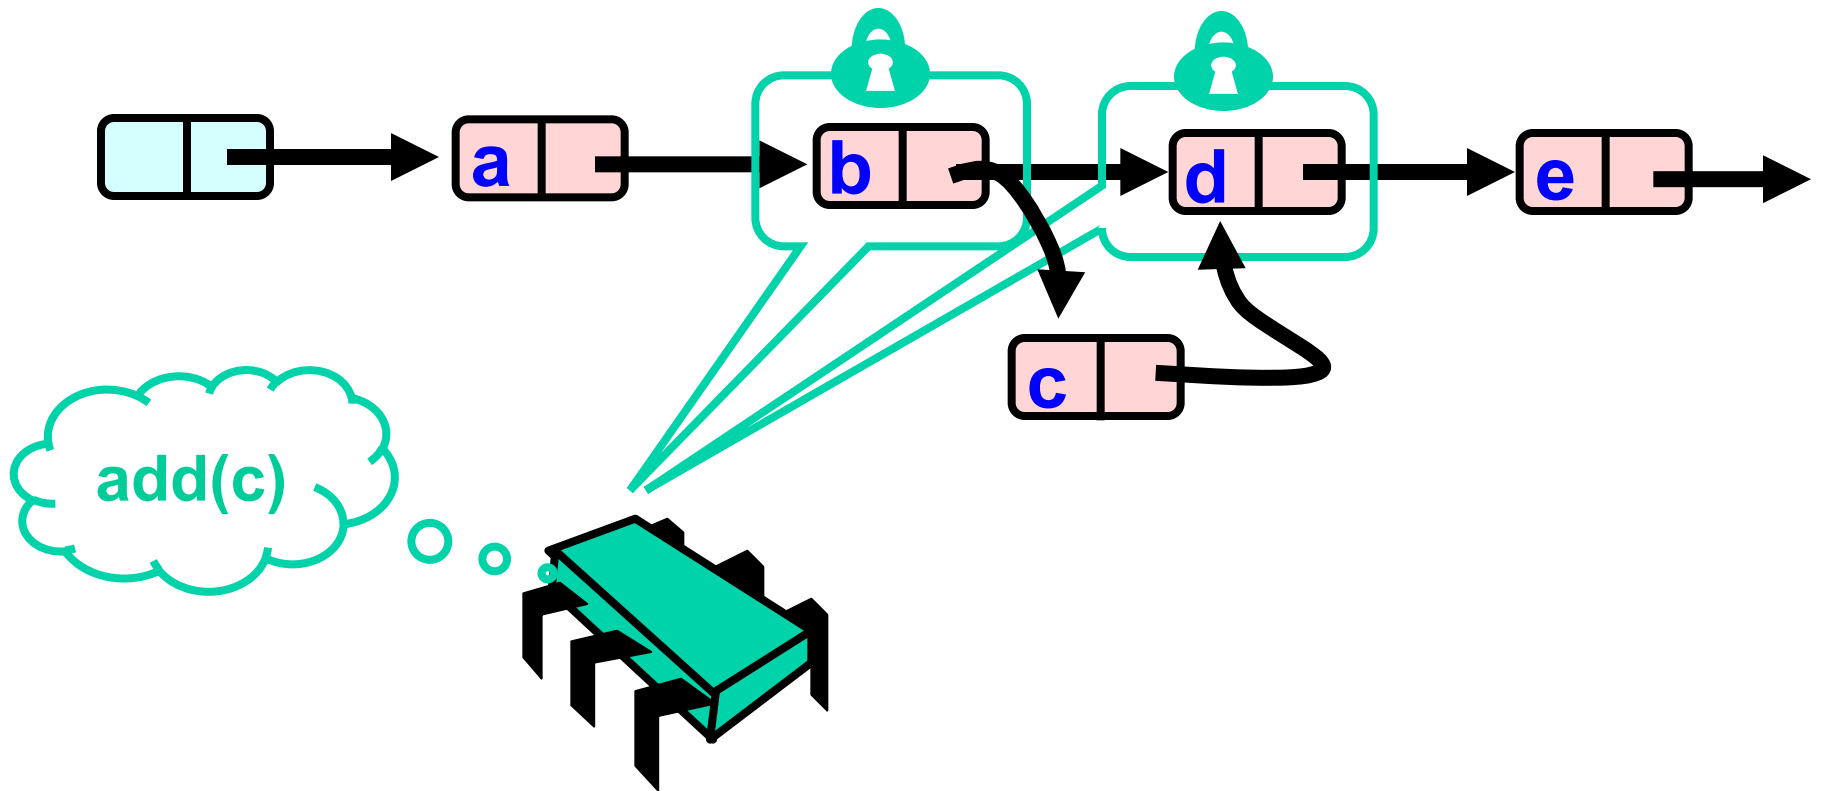
# What Else Could Go Wrong?



# Validate Part 2 (while holding locks)



# Optimistic: Linearization Point





# Same Abstraction Map

- $S(\text{head}) =$ 
  - $\{ x \mid \text{there exists } a \text{ such that}$ 
    - $a \text{ reachable from head and}$
    - $a.\text{item} = x$
  - $\}$


# Invariants

- Careful: we may traverse deleted nodes
- But we establish properties by
  - Validation
  - After we lock target nodes

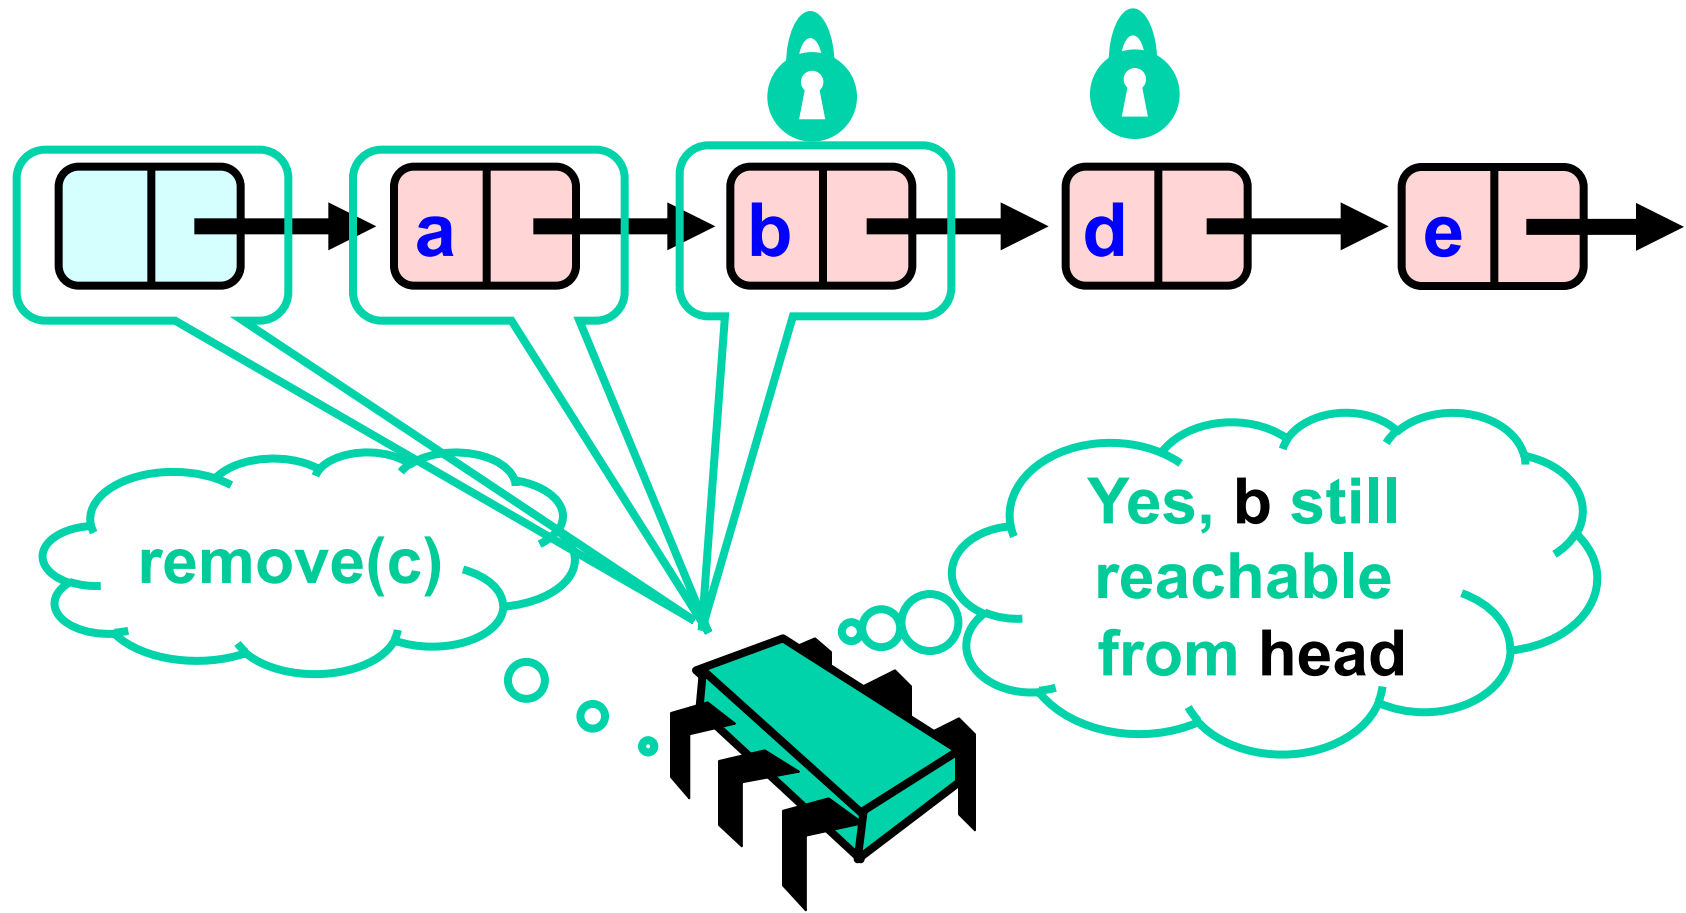
# Correctness

- If
  - Nodes b and c both locked
  - Node b still accessible
  - Node c still successor to b
- Then
  - Neither will be deleted
  - OK to delete and return true

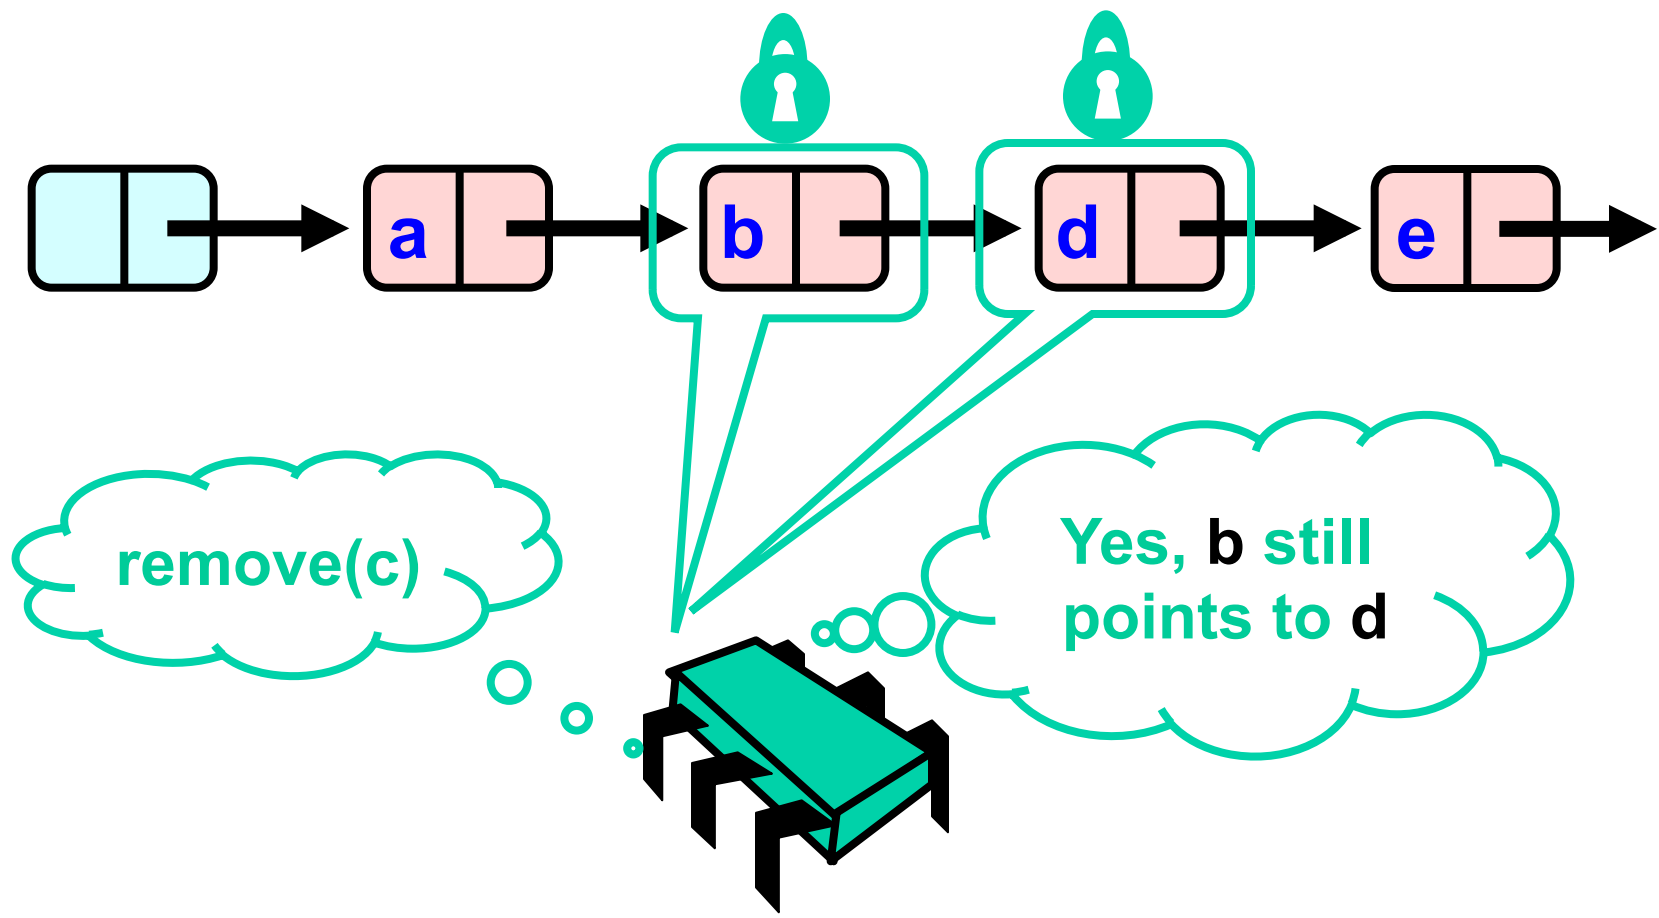
# Unsuccessful Remove



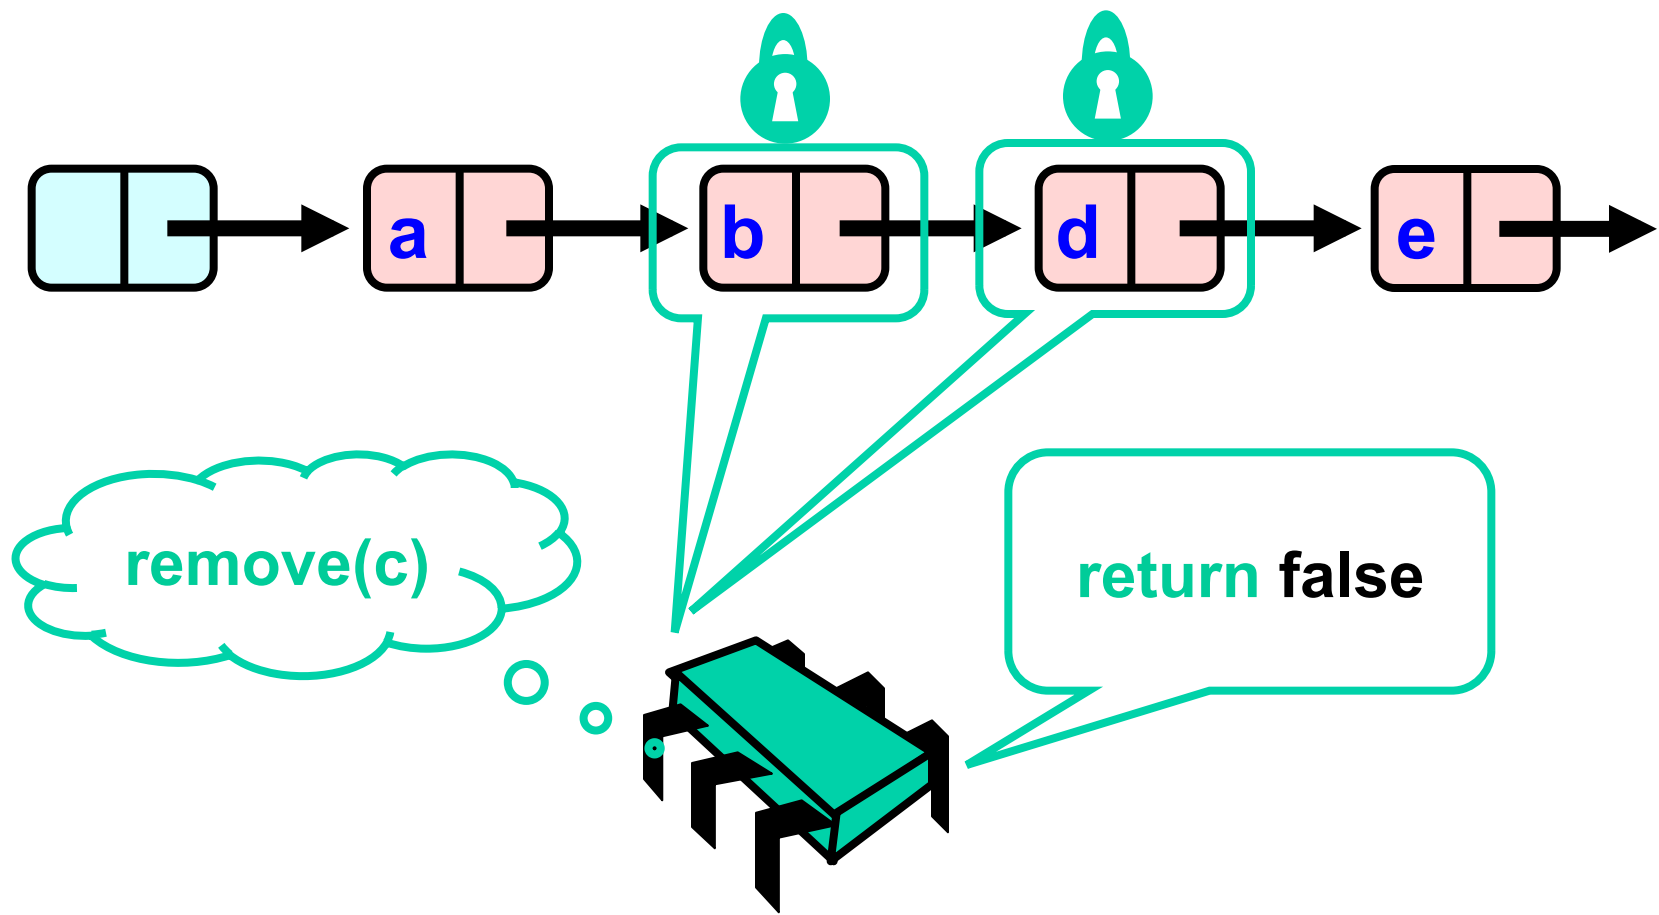
# Validate (1)



# Validate (2)



# OK Computer



# Correctness

- If
  - Nodes b and d both locked
  - Node b still accessible
  - Node d still successor to b
- Then
  - Neither will be deleted
  - No thread can add c after b
  - OK to return false



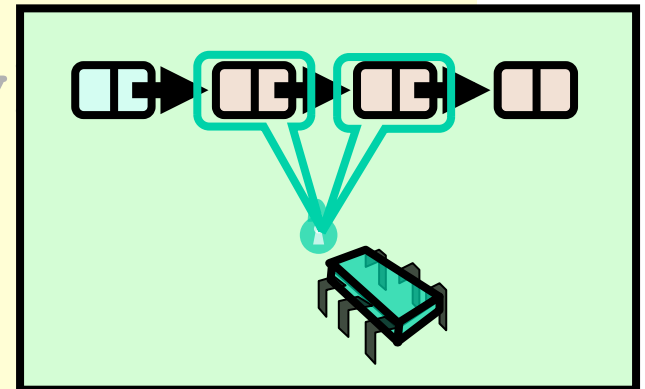
# Validation

```
private boolean
    validate(Node pred,
              Node curr) {
    Node node = head;
    while (node.key <= pred.key) {
        if (node == pred)
            return pred.next == curr;
        node = node.next;
    }
    return false;
}
```

# Validation

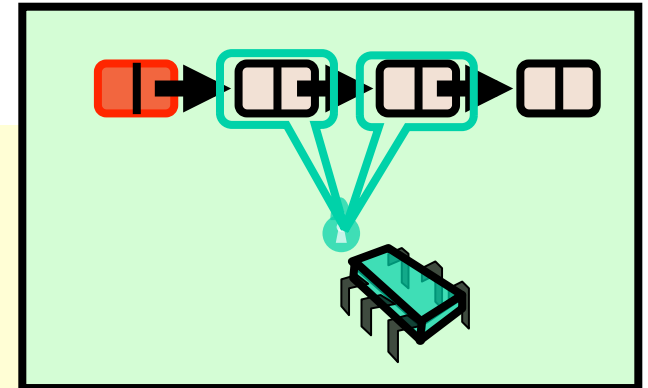
```
private boolean  
validate(Node pred,  
         Node curr) {  
    Node node = head;  
    while (node.key <= pred.key  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```

**Predecessor &  
current nodes**



# Validation

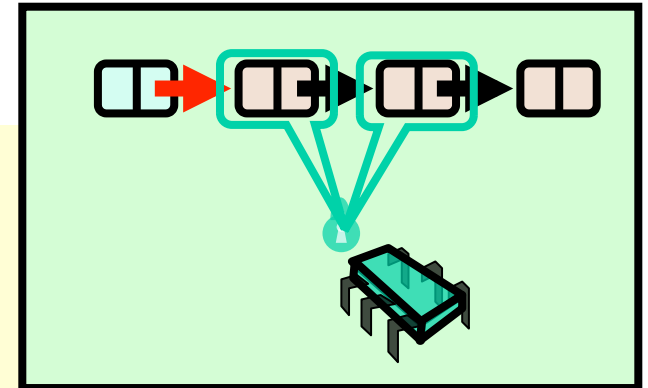
```
private boolean  
    validate(Node pred,  
              Node curr) {  
        Node node = head;  
        while (node.key <= pred.key) {  
            if (node == pred)  
                return pred.next == curr;  
            node = node.next;  
        }  
        return false;  
    }
```



**Begin at the  
beginning**

# Validation

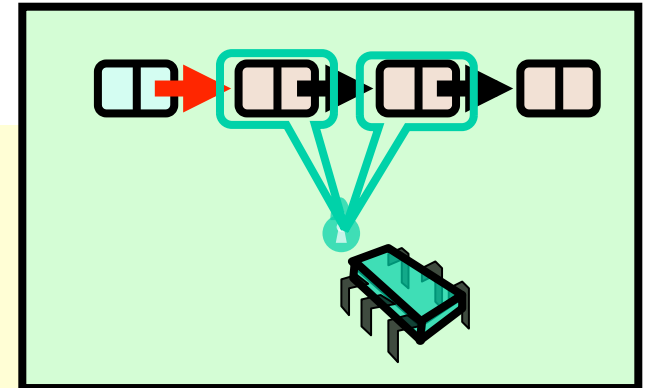
```
private boolean  
    validate(Node pred,  
              Node curr) {  
    Node node = head;  
    while (node.key <= pred.key) {  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```



**Search range of keys**

# Validation

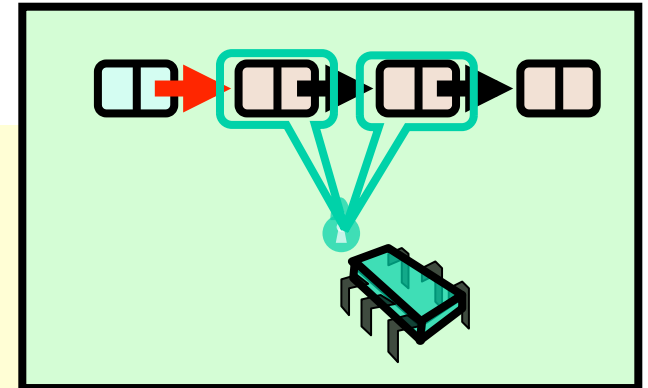
```
private boolean
  validate(Node pred,
           Node curr) {
  Node node = head;
  while (node.key <= pred.key) {
    if (node == pred)
      return pred.next == curr;
    node = node.next;
  }
  return false;
}
```



**Predecessor reachable**

# Validation

```
private boolean
validate(Node pred,
        Node curr) {
    Node node = head;
    while (node.key <= pred.key) {
        if (node == pred)
            return pred.next == curr;
        node = node.next;
    }
    return false;
}
```




**Is current node next?**

# Validation

```
private boolean  
validate(Node pred,  
         Node curr) {  
    Node node = head;  
    while (node.key <= pred.key) {  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```

Otherwise move on

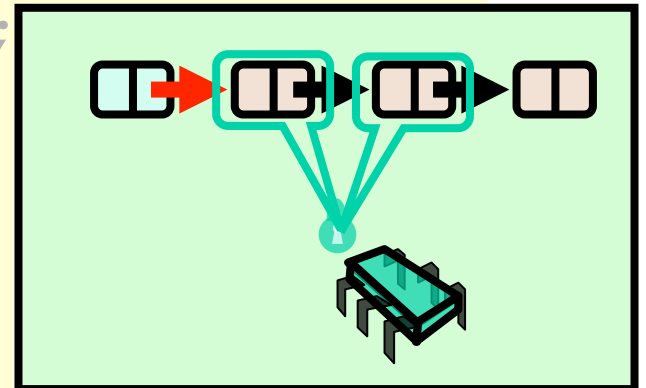
**node = node.next;**



# Validation

```
private boolean validate(Node pred,  
    Node curr) {  
    Node node = head;  
    while (node.key <= pred.key) {  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```


**Predecessor not reachable**





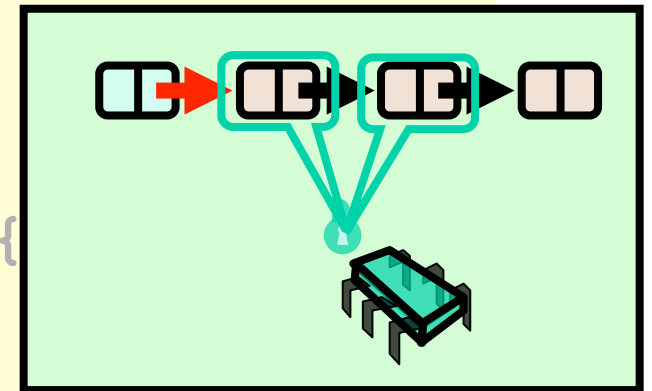
# Remove: the Traversal

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    retry: while (true) {  
        Node pred = head;  
        Node curr = pred.next;  
        while (curr.key <= key) {  
            if (item == curr.item)  
                break;  
            pred = curr;  
            curr = curr.next;  
        } ...  
    }
```



# Remove: the Traversal

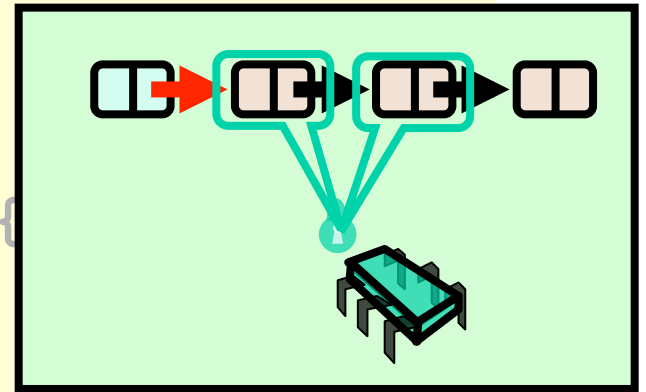
```
public boolean remove(T item) {  
    int key = item.hashCode();  
    retry: while (true) {  
        Node pred = head;  
        Node curr = pred.next;  
        while (curr.key <= key) {  
            if (item == curr.item)  
                break;  
            pred = curr;  
            curr = curr.next;  
        } ...  
    }
```



**Search key**

# Remove: the Traversal

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    retry: while (true) {  
        Node pred = head;  
        Node curr = pred.next;  
        while (curr.key <= key) {  
            if (item == curr.item)  
                break;  
            pred = curr;  
            curr = curr.next;  
        } ...  
    }
```



**Retry on synchronization conflict  
(If validation fails, we come back here.)**

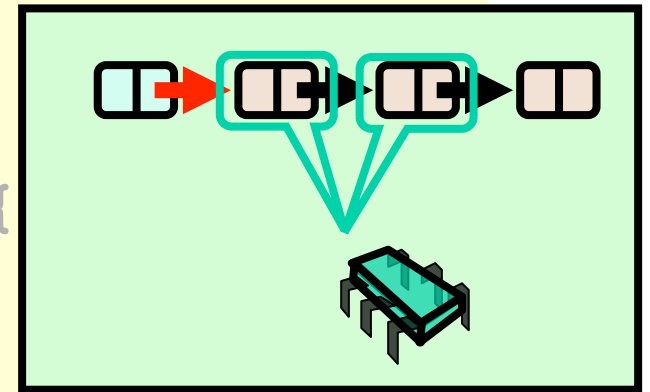
# Remove: the Traversal

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    retry: while (true) {
```

```
        Node pred = head;  
        Node curr = pred.next;
```

```
        while (curr.key <= key) {  
            if (item == curr.item)  
                break;  
            pred = curr;  
            curr = curr.next;  
        }
```

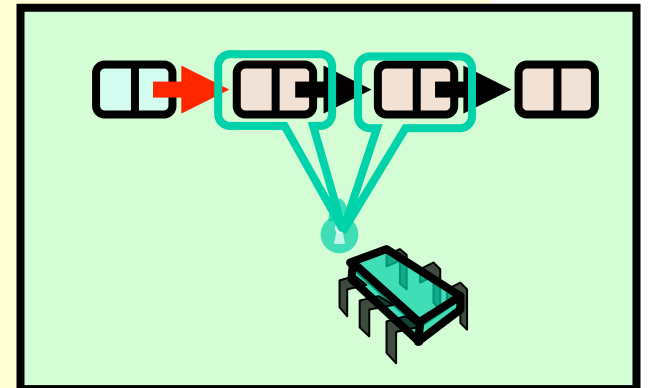
**Examine predecessor and current nodes**



# Remove: the Traversal

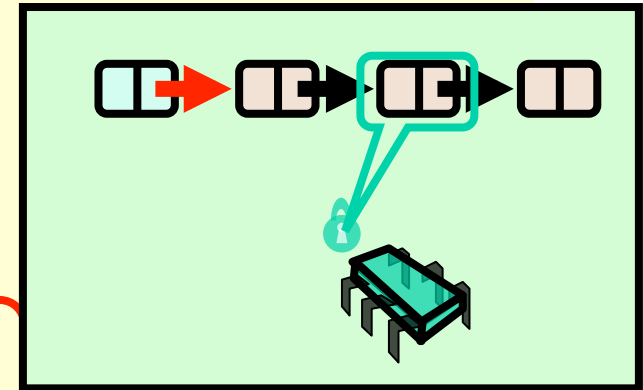
```
public boolean remove(T item) {  
    int key = item.hashCode();  
    retry: while (true) {  
        Node pred = head;  
        Node curr = pred.next;  
        while (curr.key <= key) {  
            if (item == curr.item)  
                break;  
            pred = curr;  
            curr = curr.next;  
        }  
        ...  
    }  
}
```

**Search by key**



# Remove: the Traversal

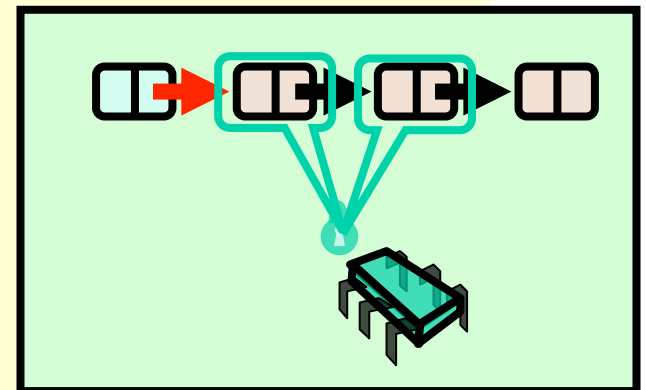
```
public boolean remove(T item) {  
    int key = item.hashCode();  
    retry: while (true) {  
        Node pred = head;  
        Node curr = pred.next;  
        while (curr.key <= key)  
            if (item == curr.item)  
                break;  
        pred = curr;  
        curr = curr.next;  
    }  
    Stop if we find item
```



# Remove: the Traversal

```
public boolean remove(T item) {  
    int key = item.hashCode();  
    retry: while (true) {  
        Node pred = head;  
        Node curr = pred.next;  
        while (curr.key <= key) {  
            if (item == curr.item)  
                break;  
            pred = curr;  
            curr = curr.next;  
        }  
        ...  
    }  
}
```

**Move along**



# On Exit from Inner Loop

- If item is present
  - curr holds item
  - pred just before curr
- If item is absent
  - curr has first higher key
  - pred just before curr
- Assuming no synchronization problems



# Remove Continue: the Deletion (After Existing Inner Loop)

```
pred.lock() ; curr.lock() ;  
try {  
    if (validate(pred,curr) {  
        if (curr.item == item) {  
            pred.next = curr.next;  
            return true;  
        } else {  
            return false;  
        }  
    } finally {  
        pred.unlock() ;  
        curr.unlock() ;  
    }  
}
```

# Remove Continue: the Deletion (After Existing Inner Loop)

```
pred.lock(); curr.lock();  
try {  
    if (validate(pred, curr) {  
        if (curr.item == item) {  
            pred.next = curr.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
} finally {  
    pred.unlock();  
    curr.unlock();  
}
```

**Always unlock**

# Remove Continue: the Deletion (After Existing Inner Loop)

```
pred.lock(); curr.lock();  
try {  
    if (validate(pred, curr) {  
        if (curr.item == item) {  
            pred.next = curr.next;  
            return true;  
        } else {  
            return false;  
        }  
    } finally {  
        pred.unlock();  
        curr.unlock();  
    }  
}
```

**Check for  
synchronization  
conflicts**

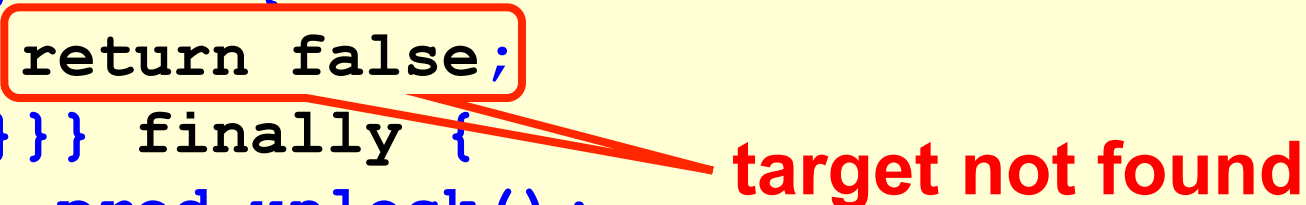
# Remove Continue: the Deletion (After Existing Inner Loop)

```
pred.lock(); curr.lock();  
try {  
    if (validate(pred, curr) {  
        if (curr.item == item) {  
            pred.next = curr.next;  
            return true;  
        } else {  
            return false;  
        }  
    } finally {  
        pred.unlock();  
        curr.unlock();  
    }  
}
```

**target found,  
remove node**

# Remove Continue: the Deletion (After Existing Inner Loop)

```
pred.lock(); curr.lock();  
try {  
    if (validate(pred, curr) {  
        if (curr.item == item) {  
            pred.next = curr.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
} finally {  
    pred.unlock();  
    curr.unlock();  
}
```



**target not found**

# Optimistic List

- Limited hot-spots
  - Holding locks only on the targets of add(), remove(), contains()
  - No contention on traversals
  - Traversals are "wait-free"  
(What's wait free?)

# Progress Conditions

- ***Deadlock-free:*** some thread trying to acquire the lock eventually succeeds.
- ***Starvation-free:*** every thread trying to acquire the lock eventually succeeds.
- ***Lock-free:*** some thread calling a method eventually returns.
- ***Wait-free:*** every thread calling a method eventually returns.

# Progress Conditions

	Non-Blocking	Blocking
Everyone makes progress	Wait-free	Starvation-free
Someone makes progress	Lock-free	Deadlock-free



# So Far, So Good

- Much less lock acquisition/release
  - Performance
  - Concurrency
- Problems
  - Need to traverse list twice
  - **contains ()** method acquires locks

# Evaluation

- Optimistic is effective if
  - cost of scanning twice without locks is less than
  - cost of scanning once with locks
- Drawback
  - **contains ()** acquires locks
  - 90% of calls in many apps

# Lazy List

- Like optimistic, except
  - Scan once
  - **contains (x)** never locks ...
- Key insight
  - Removing nodes causes trouble
  - Do it “lazily”


# Lazy List

- **remove ( )**
  - Scans list (as before)
  - Locks predecessor & current (as before)
- Logical delete
  - Marks current node as removed (new!)
- Physical delete
  - Redirects predecessor's next (as before)

# Lazy Removal

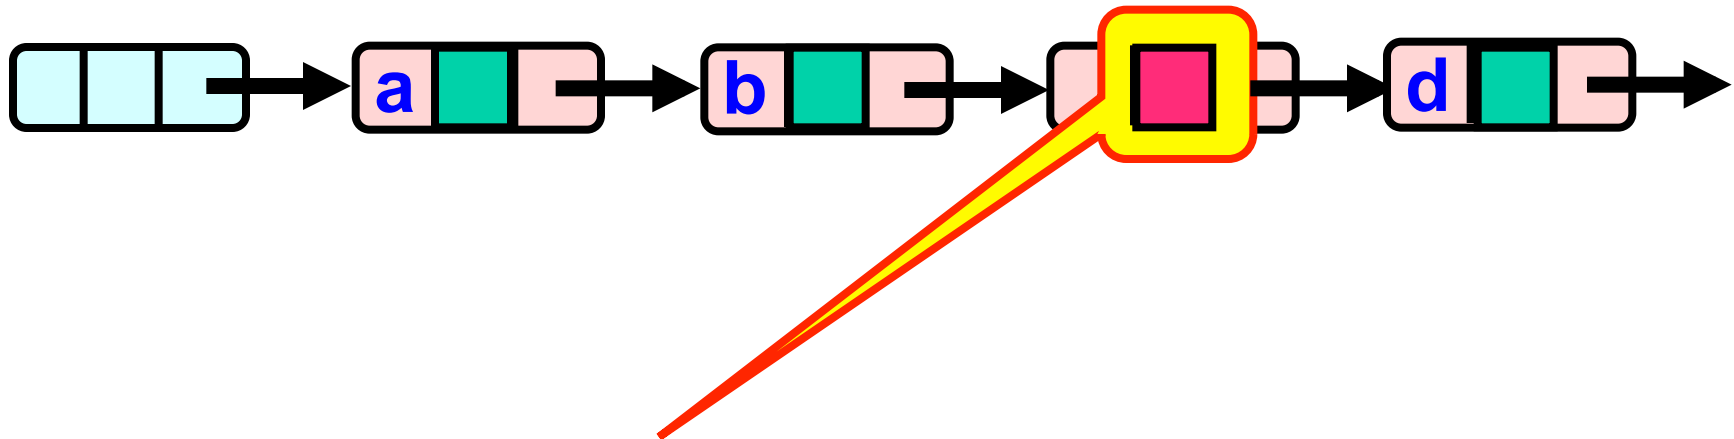


# Lazy Removal



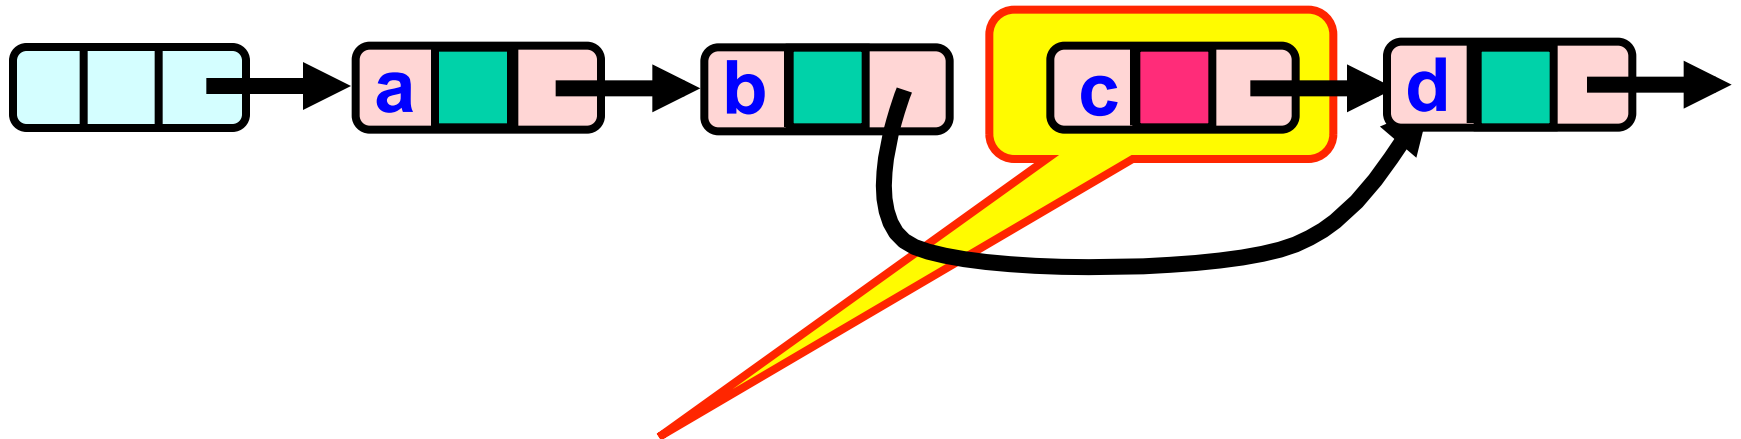
Present in list

# Lazy Removal



Logically deleted

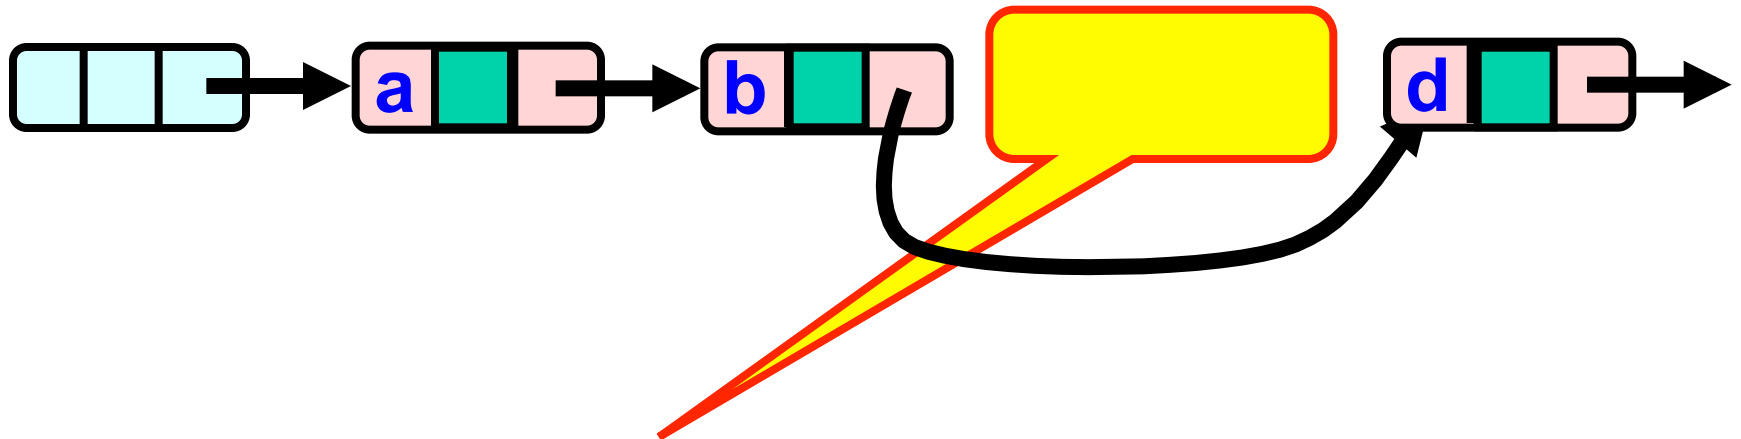
# Lazy Removal



Physically deleted



# Lazy Removal



Physically deleted

# Lazy List

- All Methods
  - Scan through locked and marked nodes
  - add and remove still locks pred and curr, but not contain
  - Adding / removing a node doesn't slow down contain() ...

# Lazy List Validation

- No need to rescan list!
- Check that pred is not marked
- Check that curr is not marked
- Check that pred points to curr

# New Abstraction Map

- $S(\text{head}) =$ 
  - $\{ x \mid \text{there exists node } a \text{ such that}$ 
    - $a$  reachable from head and
    - $a.\text{item} = x$  and
    - $a$  is unmarked
  - $\}$

# Invariant

- If an item is not marked, it is reachable from head and still in the set.
- Any unmarked reachable node remains reachable even if its predecessor is logically or physically removed

# Validation

```
private boolean  
    validate(Node pred, Node curr) {  
return  
    !pred.marked &&  
    !curr.marked &&  
    pred.next == curr) ;  
}
```

# List Validate Method

```
private boolean  
    validate(Node pred, Node curr) {  
    return  
        !pred.marked &&  
        !curr.marked &&  
        pred.next == curr);  
}
```

**Predecessor not  
Logically removed**

# List Validate Method

```
private boolean  
    validate(Node pred, Node curr) {  
    return  
        !pred.marked &&  
        !curr.marked &&  
        pred.next == curr);  
}
```



**Current not  
Logically removed**



# List Validate Method

```
private boolean  
    validate(Node pred, Node curr) {  
    return  
        !pred.marked &&  
        !curr.marked &&  
        pred.next == curr) ;  
}
```

**Predecessor still  
Points to current**

# Validation

```
private boolean  
    validate(Node pred, Node curr) {  
return  
    !pred.marked &&  
    !curr.marked &&  
    pred.next == curr) ;  
}
```

**Both the next and marked fields need  
to be volatile!**

# Remove: the Deletion

```
... // the traversal
pred.lock(); curr.lock();
try {
    if (validate(pred, curr) {
        if (curr.key == key) {
            curr.marked = true;
            pred.next = curr.next;
            return true;
        } else {
            return false;
        }
    } finally {
        pred.unlock();
        curr.unlock();
    }
}
```

# Remove: the Deletion

```
... // the traversal
pred.lock(); curr.lock();
try {
    if (validate(pred, curr) {
        if (curr.key == key) {
            curr.marked = true;
            pred.next = curr.next;
            return true;
        } else {
            return false;
        }
    } finally {
        pred.unlock();
        curr.unlock();
    }
}
```

**Validate as before**

# Remove: the Deletion

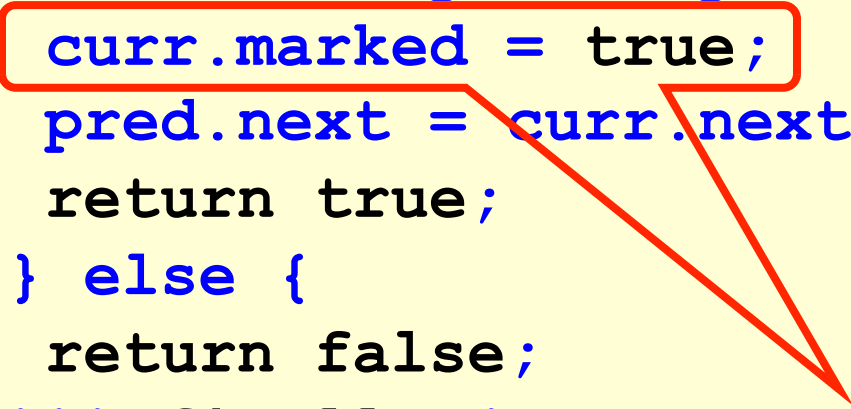
```
... // the traversal
pred.lock(); curr.lock();
try {
    if (validate(pred, curr) {
        if (curr.key == key) {
            curr.marked = true;
            pred.next = curr.next;
            return true;
        } else {
            return false;
        }
    } finally {
        pred.unlock();
        curr.unlock();
    }
}
```



**Key found**

# Remove: the Deletion

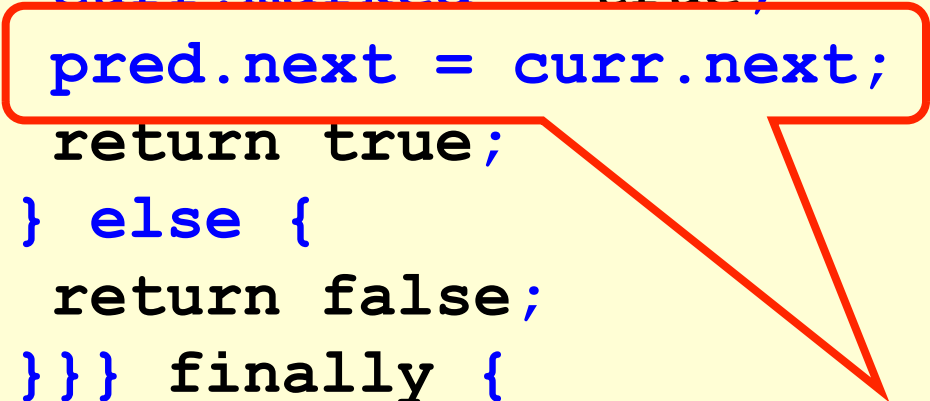
```
... // the traversal
pred.lock(); curr.lock();
try {
    if (validate(pred, curr) {
        if (curr.key == key) {
            curr.marked = true;
            pred.next = curr.next;
            return true;
        } else {
            return false;
        }
    } finally {
        pred.unlock();
        curr.unlock();
    }
}
```



**Logical remove**

# Remove: the Deletion

```
... // the traversal
pred.lock(); curr.lock();
try {
    if (validate(pred, curr) {
        if (curr.key == key) {
            curr.marked = true;
            pred.next = curr.next;
            return true;
        } else {
            return false;
        }
    } finally {
        pred.unlock();
        curr.unlock();
    }
}
```



**physical remove**

# Contains

```
public boolean contains(T item) {  
    int key = item.hashCode();  
    Node curr = head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```



# Contains

```
public boolean contains(T item) {  
    int key = item.hashCode();  
    Node curr = head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```



**Start at the head**

# Contains

```
public boolean contains(T item) {  
    int key = item.hashCode();  
    Node curr = head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```

**Search key range**

# Contains

```
public boolean contains(T item) {  
    int key = item.hashCode();  
    Node curr = head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```

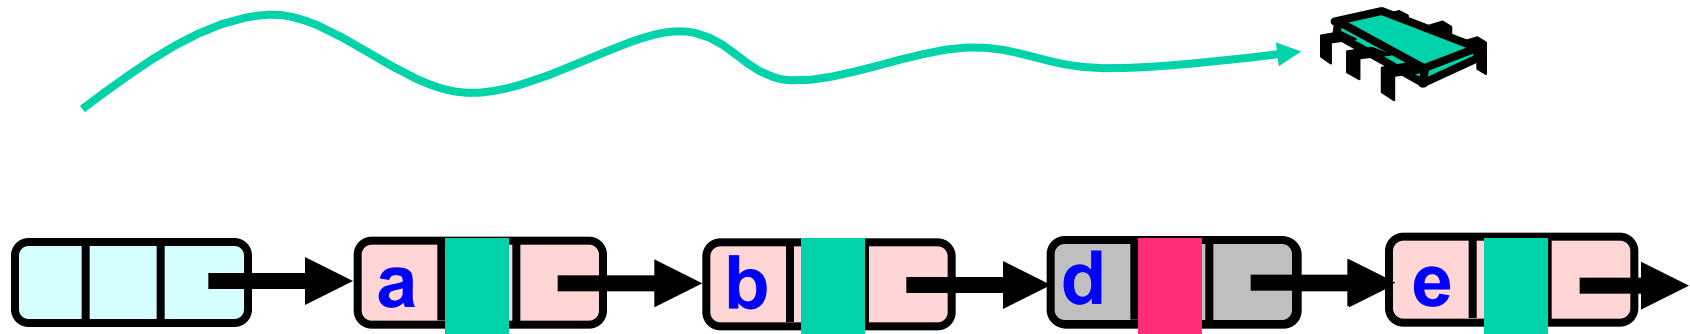
**Traverse *without locking***  
**(nodes may have been removed)**

# Contains

```
public boolean contains(T item) {  
    int key = item.hashCode();  
    Node curr = head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```

**Present and undeleted?**

# Summary: Wait-free Contains



Use Mark bit + list ordering

1. Not marked  $\rightarrow$  in the set
2. Marked or missing  $\rightarrow$  not in the set
3. Traverse the list only once!

# Evaluation

- Good:
  - **contains ()** doesn't lock
  - In fact, its wait-free!
  - Good because typically high % contains()
  - Uncontended calls to add and remove don't re-traverse
- Bad
  - Contended **add ()** and **remove ()** calls must re-traverse
  - Traffic jam if one thread delays

# Traffic Jam

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
  - Enters critical section
  - And “eats the big muffin”
    - Cache miss, page fault, descheduled ...
  - Everyone else using that lock is stuck!
  - Need to trust the scheduler....

# Lock-Free Data Structures

- No matter what ...
  - Guarantees minimal progress in any execution
  - i.e. Some thread will always complete a method call
  - Even if others halt at malicious times
  - Implies that implementation can't use locks





# Lock-free Lists

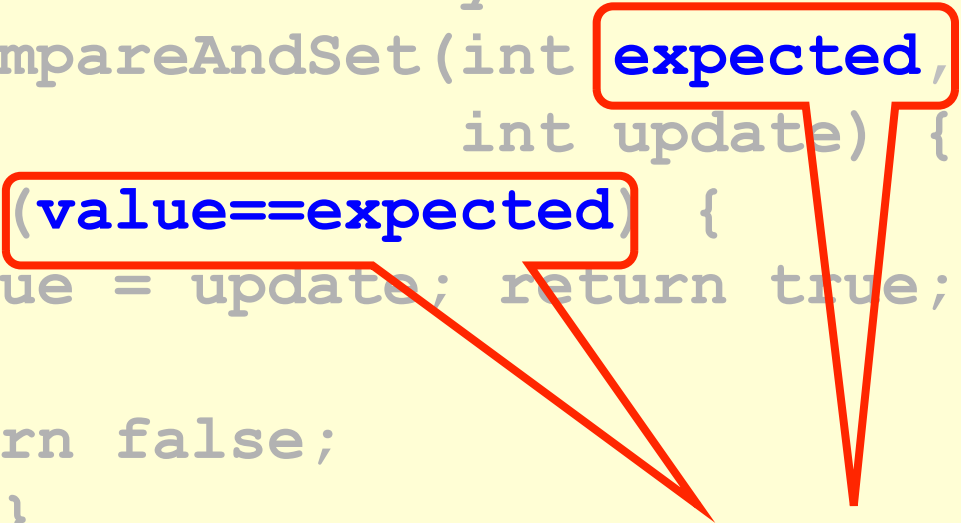
- Next logical step
  - Wait-free **contains()**
  - lock-free **add()** and **remove()**
- Use only **compareAndSet()**
  - What could go wrong?

# compareAndSet

```
public abstract class CASObject {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

# compareAndSet

```
public abstract class CASObject {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
}
```



If value is as expected, ...

# compareAndSet

```
public abstract class CASOBJECT{  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                        int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
    ... replace it
```

# compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                      int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  

```

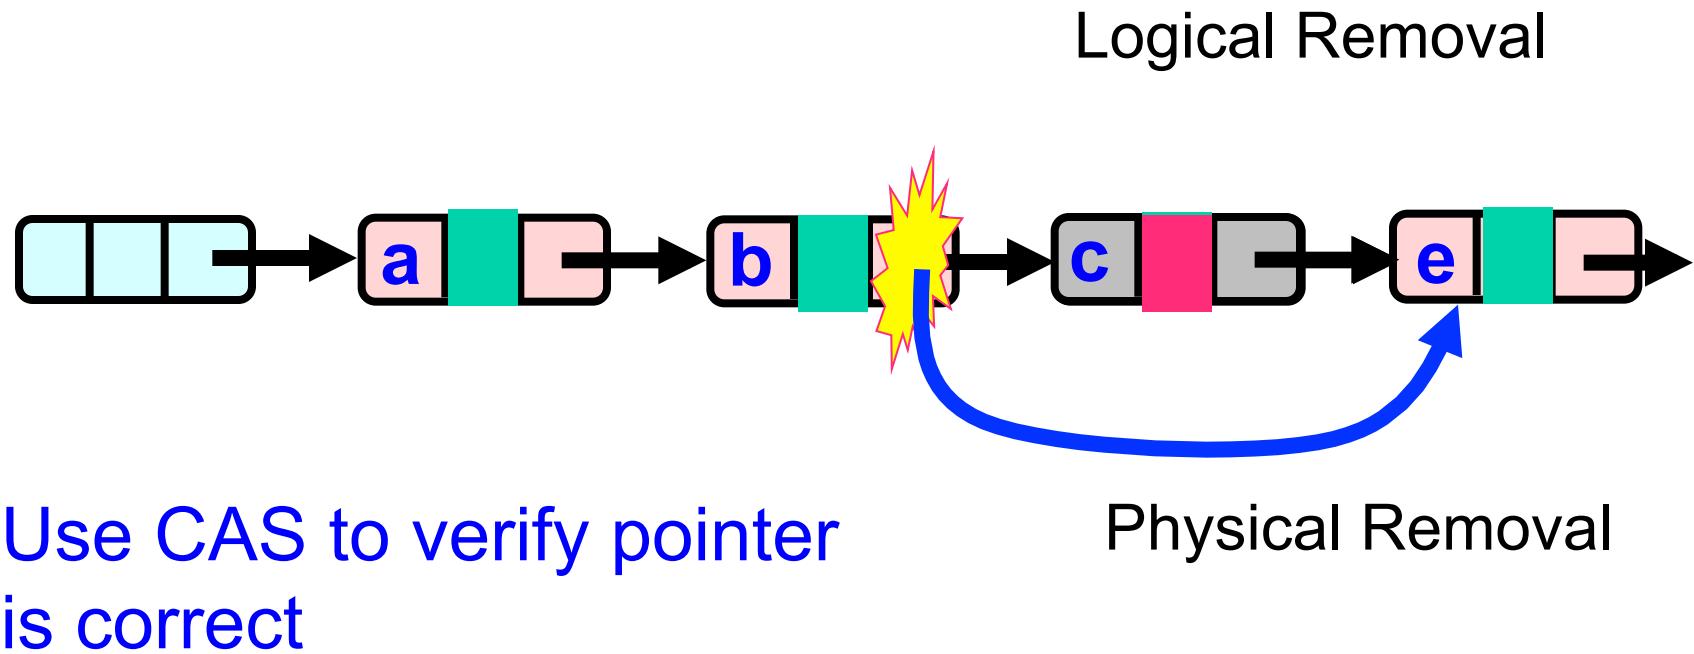
**Report success**

# compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    }  
    ...  
}
```

Otherwise report failure


# Lock-free Lists



Not enough!

# Problem...

T1: Logical Removal



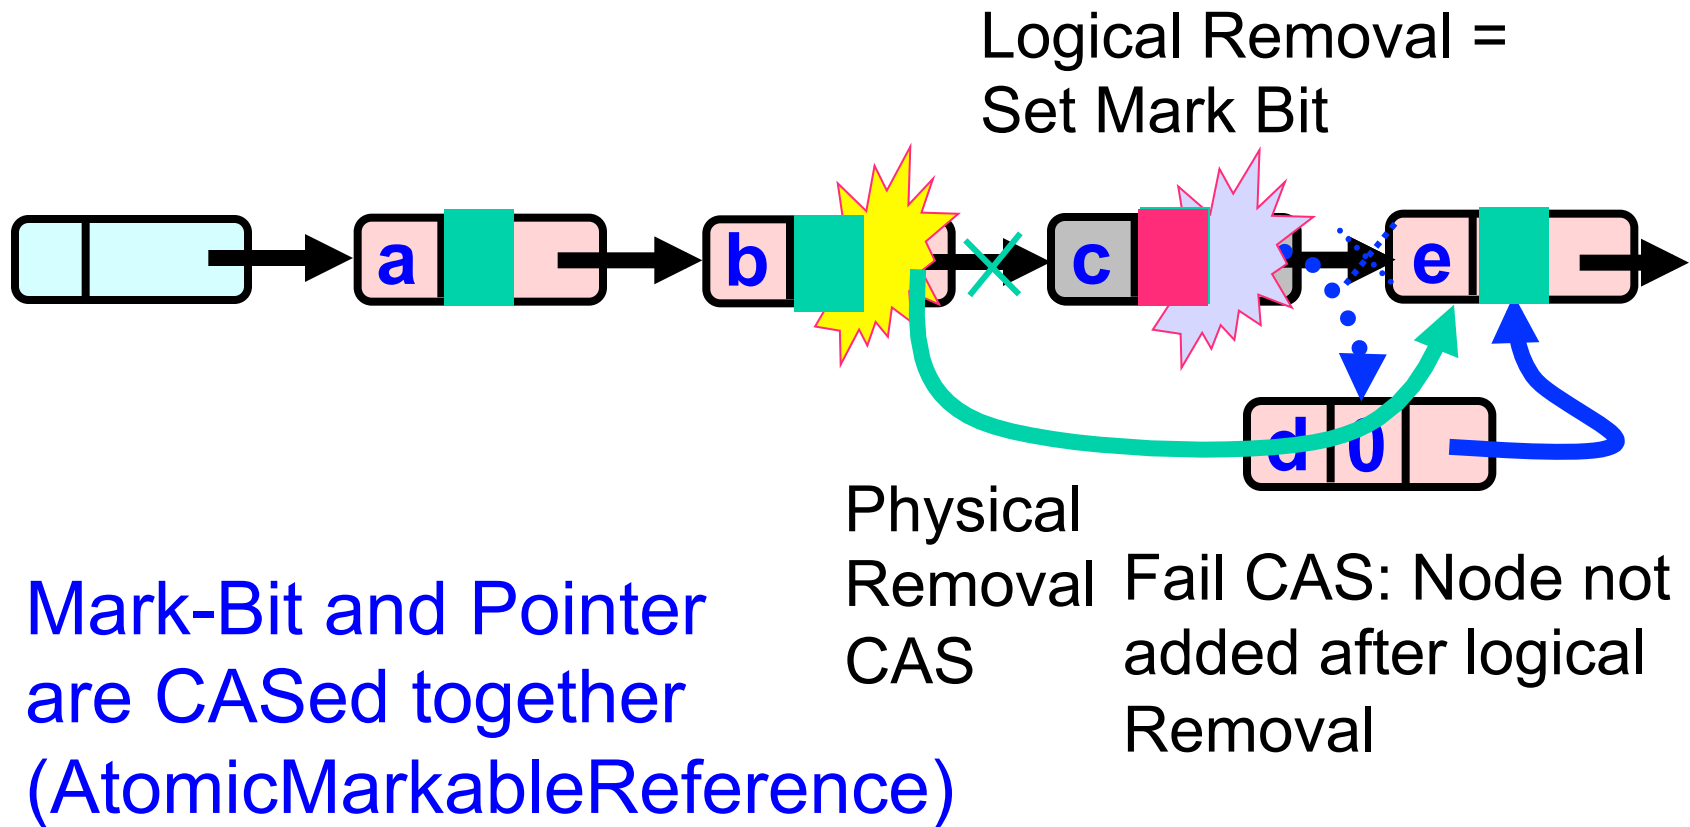
T1: Physical  
Removal

T2: Node added

**Lost Update!**

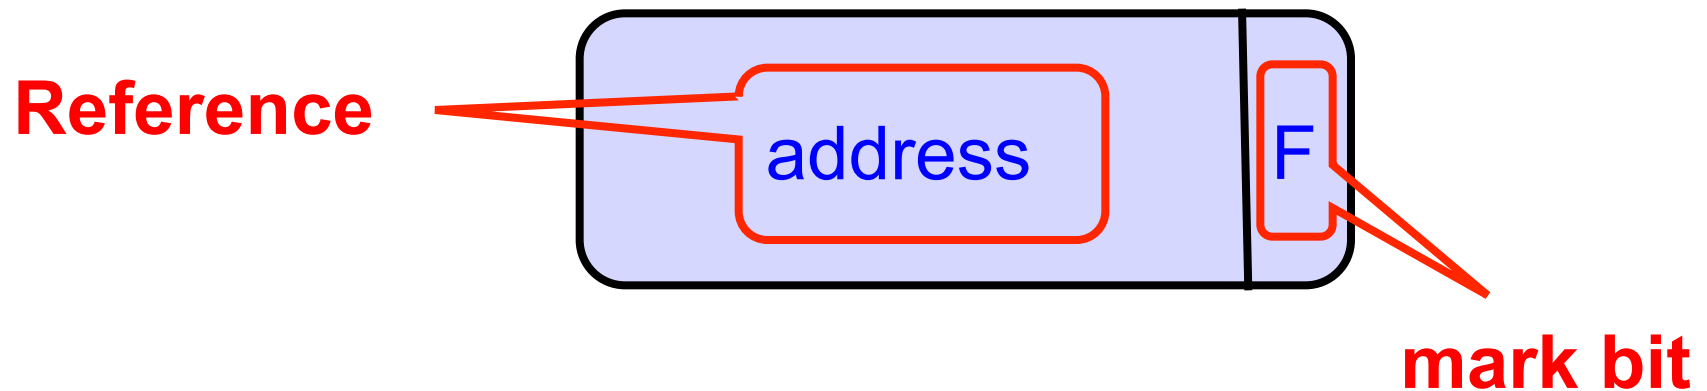


# The Solution: Combine Bit and Pointer



# Marking a Node

- **AtomicMarkableReference** class
  - `Java.util.concurrent.atomic` package



# Extracting Reference & Mark

```
public Object get(boolean[] marked) ;
```

# Extracting Reference & Mark

```
public Object get(boolean[] marked) ;
```

**Returns  
reference**

**Returns mark at  
array index 0!**

# Extracting Mark Only

```
public boolean isMarked();
```



**Value of  
mark**

# Changing State

```
public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark) ;
```

# Changing State

If this is the current  
reference ...

```
public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark);
```

And this is the  
current mark ...

# Changing State

...then change to this  
new reference ...

```
public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark) ;
```

... and this new  
mark



# Changing State

```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark) ;
```

# Changing State

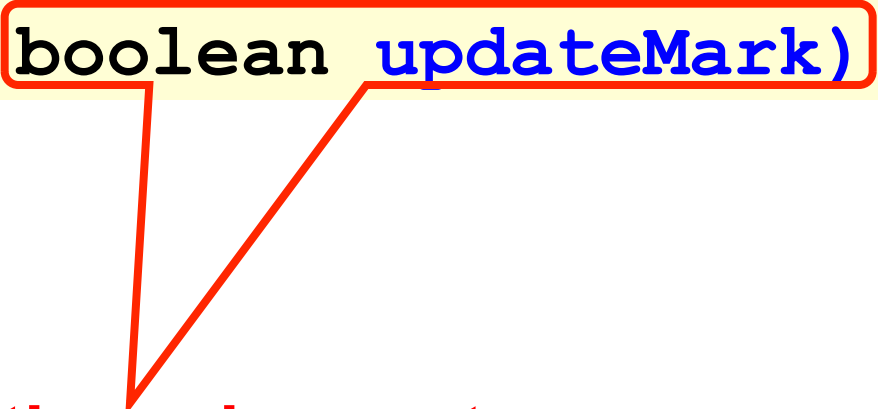
```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark) ;
```



**If this is the current  
reference ...**

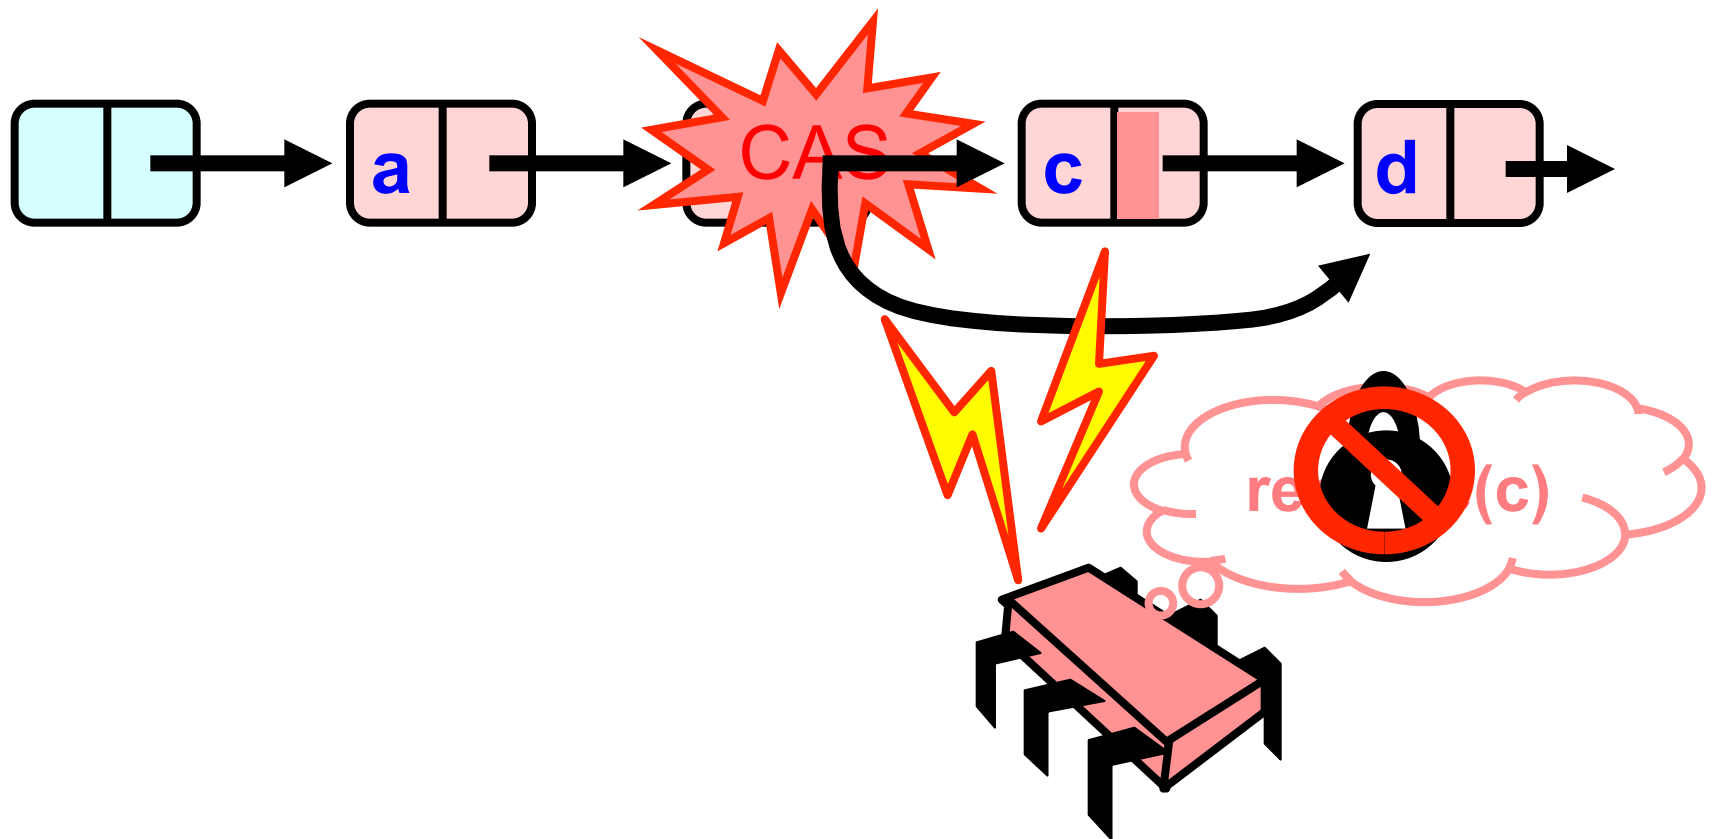
# Changing State

```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark);
```

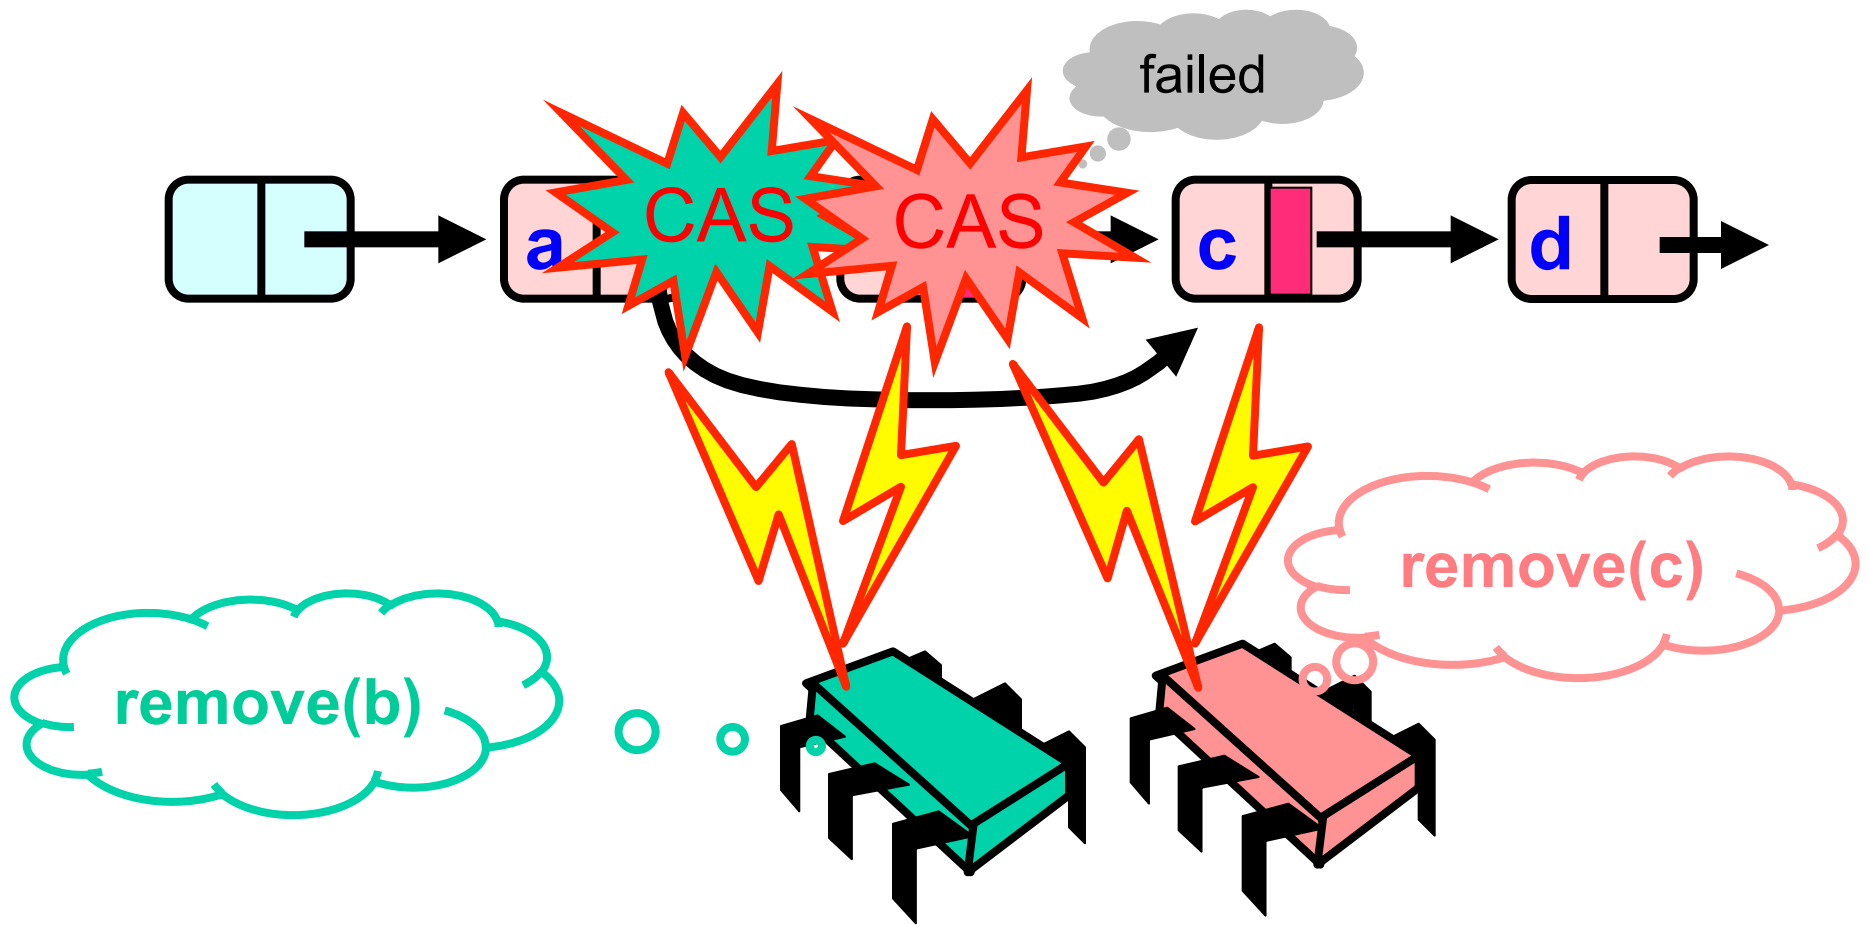


**.. then change to  
this new mark.**

# Removing a Node



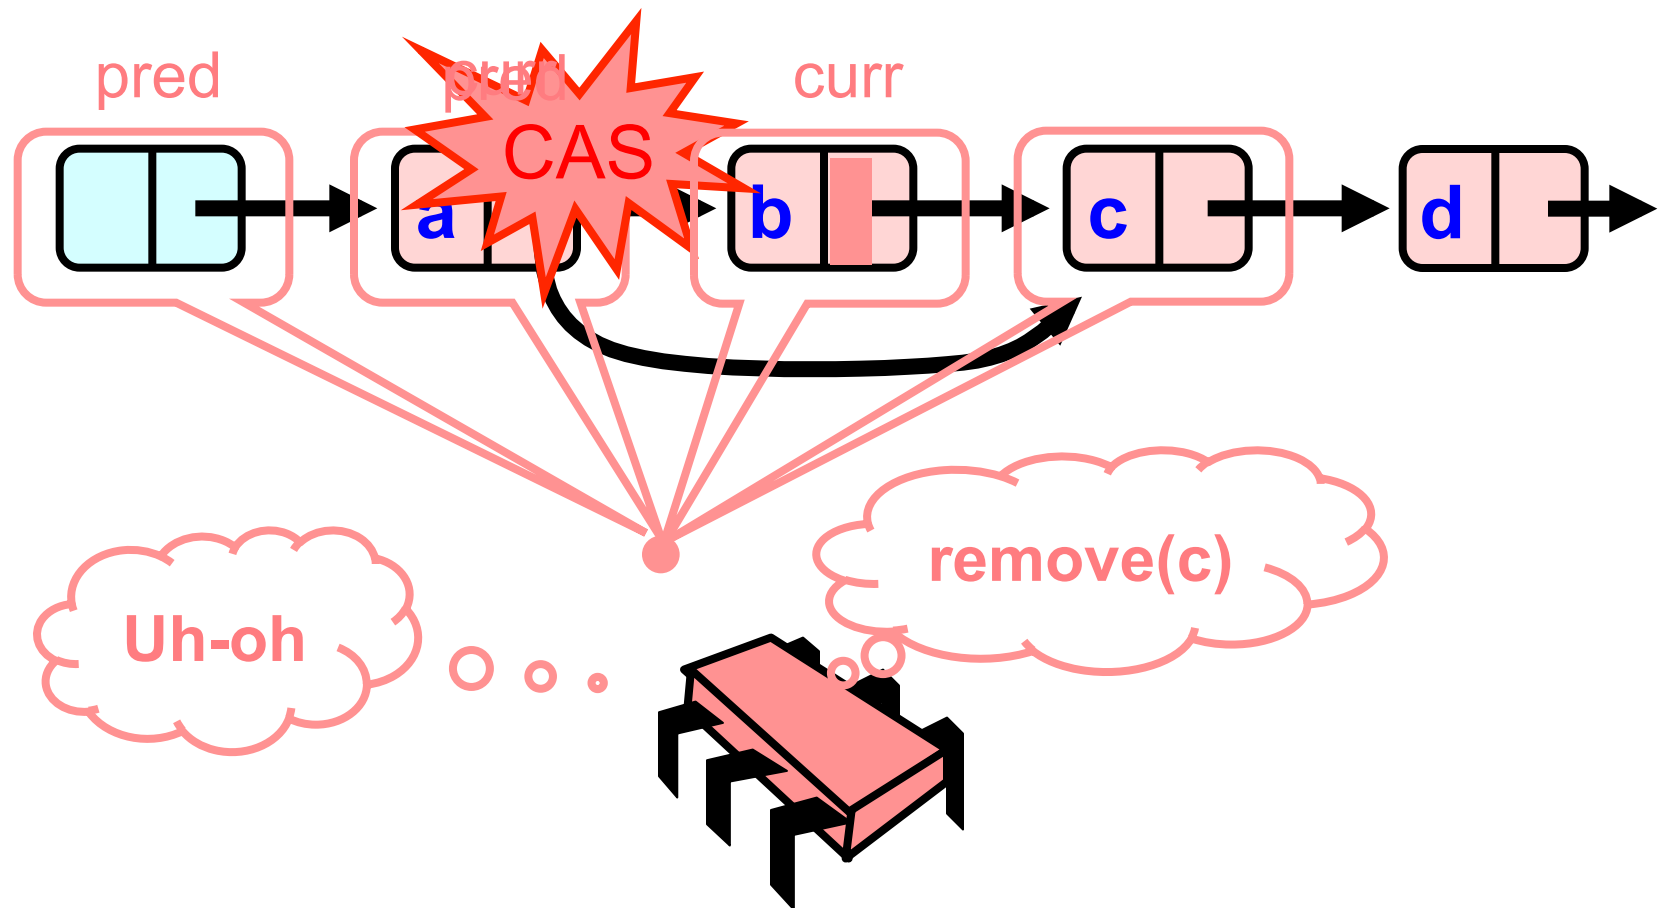
# Removing a Node



# Traversing the List

- Q: what do you do when you find a “logically” deleted node in your path?
- A: finish the job.
  - CAS the predecessor’s next field
  - Proceed (repeat as needed)

# Lock-Free Traversal (only Add and Remove)



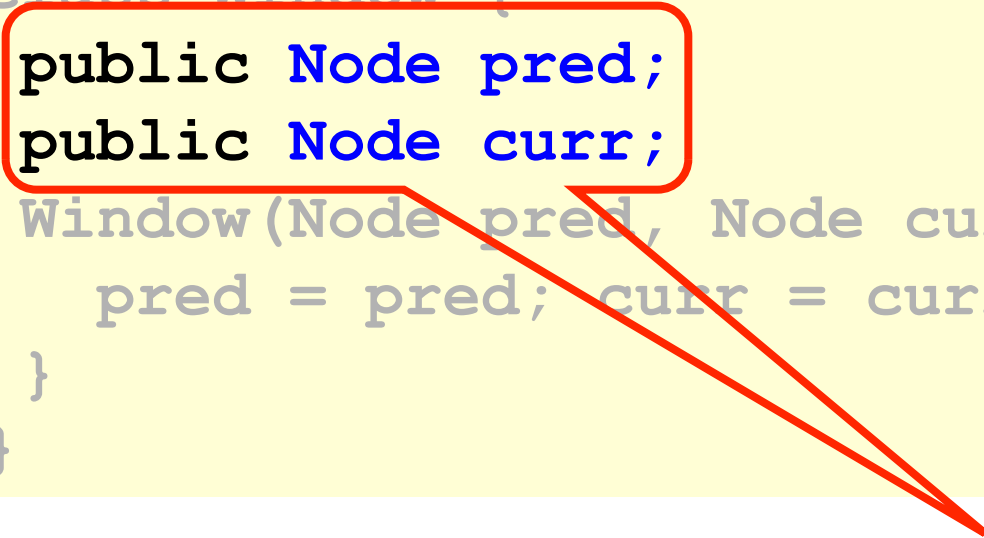
# The Window Class

```
class Window {  
    public Node pred;  
    public Node curr;  
    Window(Node pred, Node curr) {  
        pred = pred; curr = curr;  
    }  
}
```



# The Window Class

```
class Window {  
    public Node pred;  
    public Node curr;  
    Window(Node pred, Node curr) {  
        pred = pred; curr = curr;  
    }  
}
```



**A container for pred  
and current values**

# Using the Find Method

```
Window window = find(head, key);  
Node pred = window.pred;  
curr = window.curr;
```

# Using the Find Method

```
Window window = find(head, key);
```

```
Node pred = window.pred;
```

```
curr = window.curr;
```

**Find returns window**

# Using the Find Method

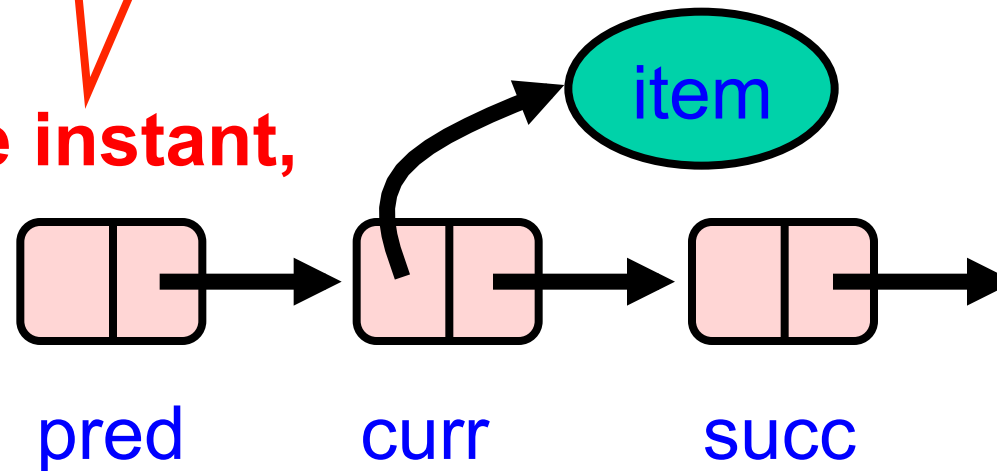
```
Window window = find(head, key);  
Node pred = window.pred;  
curr = window.curr;
```

**Extract pred and curr**

# The Find Method

```
Window window = find(item);
```

**At some instant,**

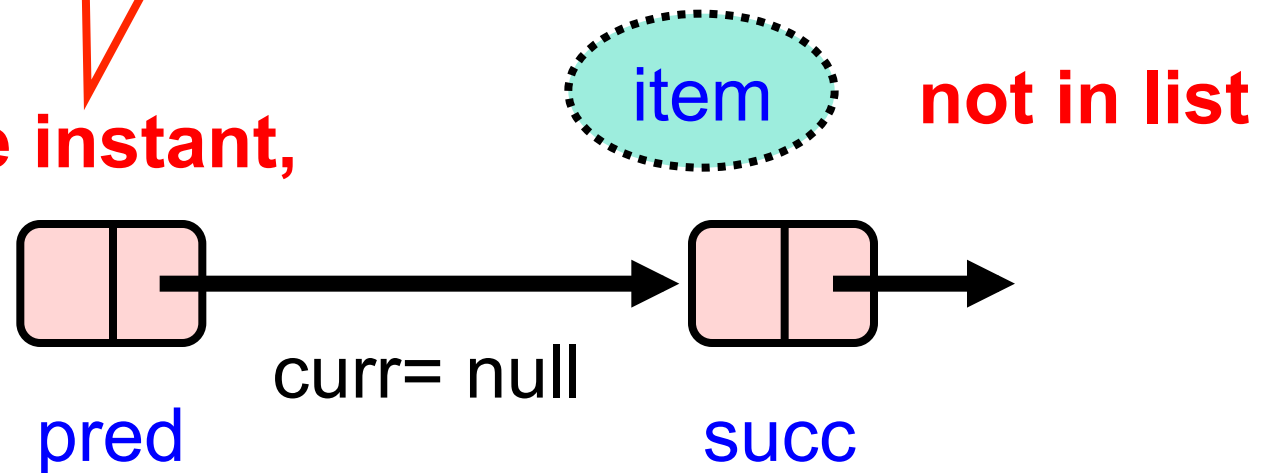


**or ...**

# The Find Method

```
Window window = find(item);
```

At some instant,



# Remove

```
public boolean remove(T item) {
    Boolean snip;
    while (true) {
        Window window = find(head, key);
        Node pred = window.pred, curr = window.curr;
        if (curr.key != key) {
            return false;
        } else {
            Node succ = curr.next.getReference();
            snip = curr.next.compareAndSet(succ, succ, false
true);
            if (!snip) continue;
            pred.next.compareAndSet(curr, succ, false, false);
            return true;
        }
    }
}
```

# Remove

```
public boolean remove(T item) {  
    Boolean snip;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key != key) {  
            return false;  
        } else {  
            Node succ = curr.next.getReference();  
            snip = curr.next.compareAndSet (succ, succ, false,  
true);  
            if (!snip) continue;  
            pred.next.compareAndSet(curr, succ, false, false);  
            return true;  
        }  
    }  
}
```

**Keep trying**



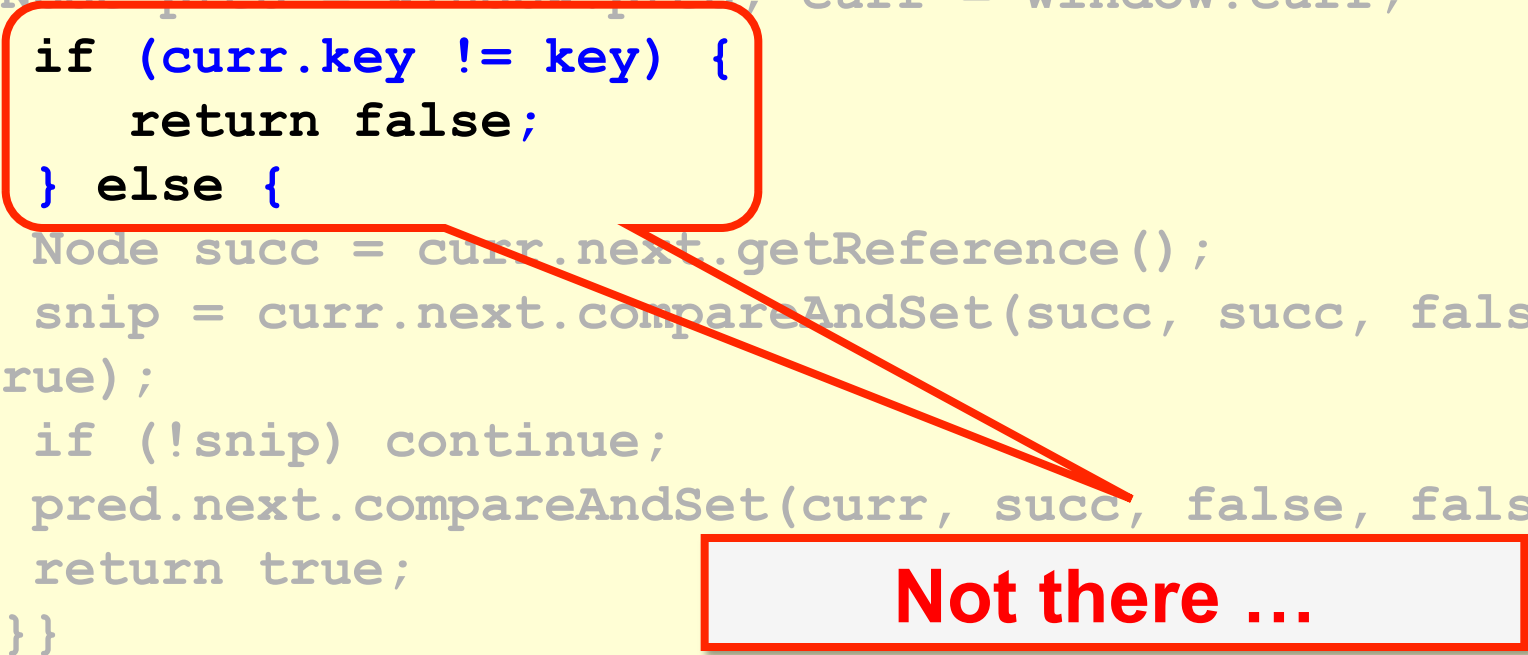
# Remove

```
public boolean remove(T item) {  
    Boolean snip;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key != key) {  
            return false;  
        } else {  
            Node succ = curr.next.getReference();  
            snip = curr.next.compareAndSet(succ, succ, false,  
true);  
            if (!snip) continue;  
            pred.next.compareAndSet(curr, succ, false, false);  
            return true;  
        }  
    }  
}
```

**Find neighbors**

# Remove

```
public boolean remove(T item) {
    Boolean snip;
    while (true) {
        Window window = find(head, key);
        Node pred = window.pred, curr = window.curr;
        if (curr.key != key) {
            return false;
        } else {
            Node succ = curr.next.getReference();
            snip = curr.next.compareAndSet(succ, succ, false,
true);
            if (!snip) continue;
            pred.next.compareAndSet(curr, succ, false, false);
            return true;
        }
    }
}
```



**Not there ...**

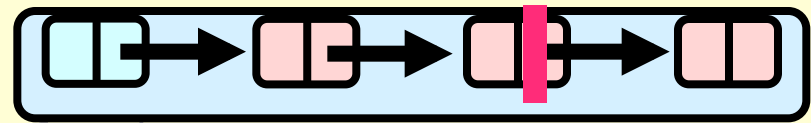
# Remove

**Try to mark node as deleted**

```
public boolean remove(T item) {  
    Boolean snip = false;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key != key) {  
            return false;  
        } else {  
            Node succ = curr.next.getReference();  
            snip = curr.next.compareAndSet(succ, succ, false,  
true);  
            if (!snip) continue;  
            pred.next.compareAndSet(curr, succ, false, false);  
            return true;  
        }  
    }  
}
```

# Remove

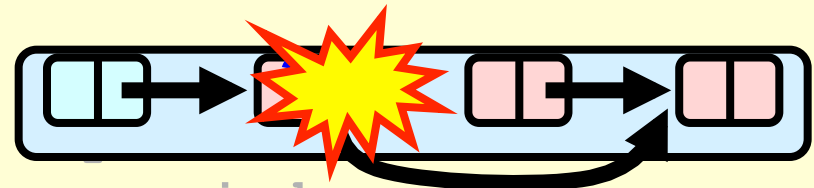
If it doesn't work,  
just retry, if it  
does, job  
essentially done



```
public boolean remove(T item) {  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (pred != null && pred.getKey() == item) {  
            return false;  
        } else {  
            Node succ = curr.next.getReference();  
            snip = curr.next.compareAndSet(succ, succ, false, true);  
            if (!snip) continue;  
            pred.next.compareAndSet(curr, succ, false, false);  
            return true;  
        }  
    }  
}
```

# Remove

```
public boolean remove(T item) {  
    Boolean snip;  
    while (true) {  
        Window window = find(head,
```



```
curr = window.curr;
```

**Try to advance reference  
(if we don't succeed,  
someone else did or will).**

```
    Node succ = curr.next.getReference();  
    snip = curr.next.compareAndSet(succ, succ, false,  
true);  
    if (!snip) continue;  
    pred.next.compareAndSet(curr, succ, false, false);  
    return true;  
}}}
```

# Remove

```
public boolean remove(T item) {
    Boolean snip;
    while (true) {
        Window window = find(head, key);
        Node pred = window.pred, curr = window.curr;
        if (curr.key != key) {
            return false;
        } else {
            Node succ = curr.next.getReference();
            snip = curr.next.compareAndSet(succ, succ, false,
true);
            if (!snip) continue;
            pred.next.compareAndSet(curr, succ, false, false);
            return true;
        }
    }
}
```

Linearization point if  
removal is successful

# Remove

```
public boolean remove(T item) {
    Boolean snip;
    while (true) {
        Window window = find(head, key);
        Node pred = window.pred, curr = window.curr;
        if (curr.key != key) {
            return false;
        } else {
            Node succ = curr.next.getReference();
            snip = curr.next.compareAndSet(succ, succ, false
true);
            if (!snip) continue;
            pred.next.compareAndSet(curr, succ, false, false);
            return true;
        }
    }
}
```

Linearization point is when we found this node (in Find()) if removal returns false.

# Add

```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key == key) {  
            return false;  
        } else {  
            Node node = new Node(item);  
            node.next = new AtomicMarkableRef(curr, false);  
            if (pred.next.compareAndSet(curr, node, false,  
false)) {return true;}  
        }  
    }  
}
```



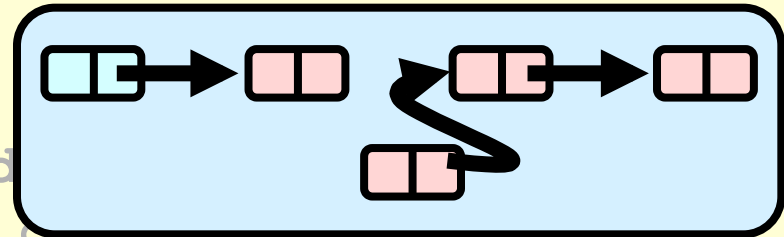
# Add

```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key == key) {  
            return false;  
        } else {  
            Node node = new Node(item);  
            node.next = new AtomicMarkableRef(curr, false);  
            if (pred.next.compareAndSet(curr, node, false,  
false)) {return true;}  
        }  
    }  
}
```

**Item already there**

# Add

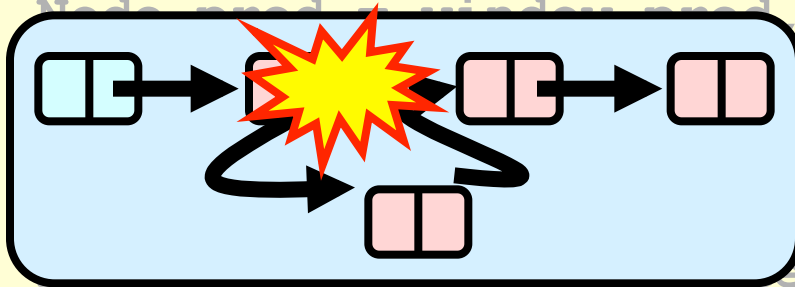
```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        Window window = find(head  
        Node pred = window.pred, curr = window.curr,  
        if (curr.key == key) {  
            return false;  
        } else {  
            Node node = new Node(item);  
            node.next = new AtomicMarkableRef(curr, false);  
            if (pred.next.compareAndSet(curr, node, false,  
            false)) {return true;}  
        }  
    }  
}
```



**create new node**

# Add

```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        Window window = find(head, item);  
        Node pred = window.pred; curr = window.curr;
```



**Install new node,  
else retry loop**

```
        node.next = new AtomicMarkableRef(curr, false);  
        if (pred.next.compareAndSet(curr, node, false,  
false)) {return true;}  
    }  
}
```

# Add

```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key == key) {  
            return false;  
        } else {  
            Node node = new Node(item);  
            node.next = new AtomicMarkableRef(curr, false);  
            if (pred.next.compareAndSet(curr, node, false,  
false)) {return true;}  
        }  
    }  
}
```

**Linearization point if  
add is successful**

# Add

```
public boolean add(T item) {  
    boolean splice;  
    while (true) {  
        Window window = find(head, key);  
        Node pred = window.pred, curr = window.curr;  
        if (curr.key == key) {  
            return false;  
        } else {  
            Node node = new Node(item);  
            node.next = new AtomicMarkableRef(curr, false);  
            if (pred.next.compareAndSet(curr, node, false,  
false)) {return true;}  
        }  
    }  
}
```

Linearization point is  
when we found this  
node (in Find()) if  
removal returns false.

# Wait-free Contains

```
public boolean contains(T item) {  
    boolean marked;  
    int key = item.hashCode();  
    Node curr = head;  
    while (curr.key < key)  
        curr = curr.next;  
    Node succ = curr.next.get(marked);  
    return (curr.key == key && !marked[0])  
}
```

# Wait-free Contains

```
public boolean contains(T item) {  
    boolean marked;  
    int key = item.hashCode();  
    Node curr = head;  
    while (curr.key < key)  
        curr = curr.next;  
    Node succ = curr.next.get(marked);  
    return (curr.key == key && !marked[0])  
}
```

**Only difference is  
that we get and  
check marked**

# Wait-free Contains

```
public boolean contains(T item) {  
    boolean marked;  
    int key = item.hashCode();  
    Node curr = head;  
    while (curr.key < key)  
        curr = curr.next;  
    Node succ = curr.next.get(marked);  
    return (curr.key == key && !marked[0])  
}
```



# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null, curr = null, succ = null;
    boolean[] marked = {false}; boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getReference();
        while (true) {
            succ = curr.next.get(marked);
            while (marked[0]) {
                ...
            }
            if (curr.key >= key)
                return new Window(pred, curr);
            pred = curr;
            curr = succ;
        }
    }
}
```

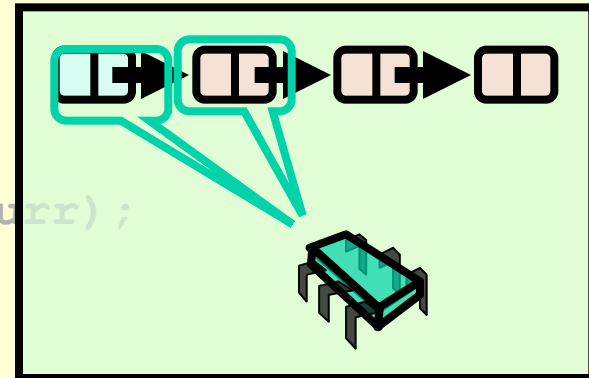
# Lock-free Find

```
public Window find(Node head, int key) {  
    Node pred = null, curr = null, succ = null;  
    boolean[] marked = {false}; boolean snip;  
    retry: while (true) {  
        pred = head;  
        curr = pred.next.getReference();  
        while (true) {  
            succ = curr.next.get(marked);  
            while (marked[0]) {  
                ...  
            }  
            if (curr.key >= key)  
                return new Window(pred, curr);  
            pred = curr;  
            curr = succ;  
        }  
    }  
}
```

**If list changes  
while traversed,  
start over**

# Lock-free Find

```
public Window find(Node head, int key) {  
    Node pred = null; Start looking from head  
    boolean[] marked = {false}; boolean snip;  
    retry: while (true) {  
        pred = head;  
        curr = pred.next.getReference();  
        while (true) {  
            succ = curr.next.get(marked);  
            while (marked[0]) {  
                ...  
            }  
            if (curr.key >= key)  
                return new Window(pred, curr);  
            pred = curr;  
            curr = succ;  
        }  
    }  
}
```



# Lock-free Find

```
public Window find(Node head, int key) {  
    Node pred = null, curr = null, succ = null;  
    boolean[] marked = {false}; boolean snip;  
    retry: while (true) { Move down the list  
        pred = head;  
        curr = pred.next.getReference();  
        while (true) {  
            succ = curr.next.get(marked);  
            while (marked[0]) {  
                ...  
            }  
            if (curr.key >= key)  
                return new Window(pred, curr);  
            pred = curr;  
            curr = succ;  
        }  
    }  
}
```

# Lock-free Find

```
public Window find(Node head, int key) {  
    Node pred = null, curr = null, succ = null;  
    boolean[] marked = {false}; boolean snip;  
    retry: while (true) {  
        pred = head;  
        curr = pred.next.getReference();  
        while (true) {  
            succ = curr.next.get(marked);  
            while (marked[0]) {  
                ...  
            }  
            if (curr.key >= key)  
                return new Window(pred, curr);  
            pred = curr;  
            curr = succ;  
        }  
    }  
}
```

**Get ref to successor and  
current deleted bit**

# Lock-free Find

```
public Window find(Node head, int key) {
    Node pred = null, curr = null, succ = null;
    boolean[] marked = {false}; boolean snip;
    retry: while (true) {
        pred = head;
        curr = pred.next.getReference();
        while (true) {
            succ = curr.next.get(marked);
            while (marked[0]) {
                ...
            }
            if (curr.key >= key)
                return new Window(pred, curr);
            pred = curr;
        }
    }
}
```

**Try to remove deleted nodes in  
path...code details soon**

# Lock-free Find

```
public Window find(Node head, int key) {  
    Node pred = null, curr = null, succ = null;  
    boolean[] marked = {false}; boolean snip;  
    retry: while (true) {  
        pred = head;  
        curr = pred.next.getReference();  
        If curr key that is greater or  
        equal, return pred and curr  
        while (marked[0]) {  
            ...  
        }  
        if (curr.key >= key)  
            return new Window(pred, curr);  
        pred = curr;  
        curr = succ;  
    }  
}
```

# Lock-free Find

```
public Window find(Node head, int key) {  
    Node pred = null, curr = null, succ = null;  
    boolean[] marked = {false}; boolean snip;  
    retry: while (true) {  
        pred = head;  
        curr = pred.next.getReference();  
        while (true) {  
            succ = curr.next.getReference();  
            while (marked[0]) {  
                ...  
            }  
            if (curr.key >= key)  
                return new Window(pred, curr);  
            pred = curr;  
            curr = succ;  
        }  
    }  
}
```

**Otherwise advance window and  
loop again**



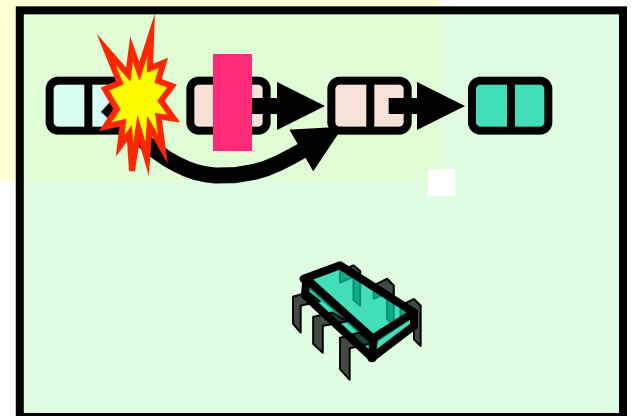
# Lock-free Find

```
retry: while (true) {  
    ...  
    while (marked[0]) {  
        snip = pred.next.compareAndSet(curr,  
                                         succ, false, false);  
        if (!snip) continue retry;  
        curr = succ;  
        succ = curr.next.get(marked);  
    }  
    ...  
}
```

# Lock-free Find

Try to snip out node

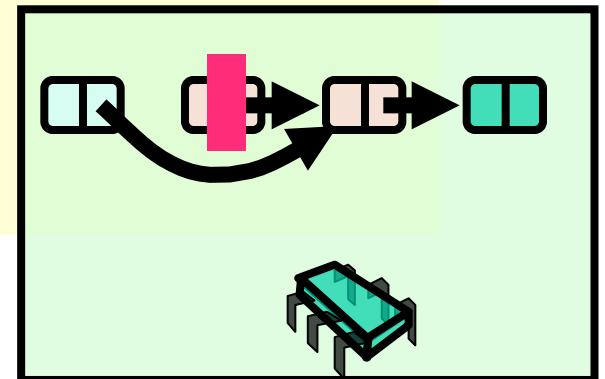
```
retry: while (true) {  
    ...  
    while (marked[0]) {  
        snip = pred.next.compareAndSet(curr,  
                                         succ, false, false);  
        if (!snip) continue retry;  
        curr = succ;  
        succ = curr.next.get(marked);  
    }  
    ...  
}
```



# Lock-free Find

if predecessor's next field changed,  
retry whole traversal

```
retry: while (true) {  
    ...  
    while (marked[0]) {  
        snip = pred.next.compareAndSet(curr,  
                                         succ, false, false);  
        if (!snip) continue retry;  
        curr = succ;  
        succ = curr.next.get(marked);  
    }  
    ...  
}
```



# Lock-free Find

Otherwise move on to check  
if next node deleted


```
retry: while (true) {  
    ...  
    while (marked[0]) {  
        snip = pred.next.compareAndSet(curr,  
                                         succ, false, false);  
        if (!snip) continue retry;  
        curr = succ;  
        succ = curr.next.get(marked) ;  
    }  
    ...  
}
```

# Performance

- Different list-based set implementations
- SunFire 6800 (bus based cache coherence)
- 16-node machine, each 1.2 GHz
- Vary percentage of **contains()** calls

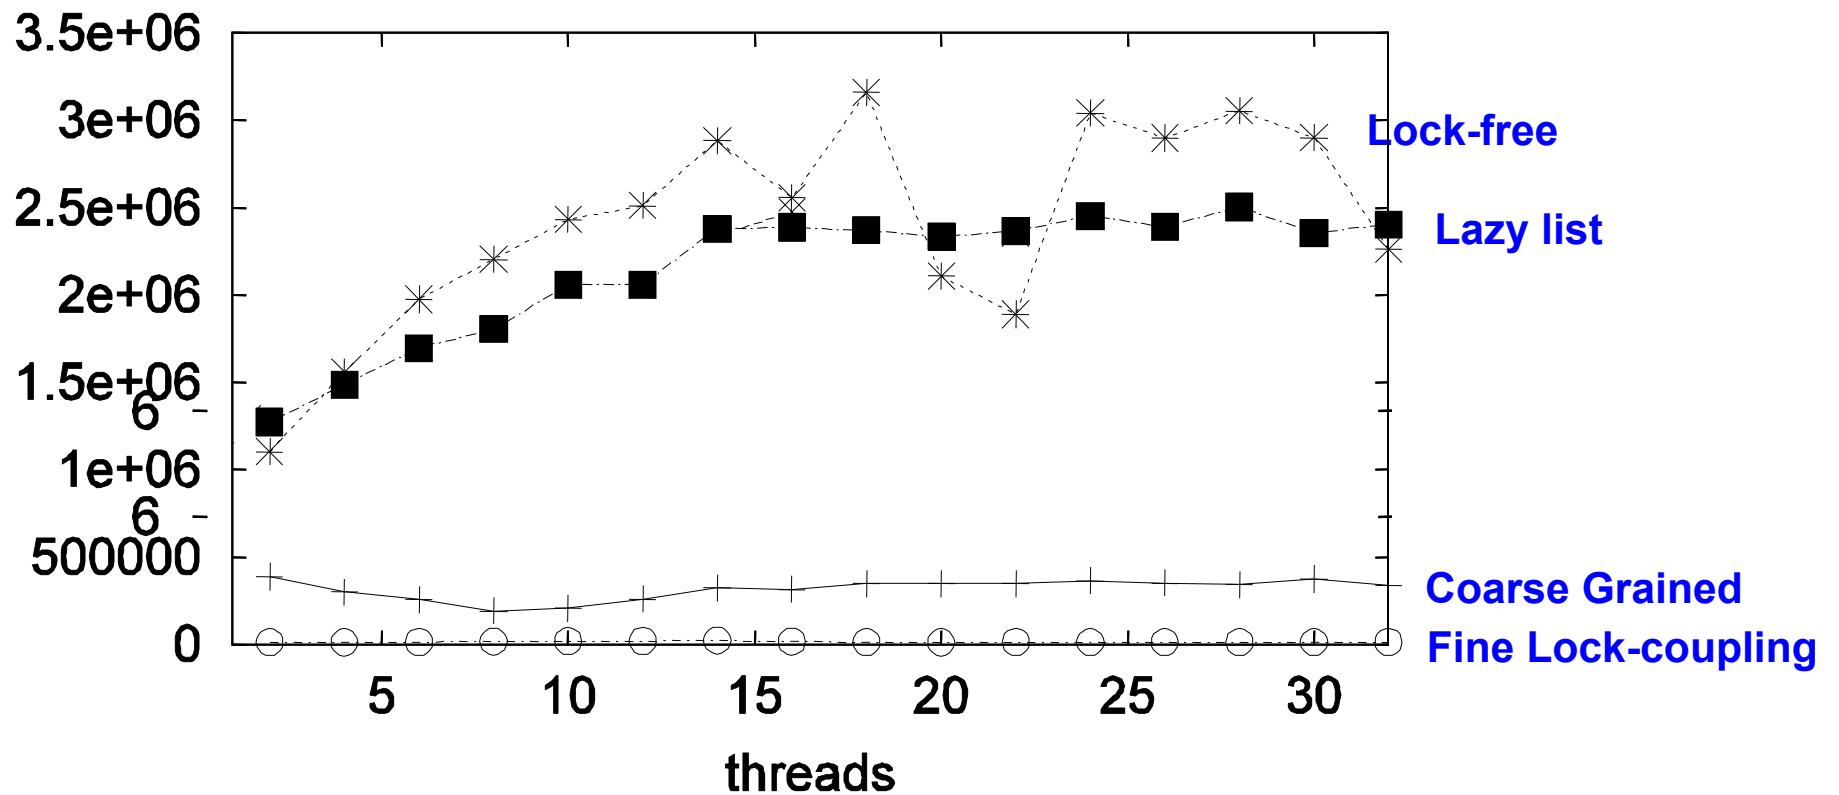
# High Contains Ratio

Ops/sec (90% contain, 9% add, 1% remove)

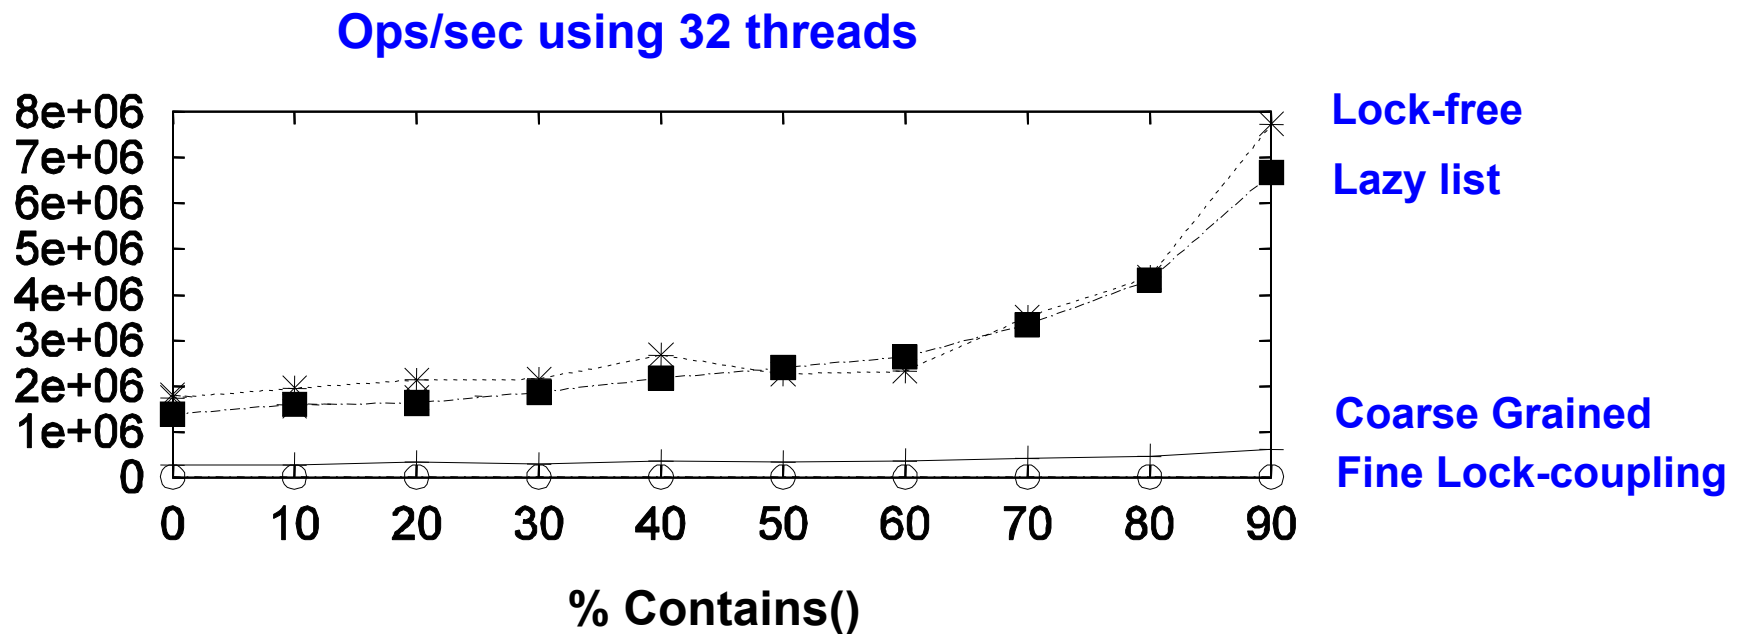


# Low Contains Ratio

Ops/sec (50% contain, 45% add, 5% remove)



# As Contains Ratio Increases





# “To Lock or Not to Lock”

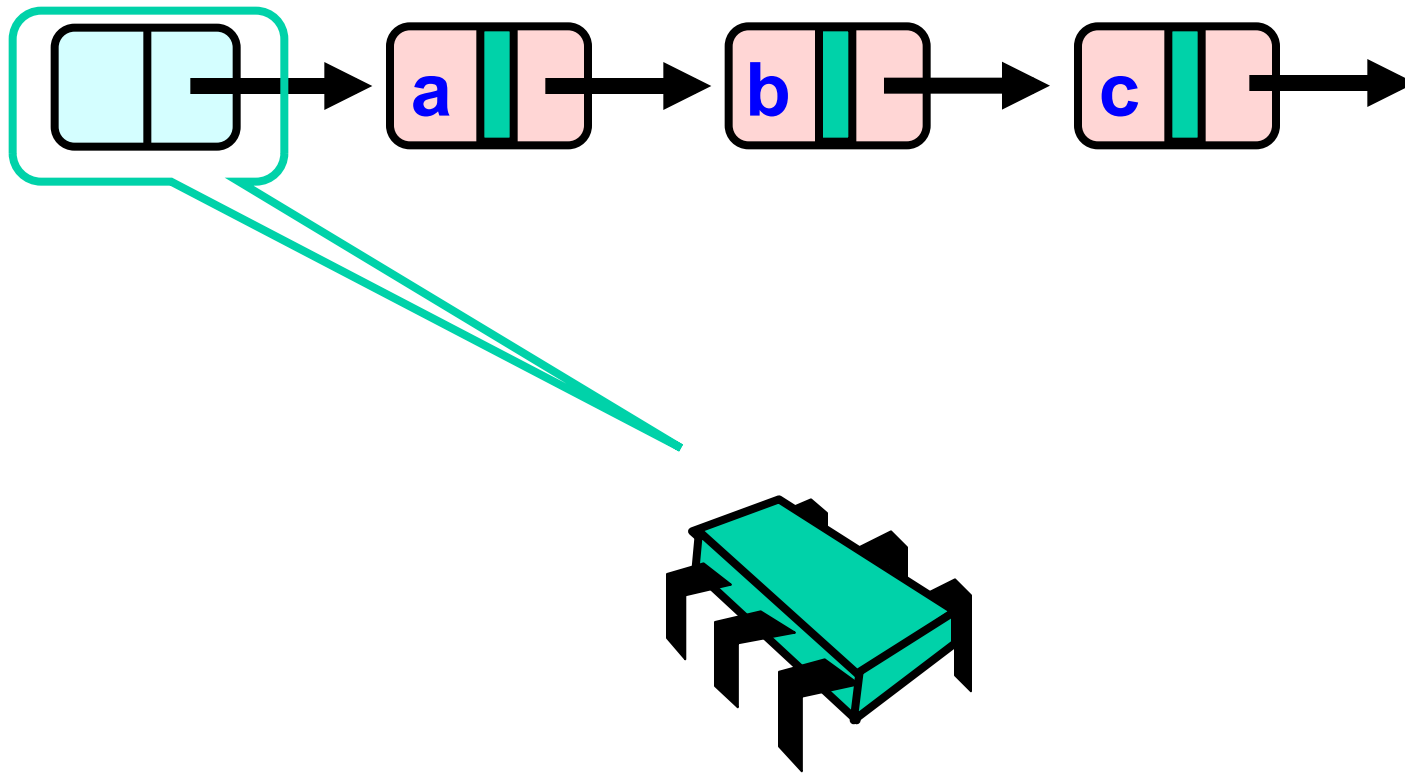
- Locking vs. Non-blocking:
  - Depending on the application usage
- The answer: nobler to compromise
  - Example: Lazy list combines blocking **add()** and **remove()** and a wait-free **contains()**
  - Remember: Blocking/non-blocking is a property of a method



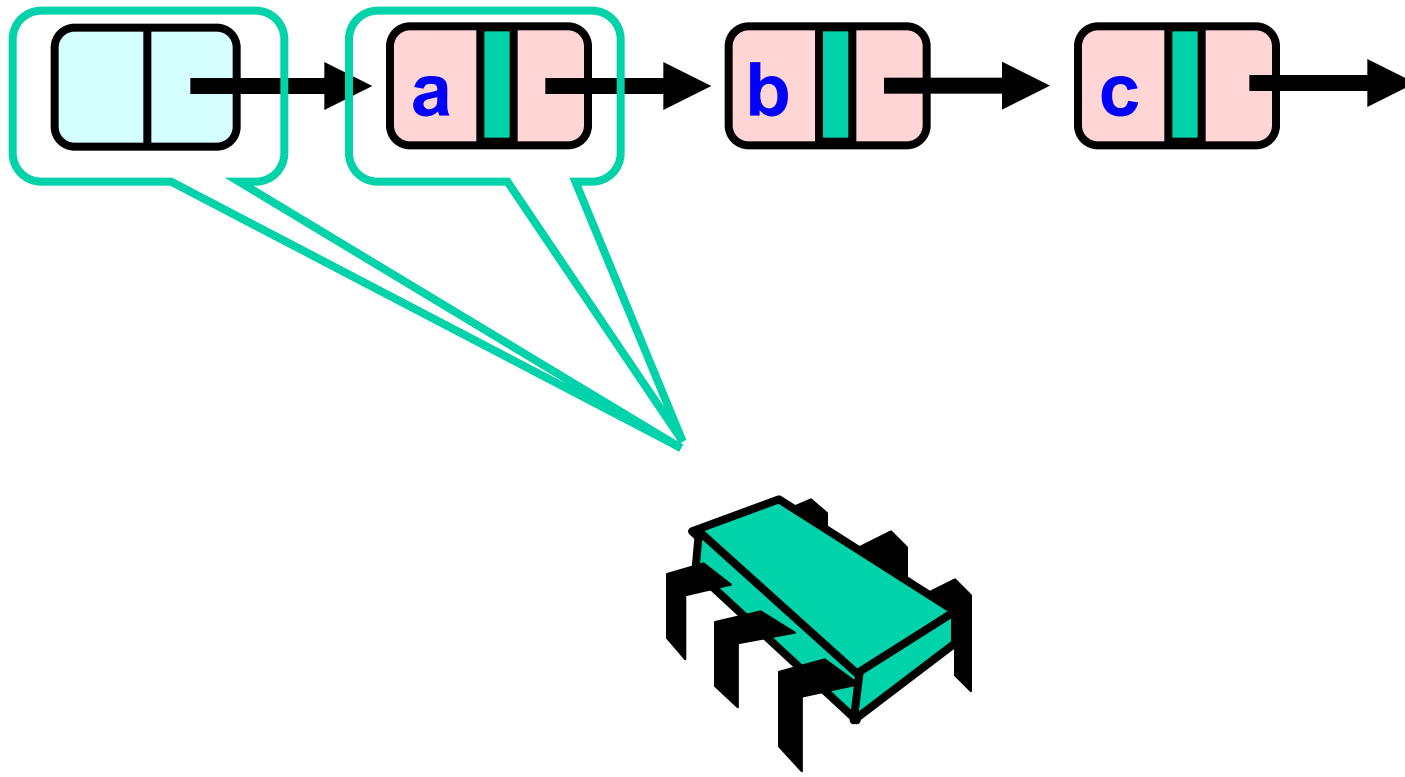
## This work is licensed under a Creative Commons Attribution-ShareAlike 2.5 License.

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.

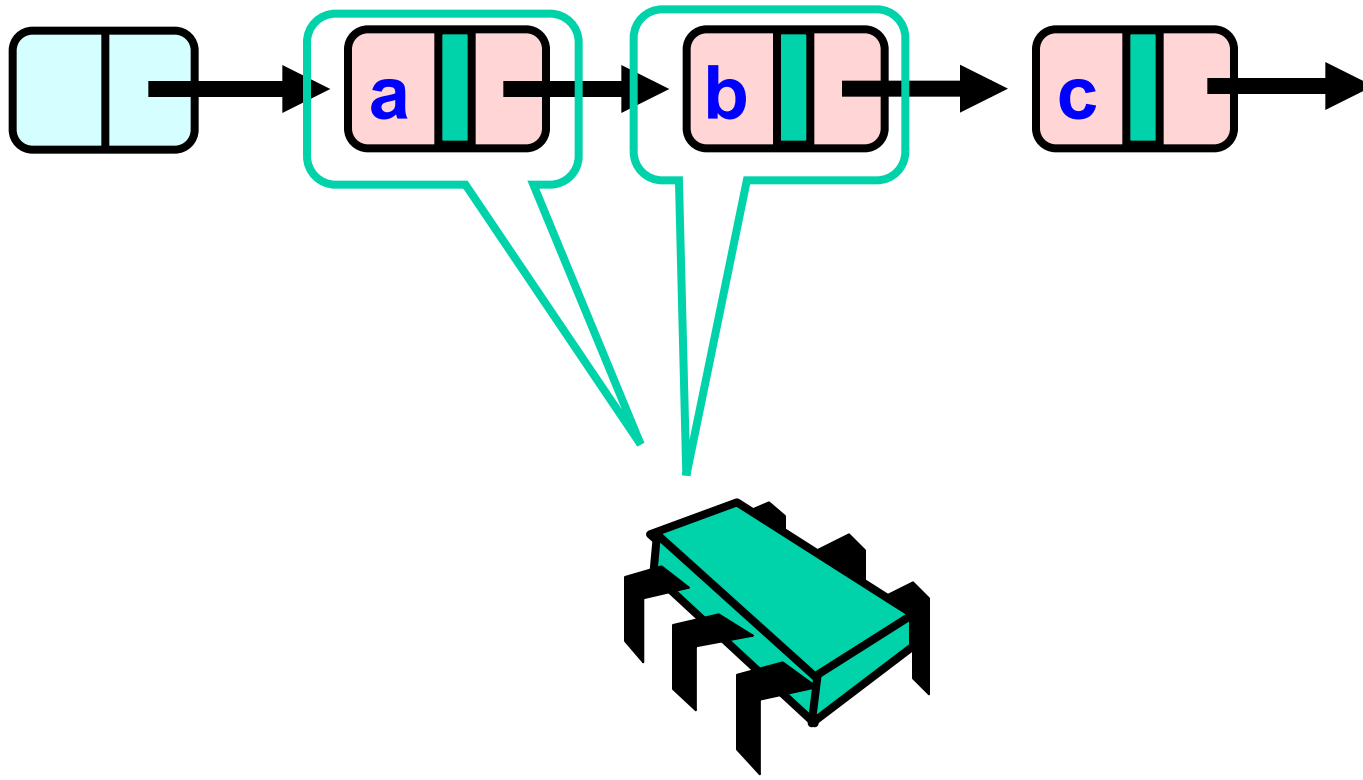
# Business as Usual



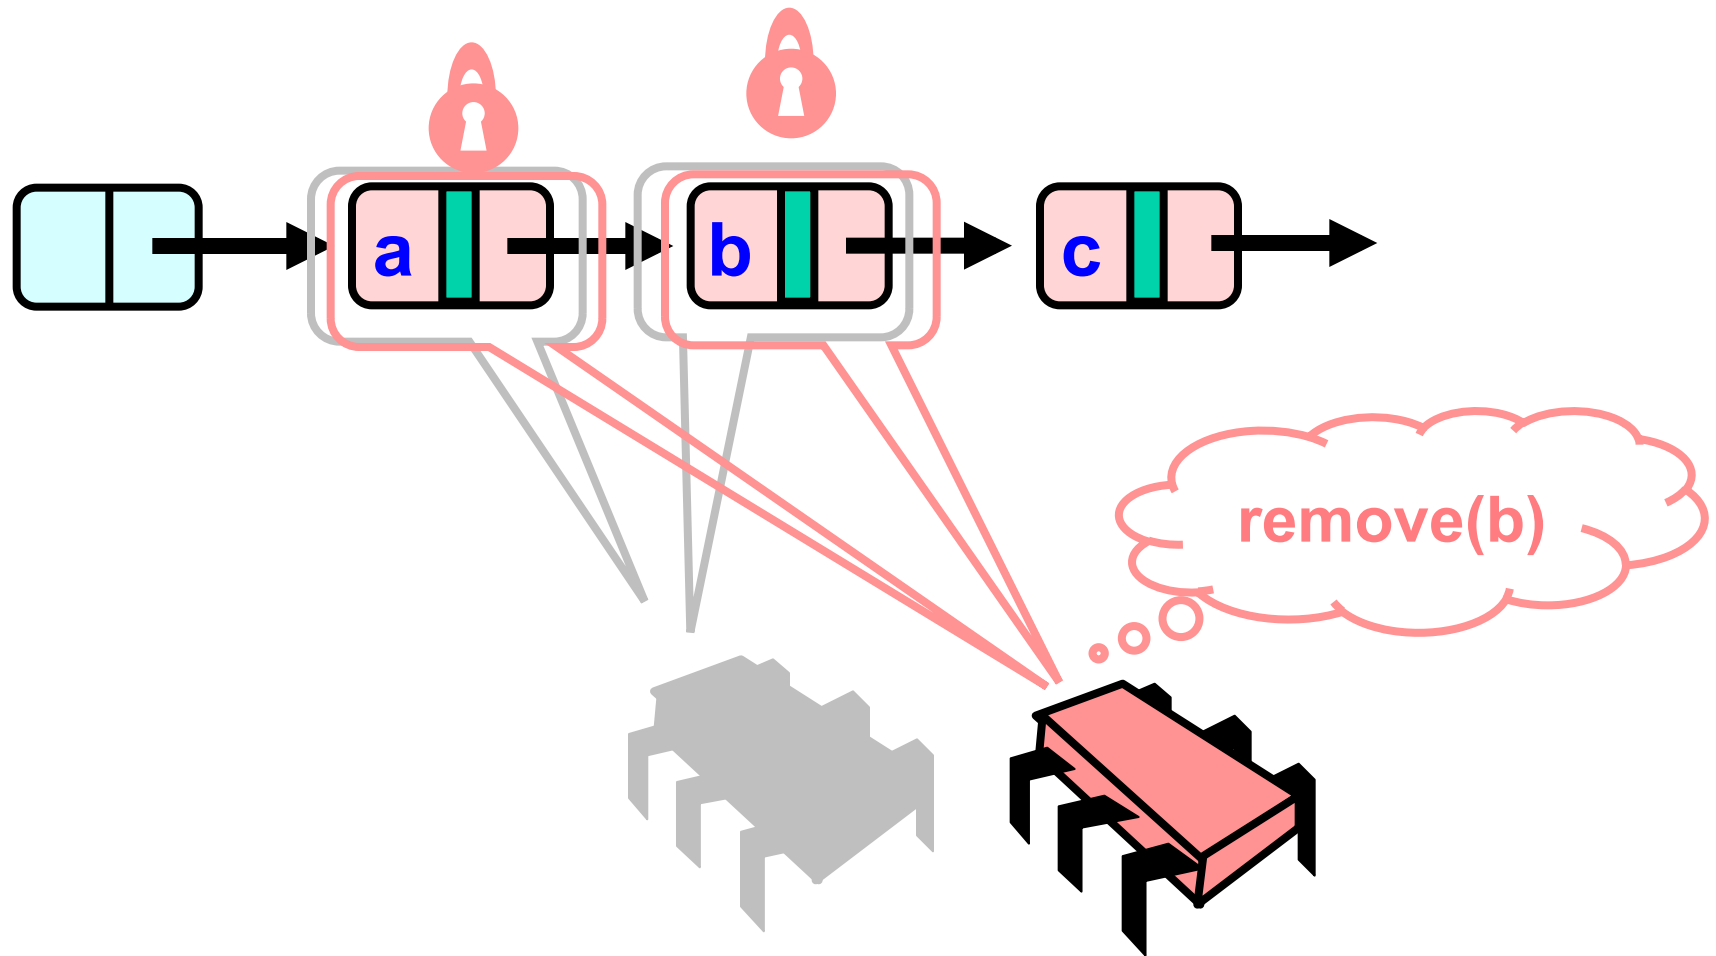
# Business as Usual



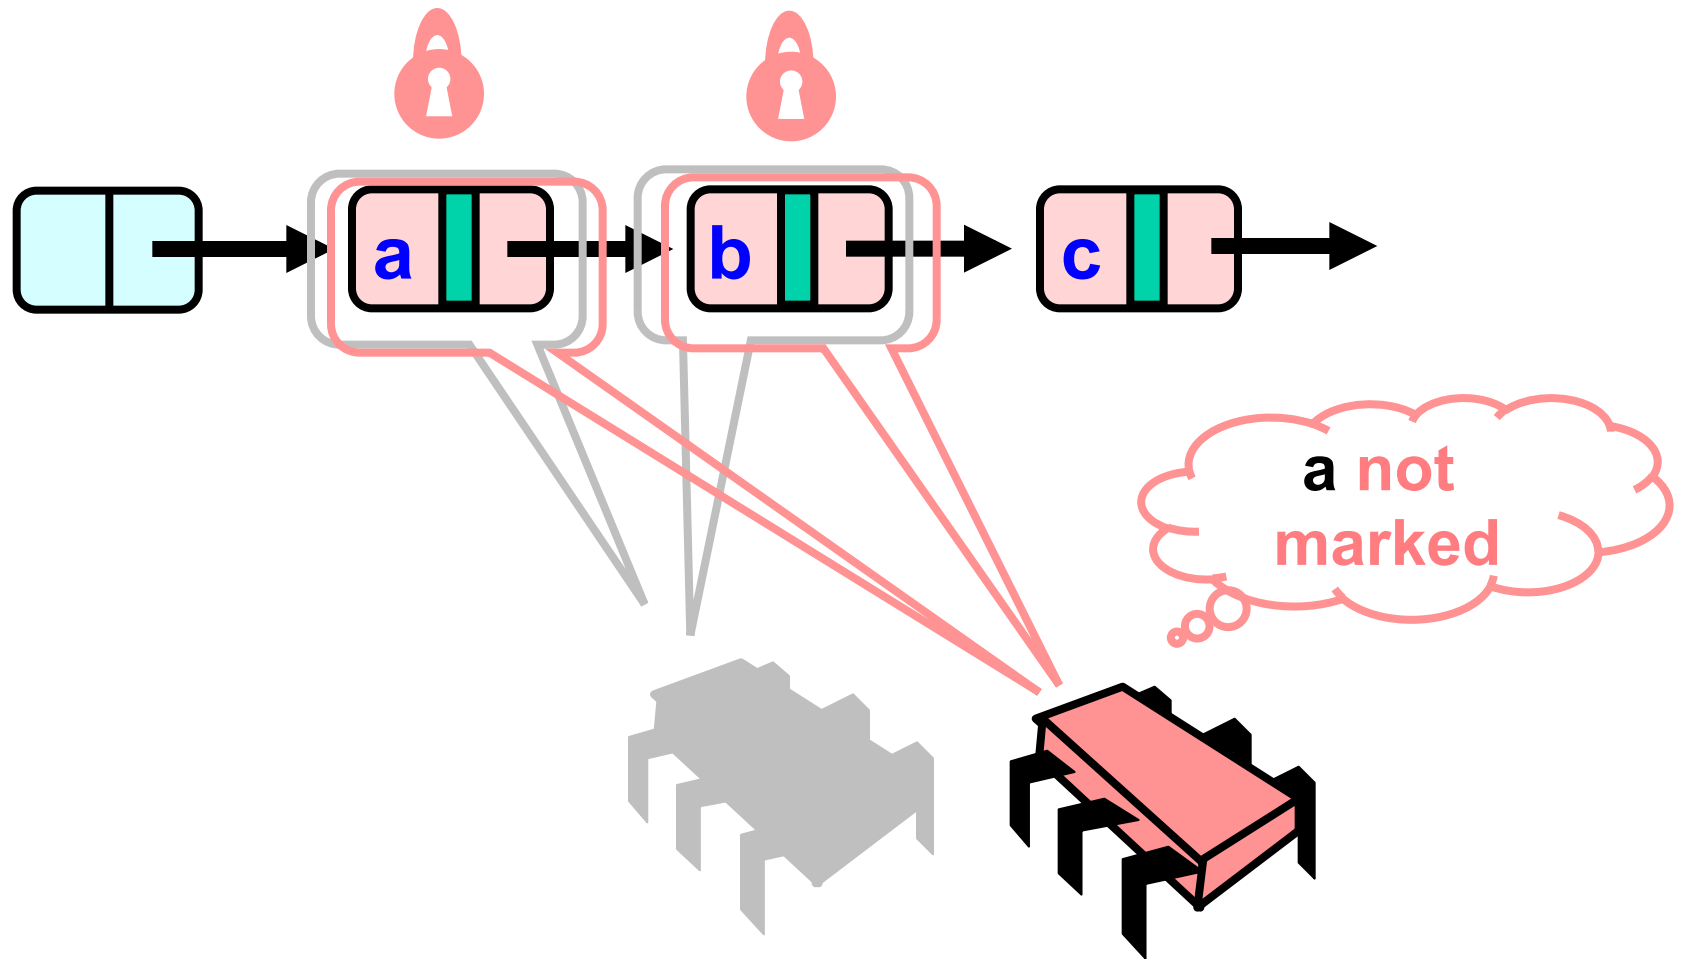
# Business as Usual




# Business as Usual



# Business as Usual




# Business as Usual

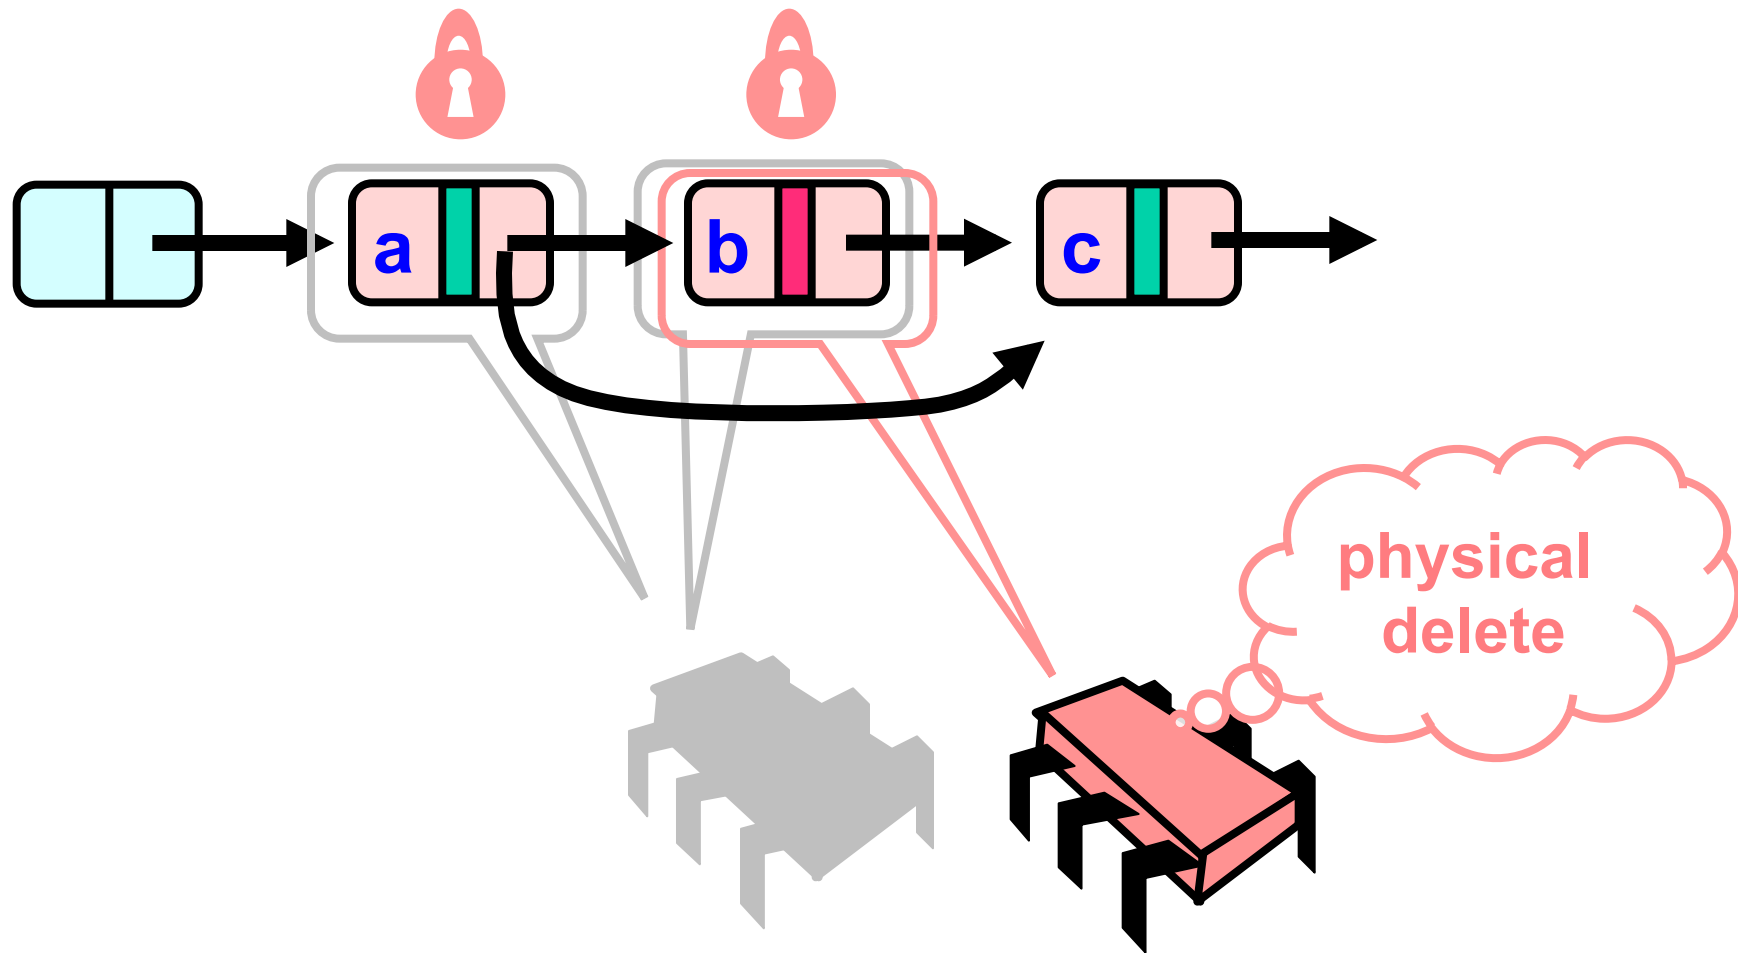




# Business as Usual



# Business as Usual



# Business as Usual

